
Non-Linear Least-Squares Minimization and Curve-Fitting for Python

Release 0.9.11

Matthew Newville, Till Stensitzki, and others

Jun 29, 2018

CONTENTS

| | | |
|----------|---|-----------|
| 1 | Getting started with Non-Linear Least-Squares Fitting | 3 |
| 2 | Downloading and Installation | 7 |
| 2.1 | Prerequisites | 7 |
| 2.2 | Downloads | 7 |
| 2.3 | Installation | 7 |
| 2.4 | Development Version | 8 |
| 2.5 | Testing | 8 |
| 2.6 | Acknowledgements | 8 |
| 2.7 | License | 9 |
| 3 | Getting Help | 11 |
| 4 | Frequently Asked Questions | 13 |
| 4.1 | What's the best way to ask for help or submit a bug report? | 13 |
| 4.2 | Why did my script break when upgrading from lmfit 0.8.3 to 0.9.0? | 13 |
| 4.3 | I get import errors from IPython | 13 |
| 4.4 | How can I fit multi-dimensional data? | 13 |
| 4.5 | How can I fit multiple data sets? | 14 |
| 4.6 | How can I fit complex data? | 14 |
| 4.7 | Can I constrain values to have integer values? | 14 |
| 4.8 | How should I cite LMFIT? | 14 |
| 5 | Parameter and Parameters | 15 |
| 5.1 | The <code>Parameter</code> class | 15 |
| 5.2 | The <code>Parameters</code> class | 17 |
| 5.3 | Simple Example | 20 |
| 6 | Performing Fits and Analyzing Outputs | 23 |
| 6.1 | The <code>minimize()</code> function | 23 |
| 6.2 | Writing a Fitting Function | 25 |
| 6.3 | Choosing Different Fitting Methods | 26 |
| 6.4 | <code>MinimizerResult</code> – the optimization result | 27 |
| 6.5 | Using a Iteration Callback Function | 30 |
| 6.6 | Using the <code>Minimizer</code> class | 30 |
| 6.7 | <code>Minimizer.emcee()</code> - calculating the posterior probability distribution of parameters | 40 |
| 6.8 | Getting and Printing Fit Reports | 43 |
| 7 | Modeling Data and Curve Fitting | 47 |
| 7.1 | Motivation and simple example: Fit data to Gaussian profile | 47 |
| 7.2 | The <code>Model</code> class | 50 |

| | | |
|-----------|--|------------|
| 7.3 | The <code>ModelResult</code> class | 60 |
| 7.4 | Composite Models : adding (or multiplying) Models | 69 |
| 8 | Built-in Fitting Models in the <code>models</code> module | 77 |
| 8.1 | Peak-like models | 77 |
| 8.2 | Linear and Polynomial Models | 87 |
| 8.3 | Step-like models | 89 |
| 8.4 | Exponential and Power law models | 91 |
| 8.5 | User-defined Models | 93 |
| 8.6 | Example 1: Fit Peak data to Gaussian, Lorentzian, and Voigt profiles | 94 |
| 8.7 | Example 2: Fit data to a Composite Model with pre-defined models | 98 |
| 8.8 | Example 3: Fitting Multiple Peaks – and using Prefixes | 99 |
| 9 | Calculation of confidence intervals | 105 |
| 9.1 | Method used for calculating confidence intervals | 105 |
| 9.2 | A basic example | 105 |
| 9.3 | An advanced example | 106 |
| 9.4 | Confidence Interval Functions | 110 |
| 10 | Bounds Implementation | 113 |
| 11 | Using Mathematical Constraints | 115 |
| 11.1 | Overview | 115 |
| 11.2 | Supported Operators, Functions, and Constants | 115 |
| 11.3 | Using Inequality Constraints | 116 |
| 11.4 | Advanced usage of Expressions in <code>lmfit</code> | 117 |
| 12 | Release Notes | 119 |
| 12.1 | Version 0.9.10 Release Notes | 119 |
| 12.2 | Version 0.9.9 Release Notes | 119 |
| 12.3 | Version 0.9.6 Release Notes | 119 |
| 12.4 | Version 0.9.5 Release Notes | 120 |
| 12.5 | Version 0.9.4 Release Notes | 120 |
| 12.6 | Version 0.9.3 Release Notes | 120 |
| 12.7 | Version 0.9.0 Release Notes | 120 |
| | Python Module Index | 123 |
| | Index | 125 |

Lmfit provides a high-level interface to non-linear optimization and curve fitting problems for Python. It builds on and extends many of the optimization methods of `scipy.optimize`. Initially inspired by (and named for) extending the [Levenberg-Marquardt](#) method from `scipy.optimize.leastsq`, Lmfit now provides a number of useful enhancements to optimization and data fitting problems, including:

- Using *Parameter* objects instead of plain floats as variables. A *Parameter* has a value that can be varied during the fit or kept at a fixed value. It can have upper and/or lower bounds. A Parameter can even have a value that is constrained by an algebraic expression of other Parameter values. As a Python object, a Parameter can also have attributes such as a standard error, after a fit that can estimate uncertainties.
- Ease of changing fitting algorithms. Once a fitting model is set up, one can change the fitting algorithm used to find the optimal solution without changing the objective function.
- Improved estimation of confidence intervals. While `scipy.optimize.leastsq` will automatically calculate uncertainties and correlations from the covariance matrix, the accuracy of these estimates is sometimes questionable. To help address this, Lmfit has functions to explicitly explore parameter space and determine confidence levels even for the most difficult cases.
- Improved curve-fitting with the *Model* class. This extends the capabilities of `scipy.optimize.curve_fit`, allowing you to turn a function that models your data into a Python class that helps you parametrize and fit data with that model.
- Many *built-in models* for common lineshapes are included and ready to use.

The Lmfit package is Free software, using an Open Source license. The software and this document are works in progress. If you are interested in participating in this effort please use the [Lmfit GitHub repository](#).

GETTING STARTED WITH NON-LINEAR LEAST-SQUARES FITTING

The `lmfit` package provides simple tools to help you build complex fitting models for non-linear least-squares problems and apply these models to real data. This section gives an overview of the concepts and describes how to set up and perform simple fits. Some basic knowledge of Python, NumPy, and modeling data are assumed – this is not a tutorial on why or how to perform a minimization or fit data, but is rather aimed at explaining how to use `lmfit` to do these things.

In order to do a non-linear least-squares fit of a model to data or for any other optimization problem, the main task is to write an *objective function* that takes the values of the fitting variables and calculates either a scalar value to be minimized or an array of values that are to be minimized, typically in the least-squares sense. For many data fitting processes, the latter approach is used, and the objective function should return an array of (data-model), perhaps scaled by some weighting factor such as the inverse of the uncertainty in the data. For such a problem, the chi-square (χ^2) statistic is often defined as:

$$\chi^2 = \sum_i^N \frac{[y_i^{\text{meas}} - y_i^{\text{model}}(\mathbf{v})]^2}{\epsilon_i^2}$$

where y_i^{meas} is the set of measured data, $y_i^{\text{model}}(\mathbf{v})$ is the model calculation, \mathbf{v} is the set of variables in the model to be optimized in the fit, and ϵ_i is the estimated uncertainty in the data.

In a traditional non-linear fit, one writes an objective function that takes the variable values and calculates the residual array $y_i^{\text{meas}} - y_i^{\text{model}}(\mathbf{v})$, or the residual array scaled by the data uncertainties, $[y_i^{\text{meas}} - y_i^{\text{model}}(\mathbf{v})]/\epsilon_i$, or some other weighting factor.

As a simple concrete example, one might want to model data with a decaying sine wave, and so write an objective function like this:

```
from numpy import exp, sin

def residual(vars, x, data, eps_data):
    amp = vars[0]
    phaseshift = vars[1]
    freq = vars[2]
    decay = vars[3]

    model = amp * sin(x*freq + phaseshift) * exp(-x*x*decay)

    return (data-model) / eps_data
```

To perform the minimization with `scipy.optimize`, one would do this:

```
from scipy.optimize import leastsq
```

(continues on next page)

(continued from previous page)

```
vars = [10.0, 0.2, 3.0, 0.007]
out = leastsq(residual, vars, args=(x, data, eps_data))
```

Though it is wonderful to be able to use Python for such optimization problems, and the SciPy library is robust and easy to use, the approach here is not terribly different from how one would do the same fit in C or Fortran. There are several practical challenges to using this approach, including:

1. The user has to keep track of the order of the variables, and their meaning – `vars[0]` is the amplitude, `vars[2]` is the frequency, and so on, although there is no intrinsic meaning to this order.
2. If the user wants to fix a particular variable (*not* vary it in the fit), the residual function has to be altered to have fewer variables, and have the corresponding constant value passed in some other way. While reasonable for simple cases, this quickly becomes a significant work for more complex models, and greatly complicates modeling for people not intimately familiar with the details of the fitting code.
3. There is no simple, robust way to put bounds on values for the variables, or enforce mathematical relationships between the variables. In fact, the optimization methods that do provide bounds, require bounds to be set for all variables with separate arrays that are in the same arbitrary order as variable values. Again, this is acceptable for small or one-off cases, but becomes painful if the fitting model needs to change.

These shortcomings are due to the use of traditional arrays to hold the variables, which matches closely the implementation of the underlying Fortran code, but does not fit very well with Python's rich selection of objects and data structures. The key concept in `lmfit` is to define and use `Parameter` objects instead of plain floating point numbers as the variables for the fit. Using `Parameter` objects (or the closely related `Parameters` – a dictionary of `Parameter` objects), allows one to:

1. forget about the order of variables and refer to `Parameters` by meaningful names.
2. place bounds on `Parameters` as attributes, without worrying about preserving the order of arrays for variables and boundaries.
3. fix `Parameters`, without having to rewrite the objective function.
4. place algebraic constraints on `Parameters`.

To illustrate the value of this approach, we can rewrite the above example for the decaying sine wave as:

```
from numpy import exp, sin
from lmfit import minimize, Parameters

def residual(params, x, data, eps_data):
    amp = params['amp']
    phaseshift = params['phase']
    freq = params['frequency']
    decay = params['decay']

    model = amp * sin(x*freq + phaseshift) * exp(-x*x*decay)

    return (data-model) / eps_data

params = Parameters()
params.add('amp', value=10)
params.add('decay', value=0.007)
params.add('phase', value=0.2)
params.add('frequency', value=3.0)
```

(continues on next page)

(continued from previous page)

```
out = minimize(residual, params, args=(x, data, eps_data))
```

At first look, we simply replaced a list of values with a dictionary, accessed by name – not a huge improvement. But each of the named `Parameter` in the `Parameters` object holds additional attributes to modify the value during the fit. For example, `Parameters` can be fixed or bounded. This can be done during definition:

```
params = Parameters()
params.add('amp', value=10, vary=False)
params.add('decay', value=0.007, min=0.0)
params.add('phase', value=0.2)
params.add('frequency', value=3.0, max=10)
```

where `vary=False` will prevent the value from changing in the fit, and `min=0.0` will set a lower bound on that parameter's value. It can also be done later by setting the corresponding attributes after they have been created:

```
params['amp'].vary = False
params['decay'].min = 0.10
```

Importantly, our objective function remains unchanged. This means the objective function can simply express the parameterized phenomenon to be modeled, and is separate from the choice of parameters to be varied in the fit.

The `params` object can be copied and modified to make many user-level changes to the model and fitting process. Of course, most of the information about how your data is modeled goes into the objective function, but the approach here allows some external control; that is, control by the **user** performing the fit, instead of by the author of the objective function.

Finally, in addition to the `Parameters` approach to fitting data, `lmfit` allows switching optimization methods without changing the objective function, provides tools for generating fitting reports, and provides a better determination of `Parameters` confidence levels.

DOWNLOADING AND INSTALLATION

2.1 Prerequisites

The lmfit package requires [Python](#), [NumPy](#), [SciPy](#), [asteval](#), [six](#), and [uncertainties](#).

Lmfit works with Python versions 2.7, 3.4, 3.5, and 3.6. Support for Python 2.6 and 3.3 ended with lmfit versions 0.9.4 and 0.9.8, respectively. Scipy version 0.17, NumPy version 1.10, asteval version 0.9.12, six version 1.10, and uncertainties version 3.0 are the minimum requirements.

In order to run the test suite, either the [nose](#) or [pytest](#) package is required. Some functionality of lmfit requires the [emcee](#) package, some functionality will make use of the [pandas](#), [Jupyter](#), [matplotlib](#), or [dill](#) packages if available. We highly recommend each of these packages.

2.2 Downloads

The latest stable version of lmfit is 0.9.11 and is available from [PyPi](#).

2.3 Installation

With `pip` now widely available, you can install lmfit with:

```
pip install lmfit
```

Alternatively, you can download the source kit, unpack it and install with:

```
python setup.py install
```

For Anaconda Python, lmfit is not an official package, but several Anaconda channels provide it, allowing installation with (for example):

```
conda install -c GSECARS lmfit
```

or:

```
conda install -c conda-forge lmfit
```

2.4 Development Version

To get the latest development version, use:

```
git clone https://github.com/lmfit/lmfit-py.git
```

and install using:

```
python setup.py install
```

2.5 Testing

A battery of tests scripts that can be run with either the `nose` or `pytest` testing framework is distributed with `lmfit` in the `tests` folder. These are automatically run as part of the development process. For any release or any master branch from the git repository, running `pytest` or `nosetests` should run all of these tests to completion without errors or failures.

Many of the examples in this documentation are distributed with `lmfit` in the `examples` folder, and should also run for you. Some of these examples assume `matplotlib` has been installed and is working correctly.

2.6 Acknowledgements

Many people have contributed to `lmfit`. The attribution of credit `in` a project such `as` this `is` very difficult to get perfect, `and` there are no doubt important contributions missing `or` under-represented here. Please consider this file `as` part of the documentation that may have bugs that need fixing.

Some of the largest `and` most important contributions (approximately `in` order of contribution `in` size to the existing code) are from:

Matthew Newville wrote the original version `and` maintains the project.

Till Stensitzki wrote the improved estimates of confidence intervals, `and` contributed many tests, bug fixes, `and` documentation.

Renee Otten wrote the brute force method, implemented basin-hopping `and` AMPGO `global` solvers, `and` has greatly improved the code, testing, `and` documentation `and` overall project.

A. R. J. Nelson added `differential_evolution`, `emcee`, `and` greatly improved the code, `docstrings`, `and` overall project.

Antonino Ingargiola wrote much of the high level Model code `and` has provided many bug fixes `and` improvements.

Daniel B. Allan wrote much of the original version of the high level Model code, `and` many improvements to the testing `and` documentation.

Austen Fox fixed many of the built-`in` model functions `and` improved the testing `and` documentation of these.

Michal Rawlik added plotting capabilities `for` Models.

(continues on next page)

(continued from previous page)

J. J. Helmus wrote the MINUT bounds **for** leastsq, originally **in** leastsqbounds.py, **and** ported to lmfit.

E. O. Le Bigot wrote the uncertainties package, a version of which was used by lmfit **for** many years, **and is** now an external dependency.

Additional patches, bug fixes, **and** suggestions have come **from** Christoph Deil, Francois Boulogne, Thomas Caswell, Colin Brosseau, nmearl, Gustavo Pasquevich, Clemens Prescher, LiCode, Ben Gamari, Yoav Roam, Alexander Stark, Alexandre Beelen, **and** many others.

The lmfit code obviously depends on, **and** owes a very large debt to the code **in** scipy.optimize. Several discussions on the scipy-user **and** lmfit mailing lists have also led to improvements **in** this code.

2.7 License

The LMFIT-py code is distribution under the following license:

Copyright, Licensing, **and** Re-distribution

The LMFIT-py code **is** distribution under the following license:

Copyright (c) 2014 Matthew Newville, The University of Chicago
 Till Stensitzki, Freie Universitat Berlin
 Daniel B. Allen, Johns Hopkins University
 Michal Rawlik, Eidgenossische Technische Hochschule, Zurich
 Antonino Ingargiola, University of California, Los Angeles
 A. R. J. Nelson, Australian Nuclear Science **and** Technology

↪Organisation

Permission to use **and** redistribute the source code **or** binary forms of this software **and** its documentation, **with or** without modification **is** hereby granted provided that the above notice of copyright, these terms of use, **and** the disclaimer of warranty below appear **in** the source code **and** documentation, **and** that none of the names of above institutions **or** authors appear **in** advertising **or** endorsement of works derived **from this** software without specific prior written permission **from all** parties.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THIS SOFTWARE.

Some code sections have been taken **from the** scipy library whose licence **is** below.

Copyright (c) 2001, 2002 Enthought, Inc.

(continues on next page)

(continued from previous page)

All rights reserved.

Copyright (c) 2003-2016 SciPy Developers.
All rights reserved.

Redistribution **and** use **in** source **and** binary forms, **with or** without modification, are permitted provided that the following conditions are met:

- a. Redistributions of source code must retain the above copyright notice, this **list** of conditions **and** the following disclaimer.
- b. Redistributions **in** binary form must reproduce the above copyright notice, this **list** of conditions **and** the following disclaimer **in** the documentation **and/or** other materials provided **with** the distribution.
- c. Neither the name of Enthought nor the names of the SciPy Developers may be used to endorse **or** promote products derived **from this** software without specific prior written permission.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT HOLDERS OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

GETTING HELP

If you have questions, comments, or suggestions for LMFIT, please use the [mailing list](#). This provides an on-line conversation that is archived well and can be searched well with standard web searches. If you find a bug in the code or documentation, use [GitHub Issues](#) to submit a report. If you have an idea for how to solve the problem and are familiar with Python and GitHub, submitting a GitHub Pull Request would be greatly appreciated.

If you are unsure whether to use the mailing list or the Issue tracker, please start a conversation on the [mailing list](#). That is, the problem you're having may or may not be due to a bug. If it is due to a bug, creating an Issue from the conversation is easy. If it is not a bug, the problem will be discussed and then the Issue will be closed. While one *can* search through closed Issues on github, these are not so easily searched, and the conversation is not easily useful to others later. Starting the conversation on the mailing list with "How do I do this?" or "Why didn't this work?" instead of "This should work and doesn't" is generally preferred, and will better help others with similar questions. Of course, there is not always an obvious way to decide if something is a Question or an Issue, and we will try our best to engage in all discussions.

FREQUENTLY ASKED QUESTIONS

A list of common questions.

4.1 What's the best way to ask for help or submit a bug report?

See *Getting Help*.

4.2 Why did my script break when upgrading from Imfit 0.8.3 to 0.9.0?

See *Version 0.9.0 Release Notes*

4.3 I get import errors from IPython

If you see something like:

```
from IPython.html.widgets import Dropdown
ImportError: No module named 'widgets'
```

then you need to install the ipywidgets package, try: `pip install ipywidgets`.

4.4 How can I fit multi-dimensional data?

The fitting routines accept data arrays that are one dimensional and double precision. So you need to convert the data and model (or the value returned by the objective function) to be one dimensional. A simple way to do this is to use `numpy.ndarray.flatten`, for example:

```
def residual(params, x, data=None):
    ....
    resid = calculate_multidim_residual()
    return resid.flatten()
```

4.5 How can I fit multiple data sets?

As above, the fitting routines accept data arrays that are one dimensional and double precision. So you need to convert the sets of data and models (or the value returned by the objective function) to be one dimensional. A simple way to do this is to use `numpy.concatenate`. As an example, here is a residual function to simultaneously fit two lines to two different arrays. As a bonus, the two lines share the ‘offset’ parameter:

```
import numpy as np
def fit_function(params, x=None, dat1=None, dat2=None):
    model1 = params['offset'] + x * params['slope1']
    model2 = params['offset'] + x * params['slope2']

    resid1 = dat1 - model1
    resid2 = dat2 - model2
    return np.concatenate((resid1, resid2))
```

4.6 How can I fit complex data?

As with working with multi-dimensional data, you need to convert your data and model (or the value returned by the objective function) to be double precision floating point numbers. The simplest approach is to use `numpy.ndarray.view`, perhaps like:

```
import numpy as np
def residual(params, x, data=None):
    ....
    resid = calculate_complex_residual()
    return resid.view(np.float)
```

4.7 Can I constrain values to have integer values?

Basically, no. None of the minimizers in lmfit support integer programming. They all (I think) assume that they can make a very small change to a floating point value for a parameters value and see a change in the value to be minimized.

4.8 How should I cite LMFIT?

See <https://dx.doi.org/10.5281/zenodo.11813>

PARAMETER AND PARAMETERS

This chapter describes the *Parameter* object, which is a key concept of lmfit.

A *Parameter* is the quantity to be optimized in all minimization problems, replacing the plain floating point number used in the optimization routines from `scipy.optimize`. A *Parameter* has a value that can either be varied in the fit or held at a fixed value, and can have upper and/or lower bounds placed on the value. It can even have a value that is constrained by an algebraic expression of other Parameter values. Since *Parameter* objects live outside the core optimization routines, they can be used in **all** optimization routines from `scipy.optimize`. By using *Parameter* objects instead of plain variables, the objective function does not have to be modified to reflect every change of what is varied in the fit, or whether bounds can be applied. This simplifies the writing of models, allowing general models that describe the phenomenon and gives the user more flexibility in using and testing variations of that model.

Whereas a *Parameter* expands on an individual floating point variable, the optimization methods actually still need an ordered group of floating point variables. In the `scipy.optimize` routines this is required to be a one-dimensional `numpy.ndarray`. In lmfit, this one-dimensional array is replaced by a *Parameters* object, which works as an ordered dictionary of *Parameter* objects with a few additional features and methods. That is, while the concept of a *Parameter* is central to lmfit, one normally creates and interacts with a *Parameters* instance that contains many *Parameter* objects. For example, the objective functions you write for lmfit will take an instance of *Parameters* as its first argument. A table of parameter values, bounds and other attributes can be printed using `Parameters.pretty_print()`.

5.1 The Parameter class

```
class Parameter(name=None, value=None, vary=True, min=-inf, max=inf, expr=None,
                brute_step=None, user_data=None)
```

A *Parameter* is an object that can be varied in a fit, or one of the controlling variables in a model. It is a central component of lmfit, and all minimization and modeling methods use *Parameter* objects.

A *Parameter* has a *name* attribute, and a scalar floating point *value*. It also has a *vary* attribute that describes whether the value should be varied during the minimization. Finite bounds can be placed on the *Parameter*'s value by setting its *min* and/or *max* attributes. A *Parameter* can also have its value determined by a mathematical expression of other *Parameter* values held in the *expr* attribute. Additional attributes include *brute_step* used as the step size in a brute-force minimization, and *user_data* reserved exclusively for user's need.

After a minimization, a *Parameter* may also gain other attributes, including *stderr* holding the estimated standard error in the *Parameter*'s value, and *correl*, a dictionary of correlation values with other *Parameters* used in the minimization.

Parameters

- **name** (*str*, optional) – Name of the *Parameter*.
- **value** (*float*, optional) – Numerical *Parameter* value.

- **vary** (*bool*, *optional*) – Whether the Parameter is varied during a fit (default is True).
- **min** (*float*, *optional*) – Lower bound for value (default is *-numpy.inf*, no lower bound).
- **max** (*float*, *optional*) – Upper bound for value (default is *numpy.inf*, no upper bound).
- **expr** (*str*, *optional*) – Mathematical expression used to constrain the value during the fit.
- **brute_step** (*float*, *optional*) – Step size for grid points in the *brute* method.
- **user_data** (*optional*) – User-definable extra attribute used for a Parameter.

stderr

float – The estimated standard error for the best-fit value.

correl

dict – A dictionary of the correlation with the other fitted Parameters of the form:

```
{'decay': 0.404, 'phase': -0.020, 'frequency': 0.102}
```

See [Bounds Implementation](#) for details on the math used to implement the bounds with **min** and **max**.

The **expr** attribute can contain a mathematical expression that will be used to compute the value for the Parameter at each step in the fit. See [Using Mathematical Constraints](#) for more details and examples of this feature.

set (*value=None*, *vary=None*, *min=None*, *max=None*, *expr=None*, *brute_step=None*)

Set or update Parameter attributes.

Parameters

- **value** (*float*, *optional*) – Numerical Parameter value.
- **vary** (*bool*, *optional*) – Whether the Parameter is varied during a fit.
- **min** (*float*, *optional*) – Lower bound for value. To remove a lower bound you must use *-numpy.inf*.
- **max** (*float*, *optional*) – Upper bound for value. To remove an upper bound you must use *numpy.inf*.
- **expr** (*str*, *optional*) – Mathematical expression used to constrain the value during the fit. To remove a constraint you must supply an empty string.
- **brute_step** (*float*, *optional*) – Step size for grid points in the *brute* method. To remove the step size you must use 0.

Notes

Each argument to *set()* has a default value of *None*, which will leave the current value for the attribute unchanged. Thus, to lift a lower or upper bound, passing in *None* will not work. Instead, you must set these to *-numpy.inf* or *numpy.inf*, as with:

```
par.set(min=None)           # leaves lower bound unchanged
par.set(min=-numpy.inf)     # removes lower bound
```

Similarly, to clear an expression, pass a blank string, (not *None*!) as with:

```
par.set(expr=None)         # leaves expression unchanged
par.set(expr='')           # removes expression
```

Explicitly setting a value or setting `vary=True` will also clear the expression.

Finally, to clear the `brute_step` size, pass 0, not None:

```
par.set(brute_step=None) # leaves brute_step unchanged
par.set(brute_step=0)   # removes brute_step
```

5.2 The Parameters class

class Parameters (*asteval=None, *args, **kws*)

An ordered dictionary of all the Parameter objects required to specify a fit model. All minimization and Model fitting routines in lmfit will use exactly one Parameters object, typically given as the first argument to the objective function.

All keys of a Parameters() instance must be strings and valid Python symbol names, so that the name must match `[a-z_][a-z0-9_]*` and cannot be a Python reserved word.

All values of a Parameters() instance must be Parameter objects.

A Parameters() instance includes an `asteval` interpreter used for evaluation of constrained Parameters.

Parameters() support copying and pickling, and have methods to convert to and from serializations using json strings.

Parameters

- **asteval** (*asteval.Interpreter, optional*) – Instance of the `asteval` Interpreter to use for constraint expressions. If None, a new interpreter will be created.
- ***args** (*optional*) – Arguments.
- ****kws** (*optional*) – Keyword arguments.

add (*name, value=None, vary=True, min=-inf, max=inf, expr=None, brute_step=None*)

Add a Parameter.

Parameters

- **name** (*str*) – Name of parameter. Must match `[a-z_][a-z0-9_]*` and cannot be a Python reserved word.
- **value** (*float, optional*) – Numerical Parameter value, typically the *initial value*.
- **vary** (*bool, optional*) – Whether the Parameter is varied during a fit (default is True).
- **min** (*float, optional*) – Lower bound for value (default is `-numpy.inf`, no lower bound).
- **max** (*float, optional*) – Upper bound for value (default is `numpy.inf`, no upper bound).
- **expr** (*str, optional*) – Mathematical expression used to constrain the value during the fit.
- **brute_step** (*float, optional*) – Step size for grid points in the *brute* method.

Examples

```
>>> params = Parameters()
>>> params.add('xvar', value=0.50, min=0, max=1)
>>> params.add('yvar', expr='1.0 - xvar')
```

which is equivalent to:

```
>>> params = Parameters()
>>> params['xvar'] = Parameter(name='xvar', value=0.50, min=0, max=1)
>>> params['yvar'] = Parameter(name='yvar', expr='1.0 - xvar')
```

`add_many(*parlist)`

Add many parameters, using a sequence of tuples.

Parameters parlist (sequence of `tuple` or `Parameter`) – A sequence of tuples, or a sequence of `Parameter` instances. If it is a sequence of tuples, then each tuple must contain at least the name. The order in each tuple must be (*name, value, vary, min, max, expr, brute_step*).

Examples

```
>>> params = Parameters()
# add with tuples: (NAME VALUE VARY MIN MAX EXPR BRUTE_STEP)
>>> params.add_many(('amp', 10, True, None, None, None, None),
...                 ('cen', 4, True, 0.0, None, None, None),
...                 ('wid', 1, False, None, None, None, None),
...                 ('frac', 0.5))
# add a sequence of Parameters
>>> f = Parameter('par_f', 100)
>>> g = Parameter('par_g', 2.)
>>> params.add_many(f, g)
```

`pretty_print(online=False, colwidth=8, precision=4, fmt='g', columns=['value', 'min', 'max', 'stderr', 'vary', 'expr', 'brute_step'])`

Pretty-print of parameters data.

Parameters

- **online** (*bool, optional*) – If True prints a one-line parameters representation (default is False).
- **colwidth** (*int, optional*) – Column width for all columns specified in `columns`.
- **precision** (*int, optional*) – Number of digits to be printed after floating point.
- **fmt** (*{'g', 'e', 'f'}, optional*) – Single-character numeric formatter. Valid values are: 'f' floating point, 'g' floating point and exponential, or 'e' exponential.
- **columns** (list of `str`, optional) – List of `Parameter` attribute names to print.

`valuesdict()`

Return an ordered dictionary of parameter values.

Returns An ordered dictionary of name:value pairs for each `Parameter`.

Return type `OrderedDict`

dumps (***kws*)

Represent Parameters as a JSON string.

Parameters ***kws* (*optional*) – Keyword arguments that are passed to *json.dumps()*.

Returns JSON string representation of Parameters.

Return type *str*

See also:

dump(), *loads()*, *load()*, *json.dumps()*

dump (*fp*, ***kws*)

Write JSON representation of Parameters to a file-like object.

Parameters

- **fp** (*file-like object*) – An open and *.write()* -supporting file-like object.
- ***kws* (*optional*) – Keyword arguments that are passed to *dumps()*.

Returns Return value from *fp.write()*. None for Python 2.7 and the number of characters written in Python 3.

Return type None or *int*

See also:

dump(), *load()*, *json.dump()*

loads (*s*, ***kws*)

Load Parameters from a JSON string.

Parameters ***kws* (*optional*) – Keyword arguments that are passed to *json.loads()*.

Returns Updated Parameters from the JSON string.

Return type *Parameters*

Notes

Current Parameters will be cleared before loading the data from the JSON string.

See also:

dump(), *dumps()*, *load()*, *json.loads()*

load (*fp*, ***kws*)

Load JSON representation of Parameters from a file-like object.

Parameters

- **fp** (*file-like object*) – An open and *.read()* -supporting file-like object.
- ***kws* (*optional*) – Keyword arguments that are passed to *loads()*.

Returns Updated Parameters loaded from *fp*.

Return type *Parameters*

See also:

dump(), *loads()*, *json.load()*

5.3 Simple Example

A basic example making use of *Parameters* and the *minimize()* function (discussed in the next chapter) might look like this:

```
#!/usr/bin/env python

# <examples/doc_parameters_basic.py>
import numpy as np

from lmfit import Minimizer, Parameters, report_fit

# create data to be fitted
x = np.linspace(0, 15, 301)
data = (5. * np.sin(2*x - 0.1) * np.exp(-x*x*0.025) +
        np.random.normal(size=len(x), scale=0.2))

# define objective function: returns the array to be minimized
def fcn2min(params, x, data):
    """Model a decaying sine wave and subtract data."""
    amp = params['amp']
    shift = params['shift']
    omega = params['omega']
    decay = params['decay']
    model = amp * np.sin(x*omega + shift) * np.exp(-x*x*decay)
    return model - data

# create a set of Parameters
params = Parameters()
params.add('amp', value=10, min=0)
params.add('decay', value=0.1)
params.add('shift', value=0.0, min=-np.pi/2., max=np.pi/2)
params.add('omega', value=3.0)

# do fit, here with leastsq model
minner = Minimizer(fcn2min, params, fcn_args=(x, data))
result = minner.minimize()

# calculate final result
final = data + result.residual

# write error report
report_fit(result)

# try to plot results
try:
    import matplotlib.pyplot as plt
    plt.plot(x, data, 'k+')
    plt.plot(x, final, 'r')
    plt.show()
except ImportError:
    pass
# <end of examples/doc_parameters_basic.py>
```

Here, the objective function explicitly unpacks each Parameter value. This can be simplified using the *Parameters* *valuesdict()* method, which would make the objective function *fcn2min* above look like:


```
def fcn2min(params, x, data):  
    """Model a decaying sine wave and subtract data."""  
    v = params.valuesdict()  
  
    model = v['amp'] * np.sin(x*v['omega'] + v['shift']) * np.exp(-x*x*v['decay'])  
    return model - data
```

The results are identical, and the difference is a stylistic choice.

PERFORMING FITS AND ANALYZING OUTPUTS

As shown in the previous chapter, a simple fit can be performed with the `minimize()` function. For more sophisticated modeling, the `Minimizer` class can be used to gain a bit more control, especially when using complicated constraints or comparing results from related fits.

6.1 The `minimize()` function

The `minimize()` function is a wrapper around `Minimizer` for running an optimization problem. It takes an objective function (the function that calculates the array to be minimized), a `Parameters` object, and several optional arguments. See *Writing a Fitting Function* for details on writing the objective function.

`minimize` (*fcn*, *params*, *method*='leastsq', *args*=None, *kws*=None, *scale_covar*=True, *iter_cb*=None, *reduce_fcn*=None, ***fit_kws*)

Perform a fit of a set of parameters by minimizing an objective (or cost) function using one of the several available methods.

The `minimize` function takes an objective function to be minimized, a dictionary (`Parameters`) containing the model parameters, and several optional arguments.

Parameters

- **`fcn`** (*callable*) – Objective function to be minimized. When method is `leastsq` or `least_squares`, the objective function should return an array of residuals (difference between model and data) to be minimized in a least-squares sense. With the scalar methods the objective function can either return the residuals array or a single scalar value. The function must have the signature: `fcn(params, *args, **kws)`
- **`params`** (`Parameters`) – Contains the `Parameters` for the model.
- **`method`** (*str*, *optional*) – Name of the fitting method to use. Valid values are:
 - `'leastsq'`: Levenberg-Marquardt (default)
 - `'least_squares'`: Least-Squares minimization, using Trust Region Reflective method by default
 - `'differential_evolution'`: differential evolution
 - `'brute'`: brute force method
 - `'basinhopping'`: basinhopping
 - `'ampgo'`: Adaptive Memory Programming for Global Optimization
 - `'nelder'`: Nelder-Mead
 - `'lbfgsb'`: L-BFGS-B

- `'powell'`: Powell
- `'cg'`: Conjugate-Gradient
- `'newton'`: Newton-CG
- `'cobyla'`: Cobyla
- `'bfgs'`: BFGS
- `'tnc'`: Truncated Newton
- `'trust-ncg'`: Newton-CG trust-region
- `'trust-exact'`: nearly exact trust-region (SciPy ≥ 1.0)
- `'trust-krylov'`: Newton GLTR trust-region (SciPy ≥ 1.0)
- `'trust-constr'`: trust-region for constrained optimization (SciPy ≥ 1.1)
- `'dogleg'`: Dog-leg trust-region
- `'slsqp'`: Sequential Linear Squares Programming

In most cases, these methods wrap and use the method of the same name from `scipy.optimize`, or use `scipy.optimize.minimize` with the same `method` argument. Thus `'leastsq'` will use `scipy.optimize.leastsq`, while `'powell'` will use `scipy.optimize.minimizer(..., method='powell')`

For more details on the fitting methods please refer to the [SciPy docs](#).

- **args** (*tuple*, *optional*) – Positional arguments to pass to `fcn`.
- **kws** (*dict*, *optional*) – Keyword arguments to pass to `fcn`.
- **iter_cb** (*callable*, *optional*) – Function to be called at each fit iteration. This function should have the signature `iter_cb(params, iter, resid, *args, **kws)`, where `params` will have the current parameter values, `iter` the iteration number, `resid` the current residual array, and `*args` and `**kws` as passed to the objective function.
- **scale_covar** (*bool*, *optional*) – Whether to automatically scale the covariance matrix (default is True, `leastsq` only).
- **reduce_fcn** (*str* or *callable*, *optional*) – Function to convert a residual array to a scalar value for the scalar minimizers. See notes in *Minimizer*.
- ****fit_kws** (*dict*, *optional*) – Options to pass to the minimizer being used.

Returns Object containing the optimized parameter and several goodness-of-fit statistics.

Return type `MinimizerResult`

Changed in version 0.9.0: Return value changed to `MinimizerResult`.

Notes

The objective function should return the value to be minimized. For the Levenberg-Marquardt algorithm from `leastsq()`, this returned value must be an array, with a length greater than or equal to the number of fitting variables in the model. For the other methods, the return value can either be a scalar or an array. If an array is returned, the sum of squares of the array will be sent to the underlying fitting method, effectively doing a least-squares optimization of the return values.

A common use for *args* and *kws* would be to pass in other data needed to calculate the residual, including such things as the data array, dependent variable, uncertainties in the data, and other data structures for the model calculation.

On output, *params* will be unchanged. The best-fit values and, where appropriate, estimated uncertainties and correlations, will all be contained in the returned *MinimizerResult*. See *MinimizerResult – the optimization result* for further details.

This function is simply a wrapper around *Minimizer* and is equivalent to:

```
fitter = Minimizer(fcn, params, fcn_args=args, fcn_kws=kws,
                  iter_cb=iter_cb, scale_covar=scale_covar, **fit_kws)
fitter.minimize(method=method)
```

6.2 Writing a Fitting Function

An important component of a fit is writing a function to be minimized – the *objective function*. Since this function will be called by other routines, there are fairly stringent requirements for its call signature and return value. In principle, your function can be any Python callable, but it must look like this:

func(params, *args, **kws) :

Calculate objective residual to be minimized from parameters.

Parameters

- **params** (*Parameters*) – Parameters.
- **args** – Positional arguments. Must match *args* argument to *minimize()*.
- **kws** – Keyword arguments. Must match *kws* argument to *minimize()*.

Returns Residual array (generally data-model) to be minimized in the least-squares sense.

Return type `numpy.ndarray`. The length of this array cannot change between calls.

A common use for the positional and keyword arguments would be to pass in other data needed to calculate the residual, including things as the data array, dependent variable, uncertainties in the data, and other data structures for the model calculation.

The objective function should return the value to be minimized. For the Levenberg-Marquardt algorithm from *leastsq()*, this returned value **must** be an array, with a length greater than or equal to the number of fitting variables in the model. For the other methods, the return value can either be a scalar or an array. If an array is returned, the sum of squares of the array will be sent to the underlying fitting method, effectively doing a least-squares optimization of the return values.

Since the function will be passed in a dictionary of *Parameters*, it is advisable to unpack these to get numerical values at the top of the function. A simple way to do this is with *Parameters.valuesdict()*, as shown below:

```
def residual(pars, x, data=None, eps=None):
    from numpy import exp, sign, sin, pi
    # unpack parameters: extract .value attribute for each parameter
    parvals = pars.valuesdict()
    period = parvals['period']
    shift = parvals['shift']
    decay = parvals['decay']

    if abs(shift) > pi/2:
        shift = shift - sign(shift)*pi
```

(continues on next page)

(continued from previous page)

```
if abs(period) < 1.e-10:
    period = sign(period)*1.e-10

model = parvals['amp'] * sin(shift + x/period) * exp(-x*x*decay*decay)

if data is None:
    return model
if eps is None:
    return model - data
return (model-data) / eps
```

In this example, x is a positional (required) argument, while the *data* array is actually optional (so that the function returns the model calculation if the data is neglected). Also note that the model calculation will divide x by the value of the *period* Parameter. It might be wise to ensure this parameter cannot be 0. It would be possible to use bounds on the Parameter to do this:

```
params['period'] = Parameter(value=2, min=1.e-10)
```

but putting this directly in the function with:

```
if abs(period) < 1.e-10:
    period = sign(period)*1.e-10
```

is also a reasonable approach. Similarly, one could place bounds on the *decay* parameter to take values only between $-\pi/2$ and $\pi/2$.

6.3 Choosing Different Fitting Methods

By default, the [Levenberg-Marquardt](#) algorithm is used for fitting. While often criticized, including the fact it finds a *local* minima, this approach has some distinct advantages. These include being fast, and well-behaved for most curve-fitting needs, and making it easy to estimate uncertainties for and correlations between pairs of fit variables, as discussed in [MinimizerResult – the optimization result](#).

Alternative algorithms can also be used by providing the *method* keyword to the `minimize()` function or `Minimizer.minimize()` class as listed in the [Table of Supported Fitting Methods](#).

Table of Supported Fitting Methods:

| Fitting Method | method <code>arg</code> to <code>minimize()</code> or <code>Minimizer.minimize()</code> |
|---|---|
| Levenberg-Marquardt | <code>leastsq</code> or <code>least_squares</code> |
| Nelder-Mead | <code>nelder</code> |
| L-BFGS-B | <code>lbfgsb</code> |
| Powell | <code>powell</code> |
| Conjugate Gradient | <code>cg</code> |
| Newton-CG | <code>newton</code> |
| COBYLA | <code>cobyla</code> |
| BFGS | <code>bfgsb</code> |
| Truncated Newton | <code>tnc</code> |
| Newton CG trust-region | <code>trust-ncg</code> |
| Exact trust-region | <code>trust-exact</code> (SciPy ≥ 1.0) |
| Newton GLTR trust-region | <code>trust-krylov</code> (SciPy ≥ 1.0) |
| Constrained trust-region | <code>trust-constr</code> (SciPy ≥ 1.1) |
| Dogleg | <code>dogleg</code> |
| Sequential Linear Squares Programming | <code>slsqp</code> |
| Differential Evolution | <code>differential_evolution</code> |
| Brute force method | <code>brute</code> |
| Basinhopping | <code>basinhopping</code> |
| Adaptive Memory Programming for Global Optimization | <code>ampgo</code> |

Note: The objective function for the Levenberg-Marquardt method **must** return an array, with more elements than variables. All other methods can return either a scalar value or an array.

Warning: Much of this documentation assumes that the Levenberg-Marquardt method is used. Many of the fit statistics and estimates for uncertainties in parameters discussed in *MinimizerResult – the optimization result* are done only for this method.

6.4 MinimizerResult – the optimization result

New in version 0.9.0.

An optimization with `minimize()` or `Minimizer.minimize()` will return a `MinimizerResult` object. This is an otherwise plain container object (that is, with no methods of its own) that simply holds the results of the minimization. These results will include several pieces of informational data such as status and error messages, fit statistics, and the updated parameters themselves.

Importantly, the parameters passed in to `Minimizer.minimize()` will be not be changed. To find the best-fit values, uncertainties and so on for each parameter, one must use the `MinimizerResult.params` attribute. For example, to print the fitted values, bounds and other parameter attributes in a well-formatted text tables you can execute:

```
result.params.pretty_print()
```

with `results` being a `MinimizerResult` object. Note that the method `pretty_print()` accepts several arguments for customizing the output (e.g., column width, numeric format, etcetera).

class MinimizerResult (***kws*)

The results of a minimization.

Minimization results include data such as status and error messages, fit statistics, and the updated (i.e., best-fit) parameters themselves in the *params* attribute.

The list of (possible) *MinimizerResult* attributes is given below:

params

Parameters – The best-fit parameters resulting from the fit.

status

int – Termination status of the optimizer. Its value depends on the underlying solver. Refer to *message* for details.

var_names

list – Ordered list of variable parameter names used in optimization, and useful for understanding the values in *init_vals* and *covar*.

covar

numpy.ndarray – Covariance matrix from minimization (*leastsq* only), with rows and columns corresponding to *var_names*.

init_vals

list – List of initial values for variable parameters using *var_names*.

init_values

dict – Dictionary of initial values for variable parameters.

nfev

int – Number of function evaluations.

success

bool – True if the fit succeeded, otherwise False.

errorbars

bool – True if uncertainties were estimated, otherwise False.

message

str – Message about fit success.

ier

int – Integer error value from *scipy.optimize.leastsq* (*leastsq* only).

lmdif_message

str – Message from *scipy.optimize.leastsq* (*leastsq* only).

nvarys

int – Number of variables in fit: N_{varys} .

ndata

int – Number of data points: N .

nfree

int – Degrees of freedom in fit: $N - N_{\text{varys}}$.

residual

numpy.ndarray – Residual array Resid_i . Return value of the objective function when using the best-fit values of the parameters.

chisqr

float – Chi-square: $\chi^2 = \sum_i^N [\text{Resid}_i]^2$.

redchi

float – Reduced chi-square: $\chi^2_\nu = \chi^2 / (N - N_{\text{vars}})$.

aic

float – Akaike Information Criterion statistic: $N \ln(\chi^2/N) + 2N_{\text{vars}}$.

bic

float – Bayesian Information Criterion statistic: $N \ln(\chi^2/N) + \ln(N)N_{\text{vars}}$.

flatchain

pandas.DataFrame – A flatchain view of the sampling chain from the *emcee* method.

show_candidates()

Pretty_print() representation of candidates from the *brute* method.

6.4.1 Goodness-of-Fit Statistics

Table of Fit Results: These values, including the standard Goodness-of-Fit statistics, are all attributes of the *MinimizerResult* object returned by *minimize()* or *Minimizer.minimize()*.

| Attribute Name | Description / Formula |
|----------------|---|
| nfev | number of function evaluations |
| nvars | number of variables in fit N_{vars} |
| ndata | number of data points: N |
| nfree | degrees of freedom in fit: $N - N_{\text{vars}}$ |
| residual | residual array, returned by the objective function: $\{\text{Resid}_i\}$ |
| chisqr | chi-square: $\chi^2 = \sum_i [\text{Resid}_i]^2$ |
| redchi | reduced chi-square: $\chi^2_\nu = \chi^2 / (N - N_{\text{vars}})$ |
| aic | Akaike Information Criterion statistic (see below) |
| bic | Bayesian Information Criterion statistic (see below) |
| var_names | ordered list of variable parameter names used for <i>init_vals</i> and <i>covar</i> |
| covar | covariance matrix (with rows/columns using <i>var_names</i>) |
| init_vals | list of initial values for variable parameters |

Note that the calculation of chi-square and reduced chi-square assume that the returned residual function is scaled properly to the uncertainties in the data. For these statistics to be meaningful, the person writing the function to be minimized **must** scale them properly.

After a fit using the *leastsq()* method has completed successfully, standard errors for the fitted variables and correlations between pairs of fitted variables are automatically calculated from the covariance matrix. The standard error (estimated 1σ error-bar) goes into the *stderr* attribute of the *Parameter*. The correlations with all other variables will be put into the *correl* attribute of the *Parameter* – a dictionary with keys for all other *Parameters* and values of the corresponding correlation.

In some cases, it may not be possible to estimate the errors and correlations. For example, if a variable actually has no practical effect on the fit, it will likely cause the covariance matrix to be singular, making standard errors impossible to estimate. Placing bounds on varied *Parameters* makes it more likely that errors cannot be estimated, as being near the maximum or minimum value makes the covariance matrix singular. In these cases, the *errorbars* attribute of the fit result (*Minimizer* object) will be *False*.

6.4.2 Akaike and Bayesian Information Criteria

The `MinimizerResult` includes the traditional chi-square and reduced chi-square statistics:

$$\chi^2 = \sum_i^N r_i^2$$
$$\chi_\nu^2 = \chi^2 / (N - N_{\text{varys}})$$

where r is the residual array returned by the objective function (likely to be `(data-model)/uncertainty` for data modeling usages), N is the number of data points (`ndata`), and N_{varys} is number of variable parameters.

Also included are the [Akaike Information Criterion](#), and [Bayesian Information Criterion](#) statistics, held in the `aic` and `bic` attributes, respectively. These give slightly different measures of the relative quality for a fit, trying to balance quality of fit with the number of variable parameters used in the fit. These are calculated as:

$$\begin{aligned} \text{aic} &= N \ln(\chi^2/N) + 2N_{\text{varys}} \\ \text{bic} &= N \ln(\chi^2/N) + \ln(N)N_{\text{varys}} \end{aligned}$$

When comparing fits with different numbers of varying parameters, one typically selects the model with lowest reduced chi-square, Akaike information criterion, and/or Bayesian information criterion. Generally, the Bayesian information criterion is considered the most conservative of these statistics.

6.5 Using a Iteration Callback Function

An iteration callback function is a function to be called at each iteration, just after the objective function is called. The iteration callback allows user-supplied code to be run at each iteration, and can be used to abort a fit.

`iter_cb(params, iter, resid, *args, **kws):`

User-supplied function to be run at each iteration.

Parameters

- **`params`** (`Parameters`) – Parameters.
- **`iter`** (`int`) – Iteration number.
- **`resid`** (`numpy.ndarray`) – Residual array.
- **`args`** – Positional arguments. Must match `args` argument to `minimize()`.
- **`kws`** – Keyword arguments. Must match `kws` argument to `minimize()`.

Returns Residual array (generally `data-model`) to be minimized in the least-squares sense.

Return type None for normal behavior, any value like `True` to abort the fit.

Normally, the iteration callback would have no return value or return `None`. To abort a fit, have this function return a value that is `True` (including any non-zero integer). The fit will also abort if any exception is raised in the iteration callback. When a fit is aborted this way, the parameters will have the values from the last iteration. The fit statistics are not likely to be meaningful, and uncertainties will not be computed.

6.6 Using the `Minimizer` class

For full control of the fitting process, you will want to create a `Minimizer` object.

```
class Minimizer (userfcn, params, fcn_args=None, fcn_kws=None, iter_cb=None, scale_covar=True,
                 nan_policy='raise', reduce_fcn=None, **kws)
```

A general minimizer for curve fitting and optimization.

Parameters

- **userfcn** (*callable*) – Objective function that returns the residual (difference between model and data) to be minimized in a least-squares sense. This function must have the signature:

```
userfcn(params, *fcn_args, **fcn_kws)
```

- **params** (*Parameters*) – Contains the Parameters for the model.
- **fcn_args** (*tuple*, *optional*) – Positional arguments to pass to *userfcn*.
- **fcn_kws** (*dict*, *optional*) – Keyword arguments to pass to *userfcn*.
- **iter_cb** (*callable*, *optional*) – Function to be called at each fit iteration. This function should have the signature:

```
iter_cb(params, iter, resid, *fcn_args, **fcn_kws)
```

where *params* will have the current parameter values, *iter* the iteration number, *resid* the current residual array, and **fcn_args* and ***fcn_kws* are passed to the objective function.

- **scale_covar** (*bool*, *optional*) – Whether to automatically scale the covariance matrix (default is True, *leastsq* only).
- **nan_policy** (*str*, *optional*) – Specifies action if *userfcn* (or a Jacobian) returns NaN values. One of:
 - 'raise' : a *ValueError* is raised
 - 'propagate' : the values returned from *userfcn* are un-altered
 - 'omit' : non-finite values are filtered
- **reduce_fcn** (*str* or *callable*, *optional*) – Function to convert a residual array to a scalar value for the scalar minimizers. Optional values are (where *r* is the residual array):
 - None : sum of squares of residual [default]

$$= (r*r).sum()$$
 - 'negentropy' : neg entropy, using normal distribution

$$= \text{rho} * \log(\text{rho}).sum(), \text{ where } \text{rho} = \exp(-r*r/2) / (\text{sqrt}(2*\text{pi}))$$
 - 'neglogcauchy' : neg log likelihood, using Cauchy distribution

$$= -\log(1/(\text{pi}*(1+r*r))).sum()$$
 - callable : must take one argument (*r*) and return a float.
- ****kws** (*dict*, *optional*) – Options to pass to the minimizer being used.

Notes

The objective function should return the value to be minimized. For the Levenberg-Marquardt algorithm from *leastsq()* or *least_squares()*, this returned value must be an array, with a length greater than or equal to the number of fitting variables in the model. For the other methods, the return value can either be a scalar or

an array. If an array is returned, the sum of squares of the array will be sent to the underlying fitting method, effectively doing a least-squares optimization of the return values. If the objective function returns non-finite values then a *ValueError* will be raised because the underlying solvers cannot deal with them.

A common use for the *fcn_args* and *fcn_kws* would be to pass in other data needed to calculate the residual, including such things as the data array, dependent variable, uncertainties in the data, and other data structures for the model calculation.

The *Minimizer* object has a few public methods:

```
Minimizer.minimize (method='leastsq', params=None, **kws)  
    Perform the minimization.
```

Parameters

- **method** (*str*, optional) – Name of the fitting method to use. Valid values are:
 - *'leastsq'*: Levenberg-Marquardt (default)
 - *'least_squares'*: Least-Squares minimization, using Trust Region Reflective method by default
 - *'differential_evolution'*: differential evolution
 - *'brute'*: brute force method
 - *'basinhopping'*: basinhopping
 - *'ampgo'*: Adaptive Memory Programming for Global Optimization
 - *'nelder'*: Nelder-Mead
 - *'lbfgsb'*: L-BFGS-B
 - *'powell'*: Powell
 - *'cg'*: Conjugate-Gradient
 - *'newton'*: Newton-CG
 - *'cobyla'*: Cobyla
 - *'bfgs'*: BFGS
 - *'tnc'*: Truncated Newton
 - *'trust-ncg'*: Newton-CG trust-region
 - *'trust-exact'*: nearly exact trust-region (SciPy >= 1.0)
 - *'trust-krylov'*: Newton GLTR trust-region (SciPy >= 1.0)
 - *'trust-constr'*: trust-region for constrained optimization (SciPy >= 1.1)
 - *'dogleg'*: Dog-leg trust-region
 - *'slsqp'*: Sequential Linear Squares Programming

In most cases, these methods wrap and use the method with the same name from *scipy.optimize*, or use *scipy.optimize.minimize* with the same *method* argument. Thus *'leastsq'* will use *scipy.optimize.leastsq*, while *'powell'* will use *scipy.optimize.minimizer(..., method='powell')*

For more details on the fitting methods please refer to the [SciPy docs](#).

- **params** (*Parameters*, optional) – Parameters of the model to use as starting values.

- ****kws** (*optional*) – Additional arguments are passed to the underlying minimization method.

Returns Object containing the optimized parameter and several goodness-of-fit statistics.

Return type *MinimizerResult*

Changed in version 0.9.0: Return value changed to *MinimizerResult*.

`Minimizer.leastsq(params=None, **kws)`

Use Levenberg-Marquardt minimization to perform a fit.

It assumes that the input Parameters have been initialized, and a function to minimize has been properly set up. When possible, this calculates the estimated uncertainties and variable correlations from the covariance matrix.

This method calls `scipy.optimize.leastsq`. By default, numerical derivatives are used, and the following arguments are set:

| <code>leastsq()</code> arg | Default Value | Description |
|----------------------------|---------------|---|
| xtol | 1.e-7 | Relative error in the approximate solution |
| ftol | 1.e-7 | Relative error in the desired sum of squares |
| maxfev | 2000*(nvar+1) | Maximum number of function calls (nvar= # of variables) |
| Dfun | None | Function to call for Jacobian calculation |

Parameters

- **params** (*Parameters*, optional) – Parameters to use as starting point.
- ****kws** (*dict*, optional) – Minimizer options to pass to `scipy.optimize.leastsq`.

Returns Object containing the optimized parameter and several goodness-of-fit statistics.

Return type *MinimizerResult*

Changed in version 0.9.0: Return value changed to *MinimizerResult*.

`Minimizer.least_squares(params=None, **kws)`

Least-squares minimization using `scipy.optimize.least_squares`.

This method wraps `scipy.optimize.least_squares`, which has inbuilt support for bounds and robust loss functions. By default it uses the Trust Region Reflective algorithm with a linear loss function (i.e., the standard least-squares problem).

Parameters

- **params** (*Parameters*, optional) – Parameters to use as starting point.
- ****kws** (*dict*, optional) – Minimizer options to pass to `scipy.optimize.least_squares`.

Returns Object containing the optimized parameter and several goodness-of-fit statistics.

Return type *MinimizerResult*

Changed in version 0.9.0: Return value changed to *MinimizerResult*.

`Minimizer.scalar_minimize(method='Nelder-Mead', params=None, **kws)`

Scalar minimization using `scipy.optimize.minimize`.

Perform fit with any of the scalar minimization algorithms supported by `scipy.optimize.minimize`. Default argument values are:

| <code>scalar_minimize()</code> arg | Default Value | Description |
|------------------------------------|---------------|---------------------------------|
| method | Nelder-Mead | fitting method |
| tol | 1.e-7 | fitting and parameter tolerance |
| hess | None | Hessian of objective function |

Parameters

- **method** (*str*, optional) – Name of the fitting method to use. One of:
 - 'Nelder-Mead' (default)
 - 'L-BFGS-B'
 - 'Powell'
 - 'CG'
 - 'Newton-CG'
 - 'COBYLA'
 - 'BFGS'
 - 'TNC'
 - 'trust-ncg'
 - 'trust-exact' (SciPy >= 1.0)
 - 'trust-krylov' (SciPy >= 1.0)
 - 'trust-constr' (SciPy >= 1.1)
 - 'dogleg'
 - 'SLSQP'
 - 'differential_evolution'
- **params** (*Parameters*, optional) – Parameters to use as starting point.
- ****kws** (*dict*, optional) – Minimizer options pass to `scipy.optimize.minimize`.

Returns Object containing the optimized parameter and several goodness-of-fit statistics.

Return type *MinimizerResult*

Changed in version 0.9.0: Return value changed to *MinimizerResult*.

Notes

If the objective function returns a NumPy array instead of the expected scalar, the sum of squares of the array will be used.

Note that bounds and constraints can be set on *Parameters* for any of these methods, so are not supported separately for those designed to use bounds. However, if you use the `differential_evolution` method you must specify finite (min, max) for each varying Parameter.

`Minimizer.prepare_fit` (*params=None*)

Prepare parameters for fitting.

Prepares and initializes model and *Parameters* for subsequent fitting. This routine prepares the conversion of *Parameters* into fit variables, organizes parameter bounds, and parses, “compiles” and checks constrain

expressions. The method also creates and returns a new instance of a `MinimizerResult` object that contains the copy of the Parameters that will actually be varied in the fit.

Parameters `params` (`Parameters`, optional) – Contains the Parameters for the model; if None, then the Parameters used to initialize the Minimizer object are used.

Returns

Return type `MinimizerResult`

Notes

This method is called directly by the fitting methods, and it is generally not necessary to call this function explicitly.

Changed in version 0.9.0: Return value changed to `MinimizerResult`.

`Minimizer.brute` (`params=None`, `Ns=20`, `keep=50`)

Use the `brute` method to find the global minimum of a function.

The following parameters are passed to `scipy.optimize.brute` and cannot be changed:

| <code>brute()</code> arg | Value | Description |
|--------------------------|-------|---|
| <code>full_output</code> | 1 | Return the evaluation grid and the objective function's values on it. |
| <code>finish</code> | None | No “polishing” function is to be used after the grid search. |
| <code>disp</code> | False | Do not print convergence messages (when finish is not None). |

It assumes that the input Parameters have been initialized, and a function to minimize has been properly set up.

Parameters

- **params** (`Parameters`, optional) – Contains the Parameters for the model. If None, then the Parameters used to initialize the Minimizer object are used.
- **Ns** (`int`, optional) – Number of grid points along the axes, if not otherwise specified (see Notes).
- **keep** (`int`, optional) – Number of best candidates from the brute force method that are stored in the `candidates` attribute. If ‘all’, then all grid points from `scipy.optimize.brute` are stored as candidates.

Returns Object containing the parameters from the brute force method. The return values (`x0`, `fval`, `grid`, `Jout`) from `scipy.optimize.brute` are stored as `brute_<parname>` attributes. The `MinimizerResult` also contains the `candidates` attribute and `show_candidates()` method. The `candidates` attribute contains the parameters and `chisqr` from the brute force method as a namedtuple, (‘Candidate’, [‘params’, ‘score’]), sorted on the (lowest) `chisqr` value. To access the values for a particular candidate one can use `result.candidate[#].params` or `result.candidate[#].score`, where a lower # represents a better candidate. The `show_candidates(#)` uses the `pretty_print()` method to show a specific candidate-# or all candidates when no number is specified.

Return type `MinimizerResult`

New in version 0.9.6.

Notes

The `brute()` method evaluates the function at each point of a multidimensional grid of points. The grid points are generated from the parameter ranges using `Ns` and (optional) `brute_step`. The implementation in

`scipy.optimize.brute` requires finite bounds and the *range* is specified as a two-tuple (*min*, *max*) or slice-object (*min*, *max*, *brute_step*). A slice-object is used directly, whereas a two-tuple is converted to a slice object that interpolates *Ns* points from *min* to *max*, inclusive.

In addition, the `brute()` method in `lmfit`, handles three other scenarios given below with their respective slice-object:

- **lower bound (min) and brute_step are specified:** `range = (min, min + Ns * brute_step, brute_step)`.
- **upper bound (max) and brute_step are specified:** `range = (max - Ns * brute_step, max, brute_step)`.
- **numerical value (value) and brute_step are specified:** `range = (value - (Ns//2) * brute_step, value + (Ns//2) * brute_step, brute_step)`.

For more information, check the examples in `examples/lmfit_brute_example.ipynb`.

`Minimizer.basinhopping` (*params=None, **kws*)

Use the *basinhopping* algorithm to find the global minimum of a function.

This method calls `scipy.optimize.basinhopping` using the default arguments. The default minimizer is *BFGS*, but since `lmfit` supports parameter bounds for all minimizers, the user can choose any of the solvers present in `scipy.optimize.minimize`.

Parameters *params* (*Parameters* object, optional) – Contains the Parameters for the model. If None, then the Parameters used to initialize the Minimizer object are used.

Returns Object containing the optimization results from the basinhopping algorithm.

Return type *MinimizerResult*

New in version 0.9.10.

`Minimizer.ampgo` (*params=None, **kws*)

Finds the global minimum of a multivariate function using the AMPGO (Adaptive Memory Programming for Global Optimization) algorithm.

Parameters

- **params** (*Parameters*, optional) – Contains the Parameters for the model. If None, then the Parameters used to initialize the Minimizer object are used.
- ****kws** (*dict*, optional) – Minimizer options to pass to the ampgo algorithm, the options are listed below:

```
local: str (default is 'L-BFGS-B')
    Name of the local minimization method. Valid options are:
    - 'L-BFGS-B'
    - 'Nelder-Mead'
    - 'Powell'
    - 'TNC'
    - 'SLSQP'
local_opts: dict (default is None)
    Options to pass to the local minimizer.
maxfunvals: int (default is None)
    Maximum number of function evaluations. If None, the
    ↪ optimization will stop
    after `totaliter` number of iterations.
totaliter: int (default is 20)
    Maximum number of global iterations.
maxiter: int (default is 5)
    Maximum number of `Tabu Tunneling` iterations during each
    ↪ global iteration.
```

(continues on next page)

(continued from previous page)

```

glbtol: float (default is 1e-5)
    Tolerance whether or not to accept a solution after a
    ↪ tunneling phase.
eps1: float (default is 0.02)
    Constant used to define an aspiration value for the objective
    ↪ function during
    the Tunneling phase.
eps2: float (default is 0.1)
    Perturbation factor used to move away from the latest local
    ↪ minimum at the
    start of a Tunneling phase.
tabulistsize: int (default is 5)
    Size of the (circular) tabu search list.
tabustrategy: str (default is 'farthest')
    Strategy to use when the size of the tabu list exceeds
    ↪ `tabulistsize`. It
    can be 'oldest' to drop the oldest point from the tabu list or
    ↪ 'farthest'
    to drop the element farthest from the last local minimum found.
disp: bool (default is False)
    Set to True to print convergence messages.

```

Returns Object containing the parameters from the `ampgo` method, with fit parameters, statistics and such. The return values (`x0`, `fval`, `eval`, `msg`, `tunnel`) are stored as `ampgo_<parname>` attributes.

Return type `MinimizerResult`

New in version 0.9.10.

Notes

The Python implementation was written by Andrea Gavana in 2014 (http://infinity77.net/global_optimization/index.html).

The details of the AMPGO algorithm are described in the paper “Adaptive Memory Programming for Constrained Global Optimization” located here:

[http://leeds-faculty.colorado.edu/glover/fred%20pubs/416%20-%20AMP%20\(TS\)%20for%20Constrained%20Global%20Opt%20w%20Lasdon%20et%20al%20.pdf](http://leeds-faculty.colorado.edu/glover/fred%20pubs/416%20-%20AMP%20(TS)%20for%20Constrained%20Global%20Opt%20w%20Lasdon%20et%20al%20.pdf)

`Minimizer.emcee` (`params=None`, `steps=1000`, `nwalkers=100`, `burn=0`, `thin=1`, `ntemps=1`, `pos=None`, `reuse_sampler=False`, `workers=1`, `float_behavior='posterior'`, `is_weighted=True`, `seed=None`)

Bayesian sampling of the posterior distribution using `emcee`.

Bayesian sampling of the posterior distribution for the parameters using the `emcee` Markov Chain Monte Carlo package. The method assumes that the prior is Uniform. You need to have `emcee` installed to use this method.

Parameters

- **params** (`Parameters`, optional) – Parameters to use as starting point. If this is not specified then the Parameters used to initialize the `Minimizer` object are used.
- **steps** (`int`, optional) – How many samples you would like to draw from the posterior distribution for each of the walkers?
- **nwalkers** (`int`, optional) – Should be set so `nwalkers` \gg `nvars`, where `nvars` are the number of parameters being varied during the fit. “Walkers are the members of the ensemble. They are almost like separate Metropolis-Hastings chains but, of course, the

proposal distribution for a given walker depends on the positions of all the other walkers in the ensemble.” - from the *emcee* webpage.

- **burn** (*int*, *optional*) – Discard this many samples from the start of the sampling regime.
- **thin** (*int*, *optional*) – Only accept 1 in every *thin* samples.
- **ntemps** (*int*, *optional*) – If *ntemps* > 1 perform a Parallel Tempering.
- **pos** (*numpy.ndarray*, *optional*) – Specify the initial positions for the sampler. If *ntemps* == 1 then *pos.shape* should be (*nwalkers*, *nvars*). Otherwise, (*ntemps*, *nwalkers*, *nvars*). You can also initialise using a previous chain that had the same *ntemps*, *nwalkers* and *nvars*. Note that *nvars* may be one larger than you expect it to be if your *userfcn* returns an array and *is_weighted* is *False*.
- **reuse_sampler** (*bool*, *optional*) – If you have already run *emcee* on a given *Minimizer* object then it possesses an internal *sampler* attribute. You can continue to draw from the same sampler (retaining the chain history) if you set this option to *True*. Otherwise a new sampler is created. The *nwalkers*, *ntemps*, *pos*, and *params* keywords are ignored with this option. **Important:** the Parameters used to create the sampler must not change in-between calls to *emcee*. Alteration of Parameters would include changed *min*, *max*, *vary* and *expr* attributes. This may happen, for example, if you use an altered Parameters object and call the *minimize* method in-between calls to *emcee*.
- **workers** (*Pool-like or int*, *optional*) – For parallelization of sampling. It can be any Pool-like object with a *map* method that follows the same calling sequence as the built-in *map* function. If *int* is given as the argument, then a multiprocessing-based pool is spawned internally with the corresponding number of parallel processes. ‘*mpi4py*’-based parallelization and ‘*joblib*’-based parallelization pools can also be used here. **Note:** because of multiprocessing overhead it may only be worth parallelising if the objective function is expensive to calculate, or if there are a large number of objective evaluations per step (*ntemps* * *nwalkers* * *nvars*).
- **float_behavior** (*str*, *optional*) – Specifies meaning of the objective function output if it returns a float. One of:
 - ‘posterior’ - objective function returns a log-posterior probability
 - ‘chi2’ - objective function returns χ^2See Notes for further details.
- **is_weighted** (*bool*, *optional*) – Has your objective function been weighted by measurement uncertainties? If *is_weighted* is *True* then your objective function is assumed to return residuals that have been divided by the true measurement uncertainty (*data - model*) / *sigma*. If *is_weighted* is *False* then the objective function is assumed to return unweighted residuals, *data - model*. In this case *emcee* will employ a positive measurement uncertainty during the sampling. This measurement uncertainty will be present in the output *params* and output chain with the name *__Insigma*. A side effect of this is that you cannot use this parameter name yourself. **Important** this parameter only has any effect if your objective function returns an array. If your objective function returns a float, then this parameter is ignored. See Notes for more details.
- **seed** (*int or numpy.random.RandomState*, *optional*) – If *seed* is an *int*, a new *numpy.random.RandomState* instance is used, seeded with *seed*. If *seed* is already a *numpy.random.RandomState* instance, then that *numpy.random.RandomState* instance is used. Specify *seed* for repeatable minimizations.

Returns `MinimizerResult` object containing updated params, statistics, etc. The updated params represent the median (50th percentile) of all the samples, whilst the parameter uncertainties are half of the difference between the 15.87 and 84.13 percentiles. The `MinimizerResult` also contains the `chain`, `flatchain` and `lnprob` attributes. The `chain` and `flatchain` attributes contain the samples and have the shape $(nwalkers, (steps - burn) // thin, nvarys)$ or $(ntemps, nwalkers, (steps - burn) // thin, nvarys)$, depending on whether Parallel tempering was used or not. `nvarys` is the number of parameters that are allowed to vary. The `flatchain` attribute is a `pandas.DataFrame` of the flattened chain, `chain.reshape(-1, nvarys)`. To access flattened chain values for a particular parameter use `result.flatchain[parname]`. The `lnprob` attribute contains the log probability for each sample in `chain`. The sample with the highest probability corresponds to the maximum likelihood estimate.

Return type `MinimizerResult`

Notes

This method samples the posterior distribution of the parameters using Markov Chain Monte Carlo. To do so it needs to calculate the log-posterior probability of the model parameters, F , given the data, D , $\ln p(F_{true}|D)$. This ‘posterior probability’ is calculated as:

$$\ln p(F_{true}|D) \propto \ln p(D|F_{true}) + \ln p(F_{true})$$

where $\ln p(D|F_{true})$ is the ‘log-likelihood’ and $\ln p(F_{true})$ is the ‘log-prior’. The default log-prior encodes prior information already known about the model. This method assumes that the log-prior probability is `-numpy.inf` (impossible) if the one of the parameters is outside its limits. The log-prior probability term is zero if all the parameters are inside their bounds (known as a uniform prior). The log-likelihood function is given by¹:

$$\ln p(D|F_{true}) = -\frac{1}{2} \sum_n \left[\frac{(g_n(F_{true}) - D_n)^2}{s_n^2} + \ln(2\pi s_n^2) \right]$$

The first summand in the square brackets represents the residual for a given datapoint (g being the generative model, D_n the data and s_n the standard deviation, or measurement uncertainty, of the datapoint). This term represents χ^2 when summed over all data points. Ideally the objective function used to create `lmfit.Minimizer` should return the log-posterior probability, $\ln p(F_{true}|D)$. However, since the in-built log-prior term is zero, the objective function can also just return the log-likelihood, unless you wish to create a non-uniform prior.

If a float value is returned by the objective function then this value is assumed by default to be the log-posterior probability, i.e. `float_behavior` is ‘posterior’. If your objective function returns χ^2 , then you should use a value of ‘chi2’ for `float_behavior`. `emcee` will then multiply your χ^2 value by -0.5 to obtain the posterior probability.

However, the default behaviour of many objective functions is to return a vector of (possibly weighted) residuals. Therefore, if your objective function returns a vector, `res`, then the vector is assumed to contain the residuals. If `is_weighted` is `True` then your residuals are assumed to be correctly weighted by the standard deviation (measurement uncertainty) of the data points (`res = (data - model) / sigma`) and the log-likelihood (and log-posterior probability) is calculated as: `-0.5 * numpy.sum(res**2)`. This ignores the second summand in the square brackets. Consequently, in order to calculate a fully correct log-posterior probability value your objective function should return a single value. If `is_weighted` is `False` then the data uncertainty, `s_n`, will be treated as a nuisance parameter and will be marginalized out. This is achieved by employing a strictly positive uncertainty (homoscedasticity) for each data point, `s_n = exp(__lnsigma)`. `__lnsigma` will be present in `MinimizerResult.params`, as well as `Minimizer.chain`, `nvarys` will also be increased by one.


¹ <http://dan.iel.fm/emcee/current/user/line/>

References

6.7 `Minimizer.emcee()` - calculating the posterior probability distribution of parameters

`Minimizer.emcee()` can be used to obtain the posterior probability distribution of parameters, given a set of experimental data. An example problem is a double exponential decay. A small amount of Gaussian noise is also added in:

```
>>> import numpy as np
>>> import lmfit
>>> import matplotlib.pyplot as plt
>>> x = np.linspace(1, 10, 250)
>>> np.random.seed(0)
>>> y = 3.0 * np.exp(-x / 2) - 5.0 * np.exp(-(x - 0.1) / 10.) + 0.1 * np.random.
↳ randn(len(x))
>>> plt.plot(x, y)
>>> plt.show()
```



Create a Parameter set for the initial guesses:

```
>>> p = lmfit.Parameters()
>>> p.add_many(('a1', 4.), ('a2', 4.), ('t1', 3.), ('t2', 3., True))

>>> def residual(p):
...     v = p.valuesdict()
...     return v['a1'] * np.exp(-x / v['t1']) + v['a2'] * np.exp(-(x - 0.1) / v['t2
↳ ']) - y
```

Solving with `minimize()` gives the Maximum Likelihood solution:


```
>>> mi = lmfit.minimize(residual, p, method='Nelder', nan_policy='omit')
>>> lmfit.printfuncs.report_fit(mi.params, min_correl=0.5)
[[Variables]]
a1:  2.98623689 (init = 4)
a2: -4.33525597 (init = 4)
t1:  1.30993186 (init = 3)
t2: 11.8240752 (init = 3)

>>> plt.plot(x, y)
```

(continues on next page)

(continued from previous page)

```
>>> plt.plot(x, residual(mi.params) + y, 'r')
>>> plt.show()
```



However, this doesn't give a probability distribution for the parameters. Furthermore, we wish to deal with the data uncertainty. This is called marginalisation of a nuisance parameter. `emcee` requires a function that returns the log-posterior probability. The log-posterior probability is a sum of the log-prior probability and log-likelihood functions. The log-prior probability is assumed to be zero if all the parameters are within their bounds and `-np.inf` if any of the parameters are outside their bounds.

```
>>> # add a noise parameter
>>> mi.params.add('noise', value=1, min=0.001, max=2)
```

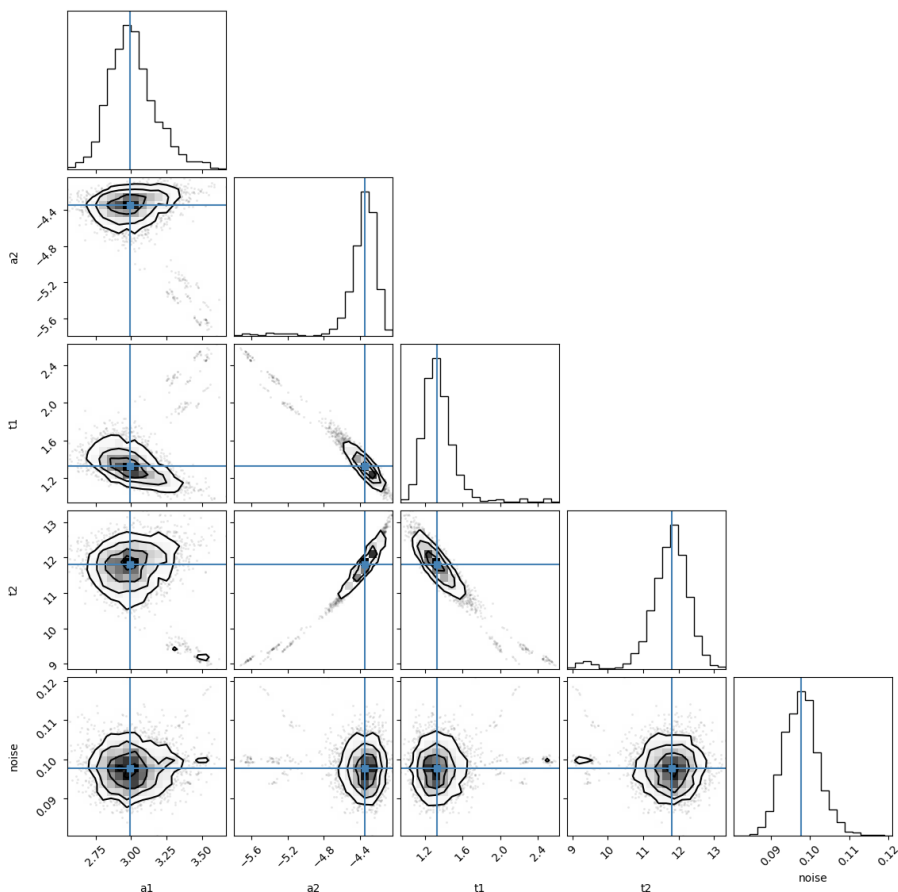
```
>>> # This is the log-likelihood probability for the sampling. We're going to
↳ estimate the
>>> # size of the uncertainties on the data as well.
>>> def lnprob(p):
...     noise = p['noise']
...     return -0.5 * np.sum((residual(p) / noise)**2 + np.log(2 * np.pi * noise**2))
```

Now we have to set up the minimizer and do the sampling:

```
>>> mini = lmfit.Minimizer(lnprob, mi.params)
>>> res = mini.emcee(burn=300, steps=1000, thin=20, params=mi.params)
```

Lets have a look at those posterior distributions for the parameters. This requires installation of the `corner` package:

```
>>> import corner
>>> corner.corner(res.flatchain, labels=res.var_names, truths=list(res.params.
↳ valuesdict().values()))
```



The values reported in the `MinimizerResult` are the medians of the probability distributions and a 1σ quantile, estimated as half the difference between the 15.8 and 84.2 percentiles. The median value is not necessarily the same as the Maximum Likelihood Estimate. We'll get that as well. You can see that we recovered the right uncertainty level on the data:

```
>>> print("median of posterior probability distribution")
>>> print('-----')
>>> lmfit.report_fit(res.params)
median of posterior probability distribution
-----
[[Variables]]
  a1:      2.99342394 +/- 0.15851315 (5.30%) (init = 2.986237)
  a2:     -4.34384999 +/- 0.12454831 (2.87%) (init = -4.335256)
  t1:      1.32338403 +/- 0.14120290 (10.67%) (init = 1.309932)
  t2:     11.7962437 +/- 0.48632272 (4.12%) (init = 11.82408)
  noise:  0.09761521 +/- 0.00431795 (4.42%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
  C(a2, t1)    = -0.965
  C(a2, t2)    =  0.959
  C(t1, t2)    = -0.927
  C(a1, a2)    = -0.241
  C(a1, t2)    = -0.168
  C(a2, noise) = -0.116
  C(t1, noise) =  0.107

>>> # find the maximum likelihood solution
>>> highest_prob = np.argmax(res.lnprob)
```

(continues on next page)

(continued from previous page)

```

>>> hp_loc = np.unravel_index(highest_prob, res.lnprob.shape)
>>> mle_soln = res.chain[hp_loc]
>>> for i, par in enumerate(p):
...     p[par].value = mle_soln[i]

>>> print("\nMaximum likelihood Estimation")
>>> print('-----')
>>> print(p)
Maximum likelihood Estimation
-----
Parameters([('a1', <Parameter 'a1', 2.9684811738216754, bounds=[-inf:inf]>),
('a2', <Parameter 'a2', -4.355238699173162, bounds=[-inf:inf]>),
('t1', <Parameter 't1', 1.3337647386777762, bounds=[-inf:inf]>),
('t2', <Parameter 't2', 11.758394302818514, bounds=[-inf:inf]>)])
>>> # Finally lets work out a 1 and 2-sigma error estimate for 't1'
>>> quantiles = np.percentile(res.flatchain['t1'], [2.28, 15.9, 50, 84.2, 97.7])
>>> print("1 sigma spread", 0.5 * (quantiles[3] - quantiles[1]))
>>> print("2 sigma spread", 0.5 * (quantiles[4] - quantiles[0]))
1 sigma spread 0.1414604069179637
2 sigma spread 0.453234685099423

```

6.8 Getting and Printing Fit Reports

fit_report (*inpars*, *modelpars*=None, *show_correl*=True, *min_correl*=0.1, *sort_pars*=False)

Generate a report of the fitting results.

The report contains the best-fit values for the parameters and their uncertainties and correlations.

Parameters

- **inpars** (*Parameters*) – Input Parameters from fit or *MinimizerResult* returned from a fit.
- **modelpars** (*Parameters*, *optional*) – Known Model Parameters.
- **show_correl** (*bool*, *optional*) – Whether to show list of sorted correlations (default is True).
- **min_correl** (*float*, *optional*) – Smallest correlation in absolute value to show (default is 0.1).
- **sort_pars** (*bool or callable*, *optional*) – Whether to show parameter names sorted in alphanumerical order. If False (default), then the parameters will be listed in the order they were added to the *Parameters* dictionary. If callable, then this (one argument) function is used to extract a comparison key from each list element.

Returns Multi-line text of fit report.

Return type string

An example using this to write out a fit report would be:

```

#!/usr/bin/env python

# <examples/doc_fitting_withreport.py>
from numpy import exp, linspace, pi, random, sign, sin

```

(continues on next page)

(continued from previous page)

```

from lmfit import Parameters, fit_report, minimize

p_true = Parameters()
p_true.add('amp', value=14.0)
p_true.add('period', value=5.46)
p_true.add('shift', value=0.123)
p_true.add('decay', value=0.032)

def residual(pars, x, data=None):
    vals = pars.valuesdict()
    amp = vals['amp']
    per = vals['period']
    shift = vals['shift']
    decay = vals['decay']

    if abs(shift) > pi/2:
        shift = shift - sign(shift)*pi
    model = amp * sin(shift + x/per) * exp(-x*x*decay*decay)
    if data is None:
        return model
    return model - data

n = 1001
xmin = 0.
xmax = 250.0

random.seed(0)

noise = random.normal(scale=0.7215, size=n)
x = linspace(xmin, xmax, n)
data = residual(p_true, x) + noise

fit_params = Parameters()
fit_params.add('amp', value=13.0)
fit_params.add('period', value=2)
fit_params.add('shift', value=0.0)
fit_params.add('decay', value=0.02)

out = minimize(residual, fit_params, args=(x,), kws={'data': data})

print(fit_report(out))
# <end of examples/doc_fitting_withreport.py>

```

which would write out:

```

[[Fit Statistics]]
  # fitting method   = leastsq
  # function evals   = 83
  # data points      = 1001
  # variables        = 4
  chi-square         = 498.811759
  reduced chi-square = 0.50031270
  Akaike info crit   = -689.222517
  Bayesian info crit = -669.587497
[[Variables]]

```

(continues on next page)

(continued from previous page)

```
amp:      13.9121945 +/- 0.14120288 (1.01%) (init = 13)
period:   5.48507045 +/- 0.02666492 (0.49%) (init = 2)
shift:    0.16203677 +/- 0.01405661 (8.67%) (init = 0)
decay:    0.03264538 +/- 3.8014e-04 (1.16%) (init = 0.02)
[[Correlations]] (unreported correlations are < 0.100)
C(period, shift) = 0.797
C(amp, decay)    = 0.582
C(amp, shift)    = -0.297
C(amp, period)   = -0.243
C(shift, decay)  = -0.182
C(period, decay) = -0.150
```


MODELING DATA AND CURVE FITTING

A common use of least-squares minimization is *curve fitting*, where one has a parametrized model function meant to explain some phenomena and wants to adjust the numerical values for the model so that it most closely matches some data. With `scipy`, such problems are typically solved with `scipy.optimize.curve_fit`, which is a wrapper around `scipy.optimize.leastsq`. Since `lmfit`'s `minimize()` is also a high-level wrapper around `scipy.optimize.leastsq` it can be used for curve-fitting problems. While it offers many benefits over `scipy.optimize.leastsq`, using `minimize()` for many curve-fitting problems still requires more effort than using `scipy.optimize.curve_fit`.

The `Model` class in `lmfit` provides a simple and flexible approach to curve-fitting problems. Like `scipy.optimize.curve_fit`, a `Model` uses a *model function* – a function that is meant to calculate a model for some phenomenon – and then uses that to best match an array of supplied data. Beyond that similarity, its interface is rather different from `scipy.optimize.curve_fit`, for example in that it uses `Parameters`, but also offers several other important advantages.

In addition to allowing you to turn any model function into a curve-fitting method, `lmfit` also provides canonical definitions for many known line shapes such as Gaussian or Lorentzian peaks and Exponential decays that are widely used in many scientific domains. These are available in the `models` module that will be discussed in more detail in the next chapter (*Built-in Fitting Models in the models module*). We mention it here as you may want to consult that list before writing your own model. For now, we focus on turning Python functions into high-level fitting models with the `Model` class, and using these to fit data.

7.1 Motivation and simple example: Fit data to Gaussian profile

Let's start with a simple and common example of fitting data to a Gaussian peak. As we will see, there is a built-in `GaussianModel` class that can help do this, but here we'll build our own. We start with a simple definition of the model function:

```
>>> from numpy import exp, linspace, random
>>>
>>> def gaussian(x, amp, cen, wid):
...     return amp * exp(-(x-cen)**2 / wid)
```

We want to use this function to fit to data $y(x)$ represented by the arrays `y` and `x`. With `scipy.optimize.curve_fit`, this would be:

```
>>> from scipy.optimize import curve_fit
>>>
>>> x = linspace(-10, 10, 101)
>>> y = gaussian(x, 2.33, 0.21, 1.51) + random.normal(0, 0.2, len(x))
>>>
>>> init_vals = [1, 0, 1] # for [amp, cen, wid]
>>> best_vals, covar = curve_fit(gaussian, x, y, p0=init_vals)
```

(continues on next page)

(continued from previous page)

```
>>> print(best_vals)
[ 2.39499424  0.19942179  1.50909389]
```

That is, we create data, make an initial guess of the model values, and run `scipy.optimize.curve_fit` with the model function, data arrays, and initial guesses. The results returned are the optimal values for the parameters and the covariance matrix. It's simple and useful, but it misses the benefits of `lmfit`.

With `lmfit`, we create a `Model` that wraps the *gaussian* model function, which automatically generates the appropriate residual function, and determines the corresponding parameter names from the function signature itself:

```
>>> from lmfit import Model
>>> gmodel = Model(gaussian)
>>> gmodel.param_names
['amp', 'cen', 'wid']
>>> gmodel.independent_vars
['x']
```

As you can see, the `Model gmodel` determined the names of the parameters and the independent variables. By default, the first argument of the function is taken as the independent variable, held in `independent_vars`, and the rest of the functions positional arguments (and, in certain cases, keyword arguments – see below) are used for Parameter names. Thus, for the *gaussian* function above, the independent variable is `x`, and the parameters are named *amp*, *cen*, and *wid*, and – all taken directly from the signature of the model function. As we will see below, you can modify the default assignment of independent variable / arguments and specify yourself what the independent variable is and which function arguments should be identified as parameter names.

The Parameters are *not* created when the model is created. The model knows what the parameters should be named, but nothing about the scale and range of your data. You will normally have to make these parameters and assign initial values and other attributes. To help you do this, each model has a `make_params()` method that will generate parameters with the expected names:

```
>>> params = gmodel.make_params()
```

This creates the *Parameters* but does not automatically give them initial values since it has no idea what the scale should be. You can set initial values for parameters with keyword arguments to `make_params()`:

```
>>> params = gmodel.make_params(cen=5, amp=200, wid=1)
```

or assign them (and other parameter properties) after the *Parameters* class has been created.

A `Model` has several methods associated with it. For example, one can use the `eval()` method to evaluate the model or the `fit()` method to fit data to this model with a `Parameter` object. Both of these methods can take explicit keyword arguments for the parameter values. For example, one could use `eval()` to calculate the predicted function:

```
>>> x = linspace(0, 10, 201)
>>> y = gmodel.eval(params, x=x)
```

or with:

```
>>> y = gmodel.eval(x=x, cen=6.5, amp=100, wid=2.0)
```

Admittedly, this is a slightly long-winded way to calculate a Gaussian function, given that you could have called your *gaussian* function directly. But now that the model is set up, we can use its `fit()` method to fit this model to data, as with:

```
>>> result = gmodel.fit(y, params, x=x)
```

or with:

```
>>> result = gmodel.fit(y, x=x, cen=6.5, amp=100, wid=2.0)
```

Putting everything together, included in the `examples` folder with the source code, is:

```
#!/usr/bin/env python

# <examples/doc_model_gaussian.py>
import matplotlib.pyplot as plt
from numpy import exp, loadtxt, pi, sqrt

from lmfit import Model

data = loadtxt('model1d_gauss.dat')
x = data[:, 0]
y = data[:, 1]

def gaussian(x, amp, cen, wid):
    """1-d gaussian: gaussian(x, amp, cen, wid)"""
    return (amp / (sqrt(2*pi) * wid)) * exp(-(x-cen)**2 / (2*wid**2))

gmodel = Model(gaussian)
result = gmodel.fit(y, x=x, amp=5, cen=5, wid=1)


print(result.fit_report())

plt.plot(x, y, 'bo')
plt.plot(x, result.init_fit, 'k--')
plt.plot(x, result.best_fit, 'r-')
plt.show()
# <end examples/doc_model_gaussian.py>
```

which is pretty compact and to the point. The returned *result* will be a *ModelResult* object. As we will see below, this has many components, including a `fit_report()` method, which will show:

```
[[Model]]
  Model(gaussian)
[[Fit Statistics]]
  # fitting method   = leastsq
  # function evals   = 35
  # data points      = 101
  # variables        = 3
  chi-square         = 3.40883599
  reduced chi-square = 0.03478404
  Akaike info crit   = -336.263713
  Bayesian info crit = -328.418352
[[Variables]]
  amp:  8.88021830 +/- 0.11359492 (1.28%) (init = 5)
  cen:  5.65866102 +/- 0.01030495 (0.18%) (init = 5)
  wid:  0.69765468 +/- 0.01030495 (1.48%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
  C(amp, wid) = 0.577
```

As the script shows, the result will also have *init_fit* for the fit with the initial parameter values and a *best_fit* for the fit with the best fit parameter values. These can be used to generate the following plot:



red line, and the initial fit as a dashed black line.

Note that the model fitting was really performed with:

```
gmodel = Model(gaussian)
result = gmodel.fit(y, params, x=x, amp=5, cen=5, wid=1)
```

These lines clearly express that we want to turn the *gaussian* function into a fitting model, and then fit the $y(x)$ data to this model, starting with values of 5 for *amp*, 5 for *cen* and 1 for *wid*. In addition, all the other features of *lmfit* are included: *Parameters* can have bounds and constraints and the result is a rich object that can be reused to explore the model fit in detail.

7.2 The Model class

The *Model* class provides a general way to wrap a pre-defined function as a fitting model.

class Model (*func*, *independent_vars*=None, *param_names*=None, *nan_policy*='raise', *missing*=None, *prefix*="", *name*=None, ***kws*)

Create a model from a user-supplied model function.

The model function will normally take an independent variable (generally, the first argument) and a series of arguments that are meant to be parameters for the model. It will return an array of data to model some data as for a curve-fitting problem.

Parameters

- **func** (*callable*) – Function to be wrapped.
- **independent_vars** (*list of str*, *optional*) – Arguments to *func* that are independent variables (default is None).
- **param_names** (*list of str*, *optional*) – Names of arguments to *func* that are to be made into parameters (default is None).
- **nan_policy** (*str*, *optional*) – How to handle NaN and missing values in data. Must be one of 'raise' (default), 'propagate', or 'omit'. See Note below.
- **missing** (*str*, *optional*) – Synonym for 'nan_policy' for backward compatibility.
- **prefix** (*str*, *optional*) – Prefix used for the model.
- **name** (*str*, *optional*) – Name for the model. When None (default) the name is the same as the model function (*func*).

- ****kws** (*dict*, *optional*) – Additional keyword arguments to pass to model function.

Notes

1. Parameter names are inferred from the function arguments, and a residual function is automatically constructed.
2. The model function must return an array that will be the same size as the data being modeled.
3. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - ‘raise’ : Raise a `ValueError` (default)
 - ‘propagate’ : do nothing
 - ‘omit’ : (was ‘drop’) drop missing data
4. The *missing* argument is deprecated in `lmfit 0.9.8` and will be removed in a later version. Use *nan_policy* instead, as it is consistent with the `Minimizer` class.

Examples

The model function will normally take an independent variable (generally, the first argument) and a series of arguments that are meant to be parameters for the model. Thus, a simple peak using a Gaussian defined as:

```
>>> import numpy as np
>>> def gaussian(x, amp, cen, wid):
...     return amp * np.exp(-(x-cen)**2 / wid)
```

can be turned into a `Model` with:

```
>>> gmodel = Model(gaussian)
```

this will automatically discover the names of the independent variables and parameters:

```
>>> print(gmodel.param_names, gmodel.independent_vars)
['amp', 'cen', 'wid'], ['x']
```

7.2.1 Model class Methods

`Model.eval` (*params=None*, ***kwargs*)

Evaluate the model with supplied parameters and keyword arguments.

Parameters

- **params** (*Parameters*, *optional*) – Parameters to use in `Model`.
- ****kwargs** (*optional*) – Additional keyword arguments to pass to model function.

Returns Value of model given the parameters and other arguments.

Return type `numpy.ndarray`

Notes

1. if *params* is None, the values for all parameters are expected to be provided as keyword arguments. If *params* is given, and a keyword argument for a parameter value is also given, the keyword argument will be used.
2. all non-parameter arguments for the model function, **including all the independent variables** will need to be passed in using keyword arguments.

`Model.fit` (*data*, *params=None*, *weights=None*, *method='leastsq'*, *iter_cb=None*, *scale_covar=True*, *verbose=False*, *fit_kws=None*, *nan_policy=None*, ***kwargs*)
Fit the model to the data using the supplied Parameters.

Parameters

- **data** (*array_like*) – Array of data to be fit.
- **params** (*Parameters*, *optional*) – Parameters to use in fit (default is None).
- **weights** (*array_like* of same size as *data*, *optional*) – Weights to use for the calculation of the fit residual (default is None).
- **method** (*str*, *optional*) – Name of fitting method to use (default is 'leastsq').
- **iter_cb** (*callable*, *optional*) – Callback function to call at each iteration (default is None).
- **scale_covar** (*bool*, *optional*) – Whether to automatically scale the covariance matrix when calculating uncertainties (default is True, *leastsq* method only).
- **verbose** (*bool*, *optional*) – Whether to print a message when a new parameter is added because of a hint (default is True).
- **nan_policy** (*str*, *optional*, one of 'raise' (default), 'propagate', or 'omit'.) – What to do when encountering NaNs when fitting Model.
- **fit_kws** (*dict*, *optional*) – Options to pass to the minimizer being used.
- ****kwargs** (*optional*) – Arguments to pass to the model function, possibly overriding *params*.

Returns

Return type *ModelResult*

Examples

Take *t* to be the independent variable and data to be the curve we will fit. Use keyword arguments to set initial guesses:

```
>>> result = my_model.fit(data, tau=5, N=3, t=t)
```

Or, for more control, pass a Parameters object.

```
>>> result = my_model.fit(data, params, t=t)
```

Keyword arguments override Parameters.

```
>>> result = my_model.fit(data, params, tau=5, t=t)
```


Notes

1. if *params* is None, the values for all parameters are expected to be provided as keyword arguments. If *params* is given, and a keyword argument for a parameter value is also given, the keyword argument will be used.
2. all non-parameter arguments for the model function, **including all the independent variables** will need to be passed in using keyword arguments.
3. Parameters (however passed in), are copied on input, so the original Parameter objects are unchanged, and the updated values are in the returned *ModelResult*.

`Model.guess` (*data*, ****kws**)

Guess starting values for the parameters of a Model.

This is not implemented for all models, but is available for many of the built-in models.

Parameters

- **data** (*array_like*) – Array of data to use to guess parameter values.
- ****kws** (*optional*) – Additional keyword arguments, passed to model function.

Returns params

Return type *Parameters*

Notes

Should be implemented for each model subclass to run `self.make_params()`, update starting values and return a *Parameters* object.

Raises `NotImplementedError`

`Model.make_params` (*verbose=False*, ****kwargs**)

Create a *Parameters* object for a Model.

Parameters

- **verbose** (*bool*, *optional*) – Whether to print out messages (default is False).
- ****kwargs** (*optional*) – Parameter names and initial values.

Returns params

Return type *Parameters*

Notes

1. The parameters may or may not have decent initial values for each parameter.
2. This applies any default values or parameter hints that may have been set.

`Model.set_param_hint` (*name*, ****kwargs**)

Set *hints* to use when creating parameters with `make_params()` for the named parameter.

This is especially convenient for setting initial values. The *name* can include the models *prefix* or not. The hint given can also include optional bounds and constraints (*value*, *vary*, *min*, *max*, *expr*), which will be used by `make_params()` when building default parameters.

Parameters

- **name** (*str*) – Parameter name.

- ****kwargs** (*optional*) – Arbitrary keyword arguments, needs to be a Parameter attribute. Can be any of the following:
 - **value** [float, optional] Numerical Parameter value.
 - **vary** [bool, optional] Whether the Parameter is varied during a fit (default is True).
 - **min** [float, optional] Lower bound for value (default is *-numpy.inf*, no lower bound).
 - **max** [float, optional] Upper bound for value (default is *numpy.inf*, no upper bound).
 - **expr** [str, optional] Mathematical expression used to constrain the value during the fit.

Example

```
>>> model = GaussianModel()  
>>> model.set_param_hint('sigma', min=0)
```

See *Using parameter hints*.

`Model.print_param_hints(colwidth=8)`

Print a nicely aligned text-table of parameter hints.

Parameters `colwidth` (*int*, *optional*) – Width of each column, except for first and last columns.

7.2.2 Model class Attributes

func

The model function used to calculate the model.

independent_vars

List of strings for names of the independent variables.

nan_policy

Describes what to do for NaNs and missing values. The choices are:

- 'raise': Raise a ValueError (default)
- 'propagate': Do not check for NaNs or missing values. The fit will try to ignore them.
- 'omit': Remove NaNs or missing observations in data. If pandas is installed, `pandas.isnull()` is used, otherwise `numpy.isnan()` is used.

name

Name of the model, used only in the string representation of the model. By default this will be taken from the model function.

opts

Extra keyword arguments to pass to model function. Normally this will be determined internally and should not be changed.

param_hints

Dictionary of parameter hints. See *Using parameter hints*.

param_names

List of strings of parameter names.

prefix

Prefix used for name-mangling of parameter names. The default is ''. If a particular *Model* has arguments *amplitude*, *center*, and *sigma*, these would become the parameter names. Using a prefix of 'gl_' would convert these parameter names to *gl_amplitude*, *gl_center*, and *gl_sigma*. This can be essential to avoid name collision in composite models.

7.2.3 Determining parameter names and independent variables for a function

The *Model* created from the supplied function *func* will create a *Parameters* object, and names are inferred from the function arguments, and a residual function is automatically constructed.

By default, the independent variable is taken as the first argument to the function. You can, of course, explicitly set this, and will need to do so if the independent variable is not first in the list, or if there is actually more than one independent variable.

If not specified, Parameters are constructed from all positional arguments and all keyword arguments that have a default value that is numerical, except the independent variable, of course. Importantly, the Parameters can be modified after creation. In fact, you will have to do this because none of the parameters have valid initial values. In addition, one can place bounds and constraints on Parameters, or fix their values.

7.2.4 Explicitly specifying independent_vars

As we saw for the Gaussian example above, creating a *Model* from a function is fairly easy. Let's try another one:

```
>>> from lmfit import Model
>>> import numpy as np
>>> def decay(t, tau, N):
...     return N*np.exp(-t/tau)
...
>>> decay_model = Model(decay)
>>> print(decay_model.independent_vars)
['t']
>>> params = decay_model.make_params()
>>> for pname, par in params.items():
...     print(pname, par)
...
tau <Parameter 'tau', -inf, bounds=[-inf:inf]>
N <Parameter 'N', -inf, bounds=[-inf:inf]>
```

Here, *t* is assumed to be the independent variable because it is the first argument to the function. The other function arguments are used to create parameters for the model.

If you want *tau* to be the independent variable in the above example, you can say so:

```
>>> decay_model = Model(decay, independent_vars=['tau'])
>>> print(decay_model.independent_vars)
['tau']
>>> params = decay_model.make_params()
>>> for pname, par in params.items():
...     print(pname, par)
...
t <Parameter 't', -inf, bounds=[-inf:inf]>
N <Parameter 'N', -inf, bounds=[-inf:inf]>
```

You can also supply multiple values for multi-dimensional functions with multiple independent variables. In fact, the meaning of *independent variable* here is simple, and based on how it treats arguments of the function you are modeling:

independent variable A function argument that is not a parameter or otherwise part of the model, and that will be required to be explicitly provided as a keyword argument for each fit with `Model.fit()` or evaluation with `Model.eval()`.

Note that independent variables are not required to be arrays, or even floating point numbers.

7.2.5 Functions with keyword arguments

If the model function had keyword parameters, these would be turned into Parameters if the supplied default value was a valid number (but not None, True, or False).

```
>>> def decay2(t, tau, N=10, check_positive=False):
...     if check_small:
...         arg = abs(t)/max(1.e-9, abs(tau))
...     else:
...         arg = t/tau
...     return N*np.exp(arg)
...
>>> mod = Model(decay2)
>>> params = mod.make_params()
>>> for pname, par in params.items():
...     print(pname, par)
...
t <Parameter 't', -inf, bounds=[-inf:inf]>
N <Parameter 'N', 10, bounds=[-inf:inf]>
```

Here, even though N is a keyword argument to the function, it is turned into a parameter, with the default numerical value as its initial value. By default, it is permitted to be varied in the fit – the 10 is taken as an initial value, not a fixed value. On the other hand, the `check_positive` keyword argument, was not converted to a parameter because it has a boolean default value. In some sense, `check_positive` becomes like an independent variable to the model. However, because it has a default value it is not required to be given for each model evaluation or fit, as independent variables are.

7.2.6 Defining a *prefix* for the Parameters

As we will see in the next chapter when combining models, it is sometimes necessary to decorate the parameter names in the model, but still have them be correctly used in the underlying model function. This would be necessary, for example, if two parameters in a composite model (see [Composite Models : adding \(or multiplying\) Models](#) or examples in the next chapter) would have the same name. To avoid this, we can add a *prefix* to the `Model` which will automatically do this mapping for us.

```
>>> def myfunc(x, amplitude=1, center=0, sigma=1):
... 
```

```
>>> mod = Model(myfunc, prefix='f1_')
>>> params = mod.make_params()
>>> for pname, par in params.items():
...     print(pname, par)
...
f1_amplitude <Parameter 'f1_amplitude', 1, bounds=[-inf:inf]>
```

(continues on next page)

(continued from previous page)

```
f1_center <Parameter 'f1_center', 0, bounds=[-inf:inf]>
f1_sigma <Parameter 'f1_sigma', 1, bounds=[-inf:inf]>
```

You would refer to these parameters as *f1_amplitude* and so forth, and the model will know to map these to the *amplitude* argument of *myfunc*.

7.2.7 Initializing model parameters

As mentioned above, the parameters created by `Model.make_params()` are generally created with invalid initial values of None. These values **must** be initialized in order for the model to be evaluated or used in a fit. There are four different ways to do this initialization that can be used in any combination:

1. You can supply initial values in the definition of the model function.
2. You can initialize the parameters when creating parameters with `Model.make_params()`.
3. You can give parameter hints with `Model.set_param_hint()`.
4. You can supply initial values for the parameters when you use the `Model.eval()` or `Model.fit()` methods.

Of course these methods can be mixed, allowing you to overwrite initial values at any point in the process of defining and using the model.

Initializing values in the function definition

To supply initial values for parameters in the definition of the model function, you can simply supply a default value:

```
>>> def myfunc(x, a=1, b=0):
>>>     ...
```

instead of using:

```
>>> def myfunc(x, a, b):
>>>     ...
```

This has the advantage of working at the function level – all parameters with keywords can be treated as options. It also means that some default initial value will always be available for the parameter.

Initializing values with `Model.make_params()`

When creating parameters with `Model.make_params()` you can specify initial values. To do this, use keyword arguments for the parameter names and initial values:

```
>>> mod = Model(myfunc)
>>> pars = mod.make_params(a=3, b=0.5)
```

Initializing values by setting parameter hints

After a model has been created, but prior to creating parameters with `Model.make_params()`, you can set parameter hints. These allows you to set not only a default initial value but also to set other parameter attributes controlling bounds, whether it is varied in the fit, or a constraint expression. To set a parameter hint, you can use `Model.set_param_hint()`, as with:

```
>>> mod = Model(myfunc)
>>> mod.set_param_hint('a', value=1.0)
>>> mod.set_param_hint('b', value=0.3, min=0, max=1.0)
>>> pars = mod.make_params()
```

Parameter hints are discussed in more detail in section *Using parameter hints*.

Initializing values when using a model

Finally, you can explicitly supply initial values when using a model. That is, as with `Model.make_params()`, you can include values as keyword arguments to either the `Model.eval()` or `Model.fit()` methods:

```
>>> y1 = mod.eval(x=x, a=7.0, b=-2.0)
>>> out = mod.fit(x=x, pars, a=3.0, b=0.0)
```

These approaches to initialization provide many opportunities for setting initial values for parameters. The methods can be combined, so that you can set parameter hints but then change the initial value explicitly with `Model.fit()`.

7.2.8 Using parameter hints

After a model has been created, you can give it hints for how to create parameters with `Model.make_params()`. This allows you to set not only a default initial value but also to set other parameter attributes controlling bounds, whether it is varied in the fit, or a constraint expression. To set a parameter hint, you can use `Model.set_param_hint()`, as with:

```
>>> mod = Model(myfunc)
>>> mod.set_param_hint('a', value=1.0)
>>> mod.set_param_hint('b', value=0.3, min=0, max=1.0)
```

Parameter hints are stored in a model's `param_hints` attribute, which is simply a nested dictionary:

```
>>> print(mod.param_hints)
{'a': {'value': 1}, 'b': {'max': 1.0, 'value': 0.3, 'min': 0}}
```

You can change this dictionary directly, or with the `Model.set_param_hint()` method. Either way, these parameter hints are used by `Model.make_params()` when making parameters.

An important feature of parameter hints is that you can force the creation of new parameters with parameter hints. This can be useful to make derived parameters with constraint expressions. For example to get the full-width at half maximum of a Gaussian model, one could use a parameter hint of:

```
>>> mod = Model(gaussian)
>>> mod.set_param_hint('fwhm', expr='2.3548*sigma')
```

7.2.9 Saving and Loading Models

New in version 0.9.8.

It is sometimes desirable to save a `Model` for later use outside of the code used to define the model. Lmfit provides a `save_model()` function that will save a `Model` to a file. There is also a companion `load_model()` function that can read this file and reconstruct a `Model` from it.

Saving a model turns out to be somewhat challenging. The main issue is that Python is not normally able to *serialize* a function (such as the model function making up the heart of the `Model`) in a way that can be reconstructed into a

callable Python object. The *dill* package can sometimes serialize functions, but with the limitation that it can be used only in the same version of Python. In addition, class methods used as model functions will not retain the rest of the class attributes and methods, and so may not be usable. With all those warnings, it should be emphasized that if you are willing to save or reuse the definition of the model function as Python code, then saving the Parameters and rest of the components that make up a model presents no problem.

If the *dill* package is installed, the model function will be saved using it. But because saving the model function is not always reliable, saving a model will always save the *name* of the model function. The `load_model()` takes an optional `funcdefs` argument that can contain a dictionary of function definitions with the function names as keys and function objects as values. If one of the dictionary keys matches the saved name, the corresponding function object will be used as the model function. With this approach, if you save a model and can provide the code used for the model function, the model can be saved and reliably reloaded and used.

save_model (*model*, *fname*)

Save a Model to a file.

Parameters

- **model** (*model instance*) – Model to be saved.
- **fname** (*str*) – Name of file for saved Model.

load_model (*fname*, *funcdefs*=None)

Load a saved Model from a file.

Parameters

- **fname** (*str*) – Name of file containing saved Model.
- **funcdefs** (*dict*, *optional*) – Dictionary of custom function names and definitions.

Returns

Return type *Model*

As a simple example, one can save a model as:

```
#!/usr/bin/env python

# <examples/doc_model_savemodel.py>
import numpy as np

from lmfit.model import Model, save_model

def mysine(x, amp, freq, shift):
    return amp * np.sin(x*freq + shift)

sinemodel = Model(mysine)
pars = sinemodel.make_params(amp=1, freq=0.25, shift=0)

save_model(sinemodel, 'sinemodel.sav')
# <end examples/doc_model_savemodel.py>
```

To load that later, one might do:

```
#!/usr/bin/env python

# <examples/doc_model_loadmodel.py>
import matplotlib.pyplot as plt
```

(continues on next page)

(continued from previous page)

```

import numpy as np

from lmfit.model import load_model

def mysine(x, amp, freq, shift):
    return amp * np.sin(x*freq + shift)

data = np.loadtxt('sinedata.dat')
x = data[:, 0]
y = data[:, 1]

model = load_model('sinemodel.sav', funcdefs={'mysine': mysine})
params = model.make_params(amp=3, freq=0.52, shift=0)
params['shift'].max = 1
params['shift'].min = -1
params['amp'].min = 0.0

result = model.fit(y, params, x=x)
print(result.fit_report())

plt.plot(x, y, 'bo')
plt.plot(x, result.best_fit, 'r-')
plt.show()
# <end examples/doc_model_loadmodel.py>

```

See also *Saving and Loading ModelResults*.

7.3 The `ModelResult` class

A *ModelResult* (which had been called *ModelFit* prior to version 0.9) is the object returned by `Model.fit()`. It is a subclass of *Minimizer*, and so contains many of the fit results. Of course, it knows the *Model* and the set of *Parameters* used in the fit, and it has methods to evaluate the model, to fit the data (or re-fit the data with changes to the parameters, or fit with different or modified data) and to print out a report for that fit.

While a *Model* encapsulates your model function, it is fairly abstract and does not contain the parameters or data used in a particular fit. A *ModelResult* *does* contain parameters and data as well as methods to alter and re-do fits. Thus the *Model* is the idealized model while the *ModelResult* is the messier, more complex (but perhaps more useful) object that represents a fit with a set of parameters to data with a model.

A *ModelResult* has several attributes holding values for fit results, and several methods for working with fits. These include statistics inherited from *Minimizer* useful for comparing different models, including *chisqr*, *redchi*, *aic*, and *bic*.

```

class ModelResult(model, params, data=None, weights=None, method='leastsq', fcn_args=None,
                  fcn_kws=None, iter_cb=None, scale_covar=True, nan_policy='raise', **fit_kws)
    Result from the Model fit.

```

This has many attributes and methods for viewing and working with the results of a fit using *Model*. It inherits from *Minimizer*, so that it can be used to modify and re-run the fit for the *Model*.

Parameters

- **model** (*Model*) – Model to use.
- **params** (*Parameters*) – Parameters with initial values for model.

- **data** (*array_like, optional*) – Data to be modeled.
- **weights** (*array_like, optional*) – Weights to multiply (data-model) for fit residual.
- **method** (*str, optional*) – Name of minimization method to use (default is *'leastsq'*).
- **fcn_args** (*sequence, optional*) – Positional arguments to send to model function.
- **fcn_dict** (*dict, optional*) – Keyword arguments to send to model function.
- **iter_cb** (*callable, optional*) – Function to call on each iteration of fit.
- **scale_covar** (*bool, optional*) – Whether to scale covariance matrix for uncertainty evaluation.
- **nan_policy** (*str, optional, one of 'raise' (default), 'propagate', or 'omit'.*) – What to do when encountering NaNs when fitting Model.
- ****fit_kws** (*optional*) – Keyword arguments to send to minimization routine.

7.3.1 ModelResult methods

`ModelResult.eval` (*params=None, **kwargs*)
Evaluate model function.

Parameters

- **params** (*Parameters, optional*) – Parameters to use.
- ****kwargs** (*optional*) – Options to send to `Model.eval()`

Returns `out` – Array for evaluated model.

Return type `numpy.ndarray`

`ModelResult.eval_components` (*params=None, **kwargs*)
Evaluate each component of a composite model function.

Parameters

- **params** (*Parameters, optional*) – Parameters, defaults to `ModelResult.params`
- ****kwargs** (*optional*) – Keyword arguments to pass to model function.

Returns Keys are prefixes of component models, and values are the estimated model value for each component of the model.

Return type `OrderedDict`

`ModelResult.fit` (*data=None, params=None, weights=None, method=None, nan_policy=None, **kwargs*)
Re-perform fit for a Model, given data and params.

Parameters

- **data** (*array_like, optional*) – Data to be modeled.
- **params** (*Parameters, optional*) – Parameters with initial values for model.
- **weights** (*array_like, optional*) – Weights to multiply (data-model) for fit residual.
- **method** (*str, optional*) – Name of minimization method to use (default is *'leastsq'*).

- **nan_policy** (*str, optional, one of 'raise' (default), 'propagate', or 'omit'.*) – What to do when encountering NaNs when fitting Model.
- ****kwargs** (*optional*) – Keyword arguments to send to minimization routine.

`ModelResult.fit_report(modelpars=None, show_correl=True, min_correl=0.1, sort_pars=False)`

Return a printable fit report.

The report contains fit statistics and best-fit values with uncertainties and correlations.

Parameters

- **modelpars** (*Parameters, optional*) – Known Model Parameters.
- **show_correl** (*bool, optional*) – Whether to show list of sorted correlations (default is True).
- **min_correl** (*float, optional*) – Smallest correlation in absolute value to show (default is 0.1).
- **sort_pars** (*callable, optional*) – Whether to show parameter names sorted in alphanumerical order (default is False). If False, then the parameters will be listed in the order as they were added to the Parameters dictionary. If callable, then this (one argument) function is used to extract a comparison key from each list element.

Returns **text** – Multi-line text of fit report.

Return type **str**

See also:

`fit_report()`

`ModelResult.conf_interval(**kwargs)`

Calculate the confidence intervals for the variable parameters.

Confidence intervals are calculated using the `confidence.conf_interval()` function and keyword arguments (****kwargs**) are passed to that function. The result is stored in the `ci_out` attribute so that it can be accessed without recalculating them.

`ModelResult.ci_report(with_offset=True, ndigits=5, **kwargs)`

Return a nicely formatted text report of the confidence intervals.

Parameters

- **with_offset** (*bool, optional*) – Whether to subtract best value from all other values (default is True).
- **ndigits** (*int, optional*) – Number of significant digits to show (default is 5).
- ****kwargs** (*optional*) – Keyword arguments that are passed to the `conf_interval` function.

Returns Text of formatted report on confidence intervals.

Return type **str**

`ModelResult.eval_uncertainty(params=None, sigma=1, **kwargs)`

Evaluate the uncertainty of the *model function* from the uncertainties for the best-fit parameters. This can be used to give confidence bands for the model.

Parameters

- **params** (*Parameters, optional*) – Parameters, defaults to `ModelResult.params`.

- **sigma** (*float*, *optional*) – Confidence level, i.e. how many sigma (default is 1).
- ****kwargs** (*optional*) – Values of options, independent variables, etcetera.

Returns Uncertainty at each value of the model.

Return type `numpy.ndarray`

Example

```
>>> out = model.fit(data, params, x=x)
>>> dely = out.eval_uncertainty(x=x)
>>> plt.plot(x, data)
>>> plt.plot(x, out.best_fit)
>>> plt.fill_between(x, out.best_fit-dely,
...                  out.best_fit+dely, color='#888888')
```

Notes

1. This is based on the excellent and clear example from <https://www.astro.rug.nl/software/kapteyn/kmpfittutorial.html#confidence-and-prediction-intervals>, which references the original work of: J. Wolberg, Data Analysis Using the Method of Least Squares, 2006, Springer
2. The value of sigma is number of *sigma* values, and is converted to a probability. Values of 1, 2, or 3 give probabilities of 0.6827, 0.9545, and 0.9973, respectively. If the sigma value is < 1, it is interpreted as the probability itself. That is, *sigma=1* and *sigma=0.6827* will give the same results, within precision errors.

`ModelResult.plot(datafmt='o', fitfmt='-', initfmt='-', xlabel=None, ylabel=None, yerr=None, numpoints=None, fig=None, data_kws=None, fit_kws=None, init_kws=None, ax_res_kws=None, ax_fit_kws=None, fig_kws=None, show_init=False)`

Plot the fit results and residuals using matplotlib, if available.

The method will produce a matplotlib figure with both results of the fit and the residuals plotted. If the fit model included weights, errorbars will also be plotted. To show the initial conditions for the fit, pass the argument `show_init=True`.

Parameters

- **datafmt** (*str*, *optional*) – Matplotlib format string for data points.
- **fitfmt** (*str*, *optional*) – Matplotlib format string for fitted curve.
- **initfmt** (*str*, *optional*) – Matplotlib format string for initial conditions for the fit.
- **xlabel** (*str*, *optional*) – Matplotlib format string for labeling the x-axis.
- **ylabel** (*str*, *optional*) – Matplotlib format string for labeling the y-axis.
- **yerr** (`numpy.ndarray`, *optional*) – Array of uncertainties for data array.
- **numpoints** (*int*, *optional*) – If provided, the final and initial fit curves are evaluated not only at data points, but refined to contain *numpoints* points in total.
- **fig** (`matplotlib.figure.Figure`, *optional*) – The figure to plot on. The default is None, which means use the current pyplot figure or create one if there is none.
- **data_kws** (*dict*, *optional*) – Keyword arguments passed on to the plot function for data points.

- **fit_kws** (*dict*, *optional*) – Keyword arguments passed on to the plot function for fitted curve.
- **init_kws** (*dict*, *optional*) – Keyword arguments passed on to the plot function for the initial conditions of the fit.
- **ax_res_kws** (*dict*, *optional*) – Keyword arguments for the axes for the residuals plot.
- **ax_fit_kws** (*dict*, *optional*) – Keyword arguments for the axes for the fit plot.
- **fig_kws** (*dict*, *optional*) – Keyword arguments for a new figure, if there is one being created.
- **show_init** (*bool*, *optional*) – Whether to show the initial conditions for the fit (default is False).

Returns

Return type A tuple with matplotlib's Figure and GridSpec objects.

Notes

The method combines `ModelResult.plot_fit` and `ModelResult.plot_residuals`.

If *yerr* is specified or if the fit model included weights, then `matplotlib.axes.Axes.errorbar` is used to plot the data. If *yerr* is not specified and the fit includes weights, *yerr* set to `1/self.weights`

If *fig* is None then `matplotlib.pyplot.figure(**fig_kws)` is called, otherwise *fig_kws* is ignored.

See also:

`ModelResult.plot_fit()` Plot the fit results using matplotlib.

`ModelResult.plot_residuals()` Plot the fit residuals using matplotlib.

`ModelResult.plot_fit` (*ax=None*, *datafmt='o'*, *fitfmt='-'*, *initfmt='-'*, *xlabel=None*, *ylabel=None*, *yerr=None*, *numpoints=None*, *data_kws=None*, *fit_kws=None*, *init_kws=None*, *ax_kws=None*, *show_init=False*)

Plot the fit results using matplotlib, if available.

The plot will include the data points, the initial fit curve (optional, with *show_init=True*), and the best-fit curve. If the fit model included weights or if *yerr* is specified, errorbars will also be plotted.

Parameters

- **ax** (*matplotlib.axes.Axes*, *optional*) – The axes to plot on. The default is None, which means use the current pyplot axis or create one if there is none.
- **datafmt** (*str*, *optional*) – Matplotlib format string for data points.
- **fitfmt** (*str*, *optional*) – Matplotlib format string for fitted curve.
- **initfmt** (*str*, *optional*) – Matplotlib format string for initial conditions for the fit.
- **xlabel** (*str*, *optional*) – Matplotlib format string for labeling the x-axis.
- **ylabel** (*str*, *optional*) – Matplotlib format string for labeling the y-axis.
- **yerr** (*numpy.ndarray*, *optional*) – Array of uncertainties for data array.
- **numpoints** (*int*, *optional*) – If provided, the final and initial fit curves are evaluated not only at data points, but refined to contain *numpoints* points in total.

- **data_kws** (*dict*, *optional*) – Keyword arguments passed on to the plot function for data points.
- **fit_kws** (*dict*, *optional*) – Keyword arguments passed on to the plot function for fitted curve.
- **init_kws** (*dict*, *optional*) – Keyword arguments passed on to the plot function for the initial conditions of the fit.
- **ax_kws** (*dict*, *optional*) – Keyword arguments for a new axis, if there is one being created.
- **show_init** (*bool*, *optional*) – Whether to show the initial conditions for the fit (default is False).

Returns

Return type matplotlib.axes.Axes

Notes

For details about plot format strings and keyword arguments see documentation of matplotlib.axes.Axes.plot.

If *yerr* is specified or if the fit model included weights, then matplotlib.axes.Axes.errorbar is used to plot the data. If *yerr* is not specified and the fit includes weights, *yerr* set to 1/self.weights

If *ax* is None then `matplotlib.pyplot.gca(**ax_kws)` is called.

See also:

ModelResult.plot_residuals() Plot the fit residuals using matplotlib.

ModelResult.plot() Plot the fit results and residuals using matplotlib.

`ModelResult.plot_residuals` (*ax=None*, *datafmt='o'*, *yerr=None*, *data_kws=None*, *fit_kws=None*,
ax_kws=None)

Plot the fit residuals using matplotlib, if available.

If *yerr* is supplied or if the model included weights, errorbars will also be plotted.

Parameters

- **ax** (*matplotlib.axes.Axes*, *optional*) – The axes to plot on. The default is None, which means use the current pyplot axis or create one if there is none.
- **datafmt** (*str*, *optional*) – Matplotlib format string for data points.
- **yerr** (*numpy.ndarray*, *optional*) – Array of uncertainties for data array.
- **data_kws** (*dict*, *optional*) – Keyword arguments passed on to the plot function for data points.
- **fit_kws** (*dict*, *optional*) – Keyword arguments passed on to the plot function for fitted curve.
- **ax_kws** (*dict*, *optional*) – Keyword arguments for a new axis, if there is one being created.

Returns

Return type matplotlib.axes.Axes

Notes

For details about plot format strings and keyword arguments see documentation of `matplotlib.axes.Axes.plot`.

If `yerr` is specified or if the fit model included weights, then `matplotlib.axes.Axes.errorbar` is used to plot the data. If `yerr` is not specified and the fit includes weights, `yerr` set to `1/self.weights`

If `ax` is `None` then `matplotlib.pyplot.gca(**ax_kws)` is called.

See also:

`ModelResult.plot_fit()` Plot the fit results using matplotlib.

`ModelResult.plot()` Plot the fit results and residuals using matplotlib.

7.3.2 ModelResult attributes

aic

Floating point best-fit Akaike Information Criterion statistic (see *MinimizerResult – the optimization result*).

best_fit

numpy.ndarray result of model function, evaluated at provided independent variables and with best-fit parameters.

best_values

Dictionary with parameter names as keys, and best-fit values as values.

bic

Floating point best-fit Bayesian Information Criterion statistic (see *MinimizerResult – the optimization result*).

chisqr

Floating point best-fit chi-square statistic (see *MinimizerResult – the optimization result*).

ci_out

Confidence interval data (see *Calculation of confidence intervals*) or `None` if the confidence intervals have not been calculated.

covar

numpy.ndarray (square) covariance matrix returned from fit.

data

numpy.ndarray of data to compare to model.

errorbars

Boolean for whether error bars were estimated by fit.

ier

Integer returned code from `scipy.optimize.leastsq`.

init_fit

numpy.ndarray result of model function, evaluated at provided independent variables and with initial parameters.

init_params

Initial parameters.

init_values

Dictionary with parameter names as keys, and initial values as values.

iter_cb

Optional callable function, to be called at each fit iteration. This must take arguments of `(params, iter, resid, *args, **kws)`, where `params` will have the current parameter values, `iter` the iteration,

resid the current residual array, and **args* and ***kwargs* as passed to the objective function. See [Using a Iteration Callback Function](#).

jacfcn

Optional callable function, to be called to calculate Jacobian array.

lmdif_message

String message returned from `scipy.optimize.leastsq`.

message

String message returned from `minimize()`.

method

String naming fitting method for `minimize()`.

model

Instance of [Model](#) used for model.

ndata

Integer number of data points.

nfev

Integer number of function evaluations used for fit.

nfree

Integer number of free parameters in fit.

nvarys

Integer number of independent, freely varying variables in fit.

params

Parameters used in fit. Will have best-fit values.

redchi

Floating point reduced chi-square statistic (see [MinimizerResult – the optimization result](#)).

residual

numpy.ndarray for residual.

scale_covar

Boolean flag for whether to automatically scale covariance matrix.

success

Boolean value of whether fit succeeded.

weights

numpy.ndarray (or None) of weighting values to be used in fit. If not None, it will be used as a multiplicative factor of the residual array, so that `weights*(data - fit)` is minimized in the least-squares sense.


7.3.3 Calculating uncertainties in the model function

We return to the first example above and ask not only for the uncertainties in the fitted parameters but for the range of values that those uncertainties mean for the model function itself. We can use the [ModelResult.eval_uncertainty\(\)](#) method of the model result object to evaluate the uncertainty in the model with a specified level for σ .

That is, adding:

```
dely = result.eval_uncertainty(sigma=3)
plt.fill_between(x, result.best_fit-dely, result.best_fit+dely, color="#ABABAB")
```

to the example fit to the Gaussian at the beginning of this chapter will give $3\text{-}\sigma$ bands for the best-fit Gaussian, and produce the figure below.



7.3.4 Saving and Loading ModelResults

New in version 0.9.8.

As with saving models (see section *Saving and Loading Models*), it is sometimes desirable to save a *ModelResult*, either for later use or to organize and compare different fit results. Lmfit provides a *save_modelresult()* function that will save a *ModelResult* to a file. There is also a companion *load_modelresult()* function that can read this file and reconstruct a *ModelResult* from it.

As discussed in section *Saving and Loading Models*, there are challenges to saving model functions that may make it difficult to restore a saved *ModelResult* in a way that can be used to perform a fit. Use of the optional *funcdefs* argument is generally the most reliable way to ensure that a loaded *ModelResult* can be used to evaluate the model function or redo the fit.

save_modelresult (*modelresult*, *fname*)

Save a ModelResult to a file.

Parameters

- **modelresult** (*ModelResult* instance) – ModelResult to be saved.
- **fname** (*str*) – Name of file for saved ModelResult.

load_modelresult (*fname*, *funcdefs*=None)

Load a saved ModelResult from a file.

Parameters

- **fname** (*str*) – Name of file containing saved ModelResult.
- **funcdefs** (*dict*, *optional*) – Dictionary of custom function names and definitions.

Returns

Return type *ModelResult*

An example of saving a *ModelResult* is:

```
#!/usr/bin/env python

# <examples/doc_model_savemodelresult.py>
import matplotlib.pyplot as plt
```

(continues on next page)

(continued from previous page)

```
import numpy as np

from lmfit.model import save_modelresult
from lmfit.models import GaussianModel

data = np.loadtxt('modelld_gauss.dat')
x = data[:, 0]
y = data[:, 1]

gmodel = GaussianModel()
result = gmodel.fit(y, x=x, amplitude=5, center=5, sigma=1)

save_modelresult(result, 'gauss_modelresult.sav')

print(result.fit_report())

plt.plot(x, y, 'bo')
plt.plot(x, result.init_fit, 'k--')
plt.plot(x, result.best_fit, 'r-')
plt.show()
# <end examples/doc_model_savemodelresult.py>
```

To load that later, one might do:

```
#!/usr/bin/env python

# <examples/doc_model_loadmodelresult.py>
import matplotlib.pyplot as plt
import numpy as np

from lmfit.model import load_modelresult

data = np.loadtxt('modelld_gauss.dat')
x = data[:, 0]
y = data[:, 1]

result = load_modelresult('gauss_modelresult.sav')
print(result.fit_report())

plt.plot(x, y, 'bo')
plt.plot(x, result.best_fit, 'r-')
plt.show()
# <end examples/doc_model_loadmodelresult.py>
```

7.4 Composite Models : adding (or multiplying) Models

One of the more interesting features of the *Model* class is that Models can be added together or combined with basic algebraic operations (add, subtract, multiply, and divide) to give a composite model. The composite model will have parameters from each of the component models, with all parameters being available to influence the whole model. This ability to combine models will become even more useful in the next chapter, when pre-built subclasses of *Model* are discussed. For now, we'll consider a simple example, and build a model of a Gaussian plus a line, as to model a peak with a background. For such a simple problem, we could just build a model that included both components:

```
def gaussian_plus_line(x, amp, cen, wid, slope, intercept):
    """line + 1-d gaussian"""

    gauss = (amp / (sqrt(2*pi) * wid)) * exp(-(x-cen)**2 / (2*wid**2))
    line = slope*x + intercept
    return gauss + line
```

and use that with:

```
mod = Model(gaussian_plus_line)
```

But we already had a function for a gaussian function, and maybe we'll discover that a linear background isn't sufficient which would mean the model function would have to be changed.

Instead, `lmfit` allows models to be combined into a *CompositeModel*. As an alternative to including a linear background in our model function, we could define a linear function:

```
def line(x, slope, intercept):
    """a line"""
    return slope*x + intercept
```

and build a composite model with just:

```
mod = Model(gaussian) + Model(line)
```

This model has parameters for both component models, and can be used as:

```
#!/usr/bin/env python

# <examples/doc_model_two_components.py>
import matplotlib.pyplot as plt
from numpy import exp, loadtxt, pi, sqrt

from lmfit import Model

data = loadtxt('model1d_gauss.dat')
x = data[:, 0]
y = data[:, 1] + 0.25*x - 1.0

def gaussian(x, amp, cen, wid):
    """1-d gaussian: gaussian(x, amp, cen, wid)"""
    return (amp / (sqrt(2*pi) * wid)) * exp(-(x-cen)**2 / (2*wid**2))

def line(x, slope, intercept):
    """a line"""
    return slope*x + intercept

mod = Model(gaussian) + Model(line)
pars = mod.make_params(amp=5, cen=5, wid=1, slope=0, intercept=1)

result = mod.fit(y, pars, x=x)

print(result.fit_report())
```

(continues on next page)


(continued from previous page)

```
plt.plot(x, y, 'bo')
plt.plot(x, result.init_fit, 'k--')
plt.plot(x, result.best_fit, 'r-')
plt.show()
# <end examples/doc_model_two_components.py>
```

which prints out the results:

```
[[Model]]
  (Model(gaussian) + Model(line))
[[Fit Statistics]]
  # fitting method   = leastsq
  # function evals   = 46
  # data points      = 101
  # variables        = 5
  chi-square         = 2.57855517
  reduced chi-square = 0.02685995
  Akaike info crit   = -360.457020
  Bayesian info crit = -347.381417
[[Variables]]
  amp:      8.45931062 +/- 0.12414515 (1.47%) (init = 5)
  cen:      5.65547873 +/- 0.00917678 (0.16%) (init = 5)
  wid:      0.67545524 +/- 0.00991686 (1.47%) (init = 1)
  slope:    0.26484404 +/- 0.00574892 (2.17%) (init = 0)
  intercept: -0.96860202 +/- 0.03352202 (3.46%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
  C(slope, intercept) = -0.795
  C(amp, wid)          = 0.666
  C(amp, intercept)    = -0.222
  C(amp, slope)        = -0.169
  C(cen, slope)        = -0.162
  C(wid, intercept)    = -0.148
  C(cen, intercept)    = 0.129
  C(wid, slope)        = -0.113
```

and shows the plot on the left.



On the left, data is shown in blue dots, the total fit is shown in solid red line, and the initial fit is shown as a black dashed line. The figure on the right shows again the data in blue dots, the Gaussian component as a black dashed line and the linear component as a red dashed line. It is created using the following code:

```
comps = result.eval_components()
```

(continues on next page)

(continued from previous page)

```
plt.plot(x, y, 'bo')
plt.plot(x, comps['gaussian'], 'k--')
plt.plot(x, comps['line'], 'r--')
```

The components were generated after the fit using the `ModelResult.eval_components()` method of the `result`, which returns a dictionary of the components, using keys of the model name (or *prefix* if that is set). This will use the parameter values in `result.params` and the independent variables (*x*) used during the fit. Note that while the `ModelResult` held in `result` does store the best parameters and the best estimate of the model in `result.best_fit`, the original model and parameters in `pars` are left unaltered.

You can apply this composite model to other data sets, or evaluate the model at other values of *x*. You may want to do this to give a finer or coarser spacing of data point, or to extrapolate the model outside the fitting range. This can be done with:

```
xwide = np.linspace(-5, 25, 3001)
predicted = mod.eval(x=xwide)
```

In this example, the argument names for the model functions do not overlap. If they had, the *prefix* argument to `Model` would have allowed us to identify which parameter went with which component model. As we will see in the next chapter, using composite models with the built-in models provides a simple way to build up complex models.

class CompositeModel (*left*, *right*, *op*[, ***kws*])

Combine two models (*left* and *right*) with a binary operator (*op*) into a CompositeModel.

Normally, one does not have to explicitly create a *CompositeModel*, but can use normal Python operators `+`, `-`, `*`, and `/` to combine components as in:

```
>>> mod = Model(fcn1) + Model(fcn2) * Model(fcn3)
```

Parameters

- **left** (*Model*) – Left-hand model.
- **right** (*Model*) – Right-hand model.
- **op** (*callable binary operator*) – Operator to combine *left* and *right* models.
- ****kws** (*optional*) – Additional keywords are passed to *Model* when creating this new model.

Notes

1. The two models must use the same independent variable.

Note that when using built-in Python binary operators, a *CompositeModel* will automatically be constructed for you. That is, doing:

```
mod = Model(fcn1) + Model(fcn2) * Model(fcn3)
```

will create a *CompositeModel*. Here, *left* will be `Model(fcn1)`, *op* will be `operator.add()`, and *right* will be another *CompositeModel* that has a *left* attribute of `Model(fcn2)`, an *op* of `operator.mul()`, and a *right* of `Model(fcn3)`.

To use a binary operator other than `+`, `-`, `*`, or `/` you can explicitly create a *CompositeModel* with the appropriate binary operator. For example, to convolve two models, you could define a simple convolution function, perhaps as:

```
import numpy as np

def convolve(dat, kernel):
    """simple convolution of two arrays"""
    npts = min(len(dat), len(kernel))
    pad = np.ones(npts)
    tmp = np.concatenate((pad*dat[0], dat, pad*dat[-1]))
    out = np.convolve(tmp, kernel, mode='valid')
    noff = int((len(out) - npts) / 2)
    return out[noff:][:npts]
```

which extends the data in both directions so that the convolving kernel function gives a valid result over the data range. Because this function takes two array arguments and returns an array, it can be used as the binary operator. A full script using this technique is here:

```
#!/usr/bin/env python

# <examples/doc_model_composite.py>
import matplotlib.pyplot as plt
import numpy as np

from lmfit import CompositeModel, Model
from lmfit.lineshapes import gaussian, step

# create data from broadened step
npts = 201
x = np.linspace(0, 10, npts)
y = step(x, amplitude=12.5, center=4.5, sigma=0.88, form='erf')
np.random.seed(0)
y = y + np.random.normal(size=npts, scale=0.35)

def jump(x, mid):
    """heaviside step function"""
    o = np.zeros(len(x))
    imid = max(np.where(x <= mid)[0])
    o[imid:] = 1.0
    return o

def convolve(arr, kernel):
    """simple convolution of two arrays"""
    npts = min(len(arr), len(kernel))
    pad = np.ones(npts)
    tmp = np.concatenate((pad*arr[0], arr, pad*arr[-1]))
    out = np.convolve(tmp, kernel, mode='valid')
    noff = int((len(out) - npts) / 2)
    return out[noff:noff+npts]

# create Composite Model using the custom convolution operator
mod = CompositeModel(Model(jump), Model(gaussian), convolve)

pars = mod.make_params(amplitude=1, center=3.5, sigma=1.5, mid=5.0)

# 'mid' and 'center' should be completely correlated, and 'mid' is
# used as an integer index, so a very poor fit variable:
```

(continues on next page)

(continued from previous page)

```

pars['mid'].vary = False

# fit this model to data array y
result = mod.fit(y, params=pars, x=x)

print(result.fit_report())

plot_components = False

# plot results
plt.plot(x, y, 'bo')
if plot_components:
    # generate components
    comps = result.eval_components(x=x)
    plt.plot(x, 10*comps['jump'], 'k--')
    plt.plot(x, 10*comps['gaussian'], 'r-')
else:
    plt.plot(x, result.init_fit, 'k--')
    plt.plot(x, result.best_fit, 'r-')
plt.show()
# <end examples/doc_model_composite.py>

```


which prints out the results:

```

[[Model]]
  (Model(jump) <function convolve at 0x114e6e1e0> Model(gaussian))
[[Fit Statistics]]
  # fitting method   = leastsq
  # function evals   = 25
  # data points      = 201
  # variables        = 3
  chi-square         = 24.7562335
  reduced chi-square = 0.12503148
  Akaike info crit   = -414.939746
  Bayesian info crit = -405.029832
[[Variables]]
  mid:      5 (fixed)
  sigma:    0.59576118 +/- 0.01348582 (2.26%) (init = 1.5)
  center:   4.50853671 +/- 0.00973231 (0.22%) (init = 3.5)
  amplitude: 0.62508459 +/- 0.00189732 (0.30%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
  C(center, amplitude) = 0.329
  C(sigma, amplitude)  = 0.268

```

and shows the plots:



Using composite models with built-in or custom operators allows you to build complex models from testable sub-components.

BUILT-IN FITTING MODELS IN THE MODELS MODULE

Lmfit provides several built-in fitting models in the `models` module. These pre-defined models each subclass from the `Model` class of the previous chapter and wrap relatively well-known functional forms, such as Gaussians, Lorentzian, and Exponentials that are used in a wide range of scientific domains. In fact, all the models are based on simple, plain Python functions defined in the `lineshapes` module. In addition to wrapping a function into a `Model`, these models also provide a `guess()` method that is intended to give a reasonable set of starting values from a data array that closely approximates the data to be fit.

As shown in the previous chapter, a key feature of the `Model` class is that models can easily be combined to give a composite `CompositeModel`. Thus, while some of the models listed here may seem pretty trivial (notably, `ConstantModel` and `LinearModel`), the main point of having these is to be able to use them in composite models. For example, a Lorentzian plus a linear background might be represented as:

```
>>> from lmfit.models import LinearModel, LorentzianModel
>>> peak = LorentzianModel()
>>> background = LinearModel()
>>> model = peak + background
```

All the models listed below are one dimensional, with an independent variable named `x`. Many of these models represent a function with a distinct peak, and so share common features. To maintain uniformity, common parameter names are used whenever possible. Thus, most models have a parameter called `amplitude` that represents the overall height (or area of) a peak or function, a `center` parameter that represents a peak centroid position, and a `sigma` parameter that gives a characteristic width. Many peak shapes also have a parameter `fwhm` (constrained by `sigma`) giving the full width at half maximum and a parameter `height` (constrained by `sigma` and `amplitude`) to give the maximum peak height.

After a list of built-in models, a few examples of their use are given.

8.1 Peak-like models

There are many peak-like models available. These include `GaussianModel`, `LorentzianModel`, `VoigtModel` and some less commonly used variations. The `guess()` methods for all of these make a fairly crude guess for the value of `amplitude`, but also set a lower bound of 0 on the value of `sigma`.

8.1.1 GaussianModel

class GaussianModel (*independent_vars=['x']*, *prefix=""*, *nan_policy='raise'*, ***kwargs*)

A model based on a Gaussian or normal distribution lineshape (see https://en.wikipedia.org/wiki/Normal_distribution), with three Parameters: `amplitude`, `center`, and `sigma`. In addition, parameters `fwhm` and

`height` are included as constraints to report full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma) = \frac{A}{\sigma\sqrt{2\pi}} e^{[-(x-\mu)^2/2\sigma^2]}$$

where the parameter `amplitude` corresponds to A , `center` to μ , and `sigma` to σ . The full width at half maximum is $2\sigma\sqrt{2\ln 2}$, approximately 2.3548σ .

Parameters

- **`independent_vars`** (`['x']`) – Arguments to `func` that are independent variables.
- **`prefix`** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **`nan_policy`** (`str`, *optional*) – How to handle NaN and missing values in data. Must be one of: ‘raise’ (default), ‘propagate’, or ‘omit’. See Notes below.
- **`missing`** (`str`, *optional*) – Synonym for ‘`nan_policy`’ for backward compatibility.
- **`**kwargs`** (*optional*) – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - ‘raise’ : Raise a `ValueError` (default)
 - ‘propagate’ : do nothing
 - ‘omit’ : (was ‘drop’) drop missing data
2. The `missing` argument is deprecated in `lmfit` 0.9.8 and will be removed in a later version. Use `nan_policy` instead, as it is consistent with the `Minimizer` class.

8.1.2 LorentzianModel

`class LorentzianModel` (`independent_vars=['x']`, `prefix=""`, `nan_policy='raise'`, `**kwargs`)

A model based on a Lorentzian or Cauchy-Lorentz distribution function (see https://en.wikipedia.org/wiki/Cauchy_distribution), with three Parameters: `amplitude`, `center`, and `sigma`. In addition, parameters `fwhm` and `height` are included as constraints to report full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma) = \frac{A}{\pi} \left[\frac{\sigma}{(x - \mu)^2 + \sigma^2} \right]$$

where the parameter `amplitude` corresponds to A , `center` to μ , and `sigma` to σ . The full width at half maximum is 2σ .

Parameters

- **`independent_vars`** (`['x']`) – Arguments to `func` that are independent variables.
- **`prefix`** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **`nan_policy`** (`str`, *optional*) – How to handle NaN and missing values in data. Must be one of: ‘raise’ (default), ‘propagate’, or ‘omit’. See Notes below.
- **`missing`** (`str`, *optional*) – Synonym for ‘`nan_policy`’ for backward compatibility.
- **`**kwargs`** (*optional*) – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - `'raise'` : Raise a `ValueError` (default)
 - `'propagate'` : do nothing
 - `'omit'` : (was `'drop'`) drop missing data
2. The `missing` argument is deprecated in `lmfit` 0.9.8 and will be removed in a later version. Use `nan_policy` instead, as it is consistent with the `Minimizer` class.

8.1.3 VoigtModel

class VoigtModel (*independent_vars=['x']*, *prefix=""*, *nan_policy='raise'*, ***kwargs*)

A model based on a Voigt distribution function (see https://en.wikipedia.org/wiki/Voigt_profile), with four Parameters: `amplitude`, `center`, `sigma`, and `gamma`. By default, `gamma` is constrained to have a value equal to `sigma`, though it can be varied independently. In addition, parameters `fwhm` and `height` are included as constraints to report full width at half maximum and maximum peak height, respectively. The definition for the Voigt function used here is

$$f(x; A, \mu, \sigma, \gamma) = \frac{\text{ARe}[w(z)]}{\sigma\sqrt{2\pi}}$$

where

$$z = \frac{x - \mu + i\gamma}{\sigma\sqrt{2}}$$

$$w(z) = e^{-z^2} \text{erfc}(-iz)$$

and `erfc()` is the complementary error function. As above, `amplitude` corresponds to A , `center` to μ , and `sigma` to σ . The parameter `gamma` corresponds to γ . If `gamma` is kept at the default value (constrained to `sigma`), the full width at half maximum is approximately 3.6013σ .

Parameters

- **`independent_vars`** (`['x']`) – Arguments to `func` that are independent variables.
- **`prefix`** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **`nan_policy`** (`str`, *optional*) – How to handle NaN and missing values in data. Must be one of: `'raise'` (default), `'propagate'`, or `'omit'`. See Notes below.
- **`missing`** (`str`, *optional*) – Synonym for `'nan_policy'` for backward compatibility.
- **`**kwargs`** (*optional*) – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - `'raise'` : Raise a `ValueError` (default)
 - `'propagate'` : do nothing
 - `'omit'` : (was `'drop'`) drop missing data
2. The `missing` argument is deprecated in `lmfit` 0.9.8 and will be removed in a later version. Use `nan_policy` instead, as it is consistent with the `Minimizer` class.

8.1.4 PseudoVoigtModel

class PseudoVoigtModel (*independent_vars*=['x'], *prefix*="", *nan_policy*='raise', ***kwargs*)

A model based on a pseudo-Voigt distribution function (see https://en.wikipedia.org/wiki/Voigt_profile#Pseudo-Voigt_Approximation), which is a weighted sum of a Gaussian and Lorentzian distribution function that share values for amplitude (A), center (μ) and full width at half maximum *fwhm* (and so have constrained values of *sigma* (σ) and *height* (maximum peak height)). A parameter *fraction* (α) controls the relative weight of the Gaussian and Lorentzian components, giving the full definition of

$$f(x; A, \mu, \sigma, \alpha) = \frac{(1 - \alpha)A}{\sigma_g \sqrt{2\pi}} e^{-(x-\mu)^2 / 2\sigma_g^2} + \frac{\alpha A}{\pi} \left[\frac{\sigma}{(x - \mu)^2 + \sigma^2} \right]$$

where $\sigma_g = \sigma / \sqrt{2 \ln 2}$ so that the full width at half maximum of each component and of the sum is 2σ . The *guess()* function always sets the starting value for *fraction* at 0.5.

Parameters

- **independent_vars** (['x']) – Arguments to func that are independent variables.
- **prefix** (*str*, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (*str*, *optional*) – How to handle NaN and missing values in data. Must be one of: 'raise' (default), 'propagate', or 'omit'. See Notes below.
- **missing** (*str*, *optional*) – Synonym for 'nan_policy' for backward compatibility.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - 'raise' : Raise a ValueError (default)
 - 'propagate' : do nothing
 - 'omit' : (was 'drop') drop missing data
2. The *missing* argument is deprecated in lmfit 0.9.8 and will be removed in a later version. Use *nan_policy* instead, as it is consistent with the Minimizer class.

8.1.5 MoffatModel

class MoffatModel (*independent_vars*=['x'], *prefix*="", *nan_policy*='raise', ***kwargs*)

A model based on the Moffat distribution function (see https://en.wikipedia.org/wiki/Moffat_distribution), with four Parameters: amplitude (A), center (μ), a width parameter *sigma* (σ) and an exponent *beta* (β). In addition, parameters *fwhm* and *height* are included as constraints to report full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma, \beta) = A \left[\left(\frac{x - \mu}{\sigma} \right)^2 + 1 \right]^{-\beta}$$

the full width have maximum is $2\sigma\sqrt{2^{1/\beta} - 1}$. The *guess()* function always sets the starting value for *beta* to 1.

Note that for ($\beta = 1$) the Moffat has a Lorentzian shape.

Parameters

- **independent_vars** (['x']) – Arguments to func that are independent variables.
- **prefix** (*str*, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (*str*, *optional*) – How to handle NaN and missing values in data. Must be one of: 'raise' (default), 'propagate', or 'omit'. See Notes below.
- **missing** (*str*, *optional*) – Synonym for 'nan_policy' for backward compatibility.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. nan_policy sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - 'raise' : Raise a ValueError (default)
 - 'propagate' : do nothing
 - 'omit' : (was 'drop') drop missing data
2. The *missing* argument is deprecated in lmfit 0.9.8 and will be removed in a later version. Use *nan_policy* instead, as it is consistent with the Minimizer class.

8.1.6 Pearson7Model

class Pearson7Model (*independent_vars*=['x'], *prefix*='', *nan_policy*='raise', ***kwargs*)

A model based on a Pearson VII distribution (see https://en.wikipedia.org/wiki/Pearson_distribution#The_Pearson_type_VII_distribution), with four parameters: amplitude (*A*), center (μ), sigma (σ), and exponent (*m*). In addition, parameters *fwhm* and *height* are included as constraints to report full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma, m) = \frac{A}{\sigma \beta(m - \frac{1}{2}, \frac{1}{2})} \left[1 + \frac{(x - \mu)^2}{\sigma^2} \right]^{-m}$$

where β is the beta function (see [scipy.special.beta](#)) The `guess()` function always gives a starting value for exponent of 1.5. In addition, parameters *fwhm* and *height* are included as constraints to report full width at half maximum and maximum peak height, respectively.

Parameters

- **independent_vars** (['x']) – Arguments to func that are independent variables.
- **prefix** (*str*, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (*str*, *optional*) – How to handle NaN and missing values in data. Must be one of: 'raise' (default), 'propagate', or 'omit'. See Notes below.
- **missing** (*str*, *optional*) – Synonym for 'nan_policy' for backward compatibility.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. nan_policy sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - 'raise' : Raise a ValueError (default)

- ‘propagate’ : do nothing
- ‘omit’ : (was ‘drop’) drop missing data

2. The *missing* argument is deprecated in lmfit 0.9.8 and will be removed in a later version. Use *nan_policy* instead, as it is consistent with the Minimizer class.

8.1.7 StudentstModel

class StudentstModel (*independent_vars*=[‘x’], *prefix*=‘’, *nan_policy*=‘raise’, ***kwargs*)

A model based on a Student’s t distribution function (see https://en.wikipedia.org/wiki/Student%27s_t-distribution), with three Parameters: amplitude (*A*), center (μ) and sigma (σ). In addition, parameters *fwhm* and *height* are included as constraints to report full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma) = \frac{A\Gamma(\frac{\sigma+1}{2})}{\sqrt{\sigma\pi}\Gamma(\frac{\sigma}{2})} \left[1 + \frac{(x - \mu)^2}{\sigma} \right]^{-\frac{\sigma+1}{2}}$$

where $\Gamma(x)$ is the gamma function.

Parameters

- **independent_vars** ([‘x’]) – Arguments to func that are independent variables.
- **prefix** (*str*, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (*str*, *optional*) – How to handle NaN and missing values in data. Must be one of: ‘raise’ (default), ‘propagate’, or ‘omit’. See Notes below.
- **missing** (*str*, *optional*) – Synonym for ‘nan_policy’ for backward compatibility.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:

- ‘raise’ : Raise a ValueError (default)
- ‘propagate’ : do nothing
- ‘omit’ : (was ‘drop’) drop missing data

2. The *missing* argument is deprecated in lmfit 0.9.8 and will be removed in a later version. Use *nan_policy* instead, as it is consistent with the Minimizer class.

8.1.8 BreitWignerModel

class BreitWignerModel (*independent_vars*=[‘x’], *prefix*=‘’, *nan_policy*=‘raise’, ***kwargs*)

A model based on a Breit-Wigner-Fano function (see https://en.wikipedia.org/wiki/Fano_resonance), with four Parameters: amplitude (*A*), center (μ), sigma (σ), and q (*q*). In addition, parameters *fwhm* and *height* are included as constraints to report full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma, q) = \frac{A(q\sigma/2 + x - \mu)^2}{(\sigma/2)^2 + (x - \mu)^2}$$

Parameters

- **independent_vars** (`['x']`) – Arguments to func that are independent variables.
- **prefix** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`str`, *optional*) – How to handle NaN and missing values in data. Must be one of: 'raise' (default), 'propagate', or 'omit'. See Notes below.
- **missing** (`str`, *optional*) – Synonym for 'nan_policy' for backward compatibility.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - 'raise' : Raise a `ValueError` (default)
 - 'propagate' : do nothing
 - 'omit' : (was 'drop') drop missing data
2. The *missing* argument is deprecated in `lmfit` 0.9.8 and will be removed in a later version. Use *nan_policy* instead, as it is consistent with the `Minimizer` class.

8.1.9 LognormalModel

class LognormalModel (*independent_vars=['x'], prefix="", nan_policy='raise', **kwargs*)

A model based on the Log-normal distribution function (see <https://en.wikipedia.org/wiki/Lognormal>), with three Parameters amplitude (A), center (μ) and sigma (σ). In addition, parameters `fwhm` and `height` are included as constraints to report full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma) = \frac{A}{\sigma\sqrt{2\pi}} \frac{e^{-(\ln(x)-\mu)^2/2\sigma^2}}{x}$$

Parameters

- **independent_vars** (`['x']`) – Arguments to func that are independent variables.
- **prefix** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`str`, *optional*) – How to handle NaN and missing values in data. Must be one of: 'raise' (default), 'propagate', or 'omit'. See Notes below.
- **missing** (`str`, *optional*) – Synonym for 'nan_policy' for backward compatibility.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - 'raise' : Raise a `ValueError` (default)
 - 'propagate' : do nothing
 - 'omit' : (was 'drop') drop missing data
2. The *missing* argument is deprecated in `lmfit` 0.9.8 and will be removed in a later version. Use *nan_policy* instead, as it is consistent with the `Minimizer` class.

8.1.10 DampedOscillatorModel

class DampedOscillatorModel (*independent_vars=['x'], prefix="", nan_policy='raise', **kwargs*)

A model based on the Damped Harmonic Oscillator Amplitude (see https://en.wikipedia.org/wiki/Harmonic_oscillator#Amplitude_part), with three Parameters: amplitude (A), center (μ) and sigma (σ). In addition, parameters `fwhm` and `height` are included as constraints to report full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma) = \frac{A}{\sqrt{[1 - (x/\mu)^2]^2 + (2\sigma x/\mu)^2}}$$

Parameters

- **independent_vars** (`['x']`) – Arguments to func that are independent variables.
- **prefix** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`str`, *optional*) – How to handle NaN and missing values in data. Must be one of: 'raise' (default), 'propagate', or 'omit'. See Notes below.
- **missing** (`str`, *optional*) – Synonym for 'nan_policy' for backward compatibility.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - 'raise' : Raise a ValueError (default)
 - 'propagate' : do nothing
 - 'omit' : (was 'drop') drop missing data
2. The `missing` argument is deprecated in lmfit 0.9.8 and will be removed in a later version. Use `nan_policy` instead, as it is consistent with the Minimizer class.

8.1.11 DampedHarmonicOscillatorModel

class DampedHarmonicOscillatorModel (*independent_vars=['x'], prefix="", nan_policy='raise', **kwargs*)

A model based on a variation of the Damped Harmonic Oscillator (see https://en.wikipedia.org/wiki/Harmonic_oscillator), following the definition given in DAVE/PAN (see <https://www.ncnr.nist.gov/dave/>) with four Parameters: amplitude (A), center (μ), sigma (σ), and gamma (γ). In addition, parameters `fwhm` and `height` are included as constraints to report full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma, \gamma) = \frac{A\sigma}{\pi[1 - \exp(-x/\gamma)]} \left[\frac{1}{(x - \mu)^2 + \sigma^2} - \frac{1}{(x + \mu)^2 + \sigma^2} \right]$$

where $\gamma = kT$ k is the Boltzmann constant in evK^{-1} and T is the temperature in K .

Parameters

- **independent_vars** (`['x']`) – Arguments to func that are independent variables.
- **prefix** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.

- **nan_policy**(*str*, *optional*) – How to handle NaN and missing values in data. Must be one of: ‘raise’ (default), ‘propagate’, or ‘omit’. See Notes below.
- **missing**(*str*, *optional*) – Synonym for ‘nan_policy’ for backward compatibility.
- ****kwargs**(*optional*) – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - ‘raise’ : Raise a `ValueError` (default)
 - ‘propagate’ : do nothing
 - ‘omit’ : (was ‘drop’) drop missing data
2. The `missing` argument is deprecated in `lmfit` 0.9.8 and will be removed in a later version. Use `nan_policy` instead, as it is consistent with the `Minimizer` class.

8.1.12 ExponentialGaussianModel

class ExponentialGaussianModel(*independent_vars*=[‘x’], *prefix*=”, *nan_policy*=‘raise’,
***kwargs*)

A model of an Exponentially modified Gaussian distribution (see https://en.wikipedia.org/wiki/Exponentially_modified_Gaussian_distribution) with four Parameters `amplitude` (A), `center` (μ), `sigma` (σ), and `gamma` (γ). In addition, parameters `fwhm` and `height` are included as constraints to report full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma, \gamma) = \frac{A\gamma}{2} \exp[\gamma(\mu - x + \gamma\sigma^2/2)] \operatorname{erfc}\left(\frac{\mu + \gamma\sigma^2 - x}{\sqrt{2}\sigma}\right)$$

where `erfc()` is the complementary error function.

Parameters

- **independent_vars** ([‘x’]) – Arguments to `func` that are independent variables.
- **prefix**(*str*, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy**(*str*, *optional*) – How to handle NaN and missing values in data. Must be one of: ‘raise’ (default), ‘propagate’, or ‘omit’. See Notes below.
- **missing**(*str*, *optional*) – Synonym for ‘nan_policy’ for backward compatibility.
- ****kwargs**(*optional*) – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - ‘raise’ : Raise a `ValueError` (default)
 - ‘propagate’ : do nothing
 - ‘omit’ : (was ‘drop’) drop missing data
2. The `missing` argument is deprecated in `lmfit` 0.9.8 and will be removed in a later version. Use `nan_policy` instead, as it is consistent with the `Minimizer` class.

8.1.13 SkewedGaussianModel

class SkewedGaussianModel (*independent_vars=['x']*, *prefix=""*, *nan_policy='raise'*, ***kwargs*)

A variation of the Exponential Gaussian, this uses a skewed normal distribution (see https://en.wikipedia.org/wiki/Skewed_normal_distribution), with Parameters amplitude (A), center (μ), sigma (σ), and gamma (γ). In addition, parameters *fwhm* and *height* are included as constraints to report full width at half maximum and maximum peak height, respectively.

$$f(x; A, \mu, \sigma, \gamma) = \frac{A}{\sigma\sqrt{2\pi}} e^{[-(x-\mu)^2/2\sigma^2]} \left\{ 1 + \operatorname{erf}\left[\frac{\gamma(x-\mu)}{\sigma\sqrt{2}}\right] \right\}$$

where $\operatorname{erf}()$ is the error function.

Parameters

- **independent_vars** (*['x']*) – Arguments to func that are independent variables.
- **prefix** (*str*, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (*str*, *optional*) – How to handle NaN and missing values in data. Must be one of: 'raise' (default), 'propagate', or 'omit'. See Notes below.
- **missing** (*str*, *optional*) – Synonym for 'nan_policy' for backward compatibility.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:

- 'raise' : Raise a ValueError (default)
- 'propagate' : do nothing
- 'omit' : (was 'drop') drop missing data

2. The *missing* argument is deprecated in lmfit 0.9.8 and will be removed in a later version. Use *nan_policy* instead, as it is consistent with the Minimizer class.

8.1.14 DoniachModel

class DoniachModel (*independent_vars=['x']*, *prefix=""*, *nan_policy='raise'*, ***kwargs*)

A model of an Doniach Sunjic asymmetric lineshape (see https://www.casaxps.com/help_manual/line_shapes.htm), used in photo-emission, with four Parameters amplitude (A), center (μ), sigma (σ), and gamma (γ). In addition, parameter *height* is included as a constraint.

$$f(x; A, \mu, \sigma, \gamma) = \frac{A}{\sigma^{1-\gamma}} \frac{\cos[\pi\gamma/2 + (1-\gamma) \arctan((x-\mu)/\sigma)]}{[1 + (x-\mu)/\sigma]^{(1-\gamma)/2}}$$

Parameters

- **independent_vars** (*['x']*) – Arguments to func that are independent variables.
- **prefix** (*str*, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (*str*, *optional*) – How to handle NaN and missing values in data. Must be one of: 'raise' (default), 'propagate', or 'omit'. See Notes below.

- **missing**(*str*, *optional*) – Synonym for ‘nan_policy’ for backward compatibility.
- ****kwargs**(*optional*) – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - ‘raise’ : Raise a `ValueError` (default)
 - ‘propagate’ : do nothing
 - ‘omit’ : (was ‘drop’) drop missing data
2. The *missing* argument is deprecated in `lmfit` 0.9.8 and will be removed in a later version. Use *nan_policy* instead, as it is consistent with the `Minimizer` class.

8.2 Linear and Polynomial Models

These models correspond to polynomials of some degree. Of course, `lmfit` is a very inefficient way to do linear regression (see `numpy.polyfit` or `scipy.stats.linregress`), but these models may be useful as one of many components of a composite model.

8.2.1 ConstantModel

class ConstantModel(*independent_vars*=[‘x’], *prefix*=”, *nan_policy*=‘raise’, ***kwargs*)
Constant model, with a single Parameter: *c*.

Note that this is ‘constant’ in the sense of having no dependence on the independent variable *x*, not in the sense of being non-varying. To be clear, *c* will be a Parameter that will be varied in the fit (by default, of course).

Parameters

- **independent_vars**([‘x’]) – Arguments to `func` that are independent variables.
- **prefix**(*str*, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy**(*str*, *optional*) – How to handle NaN and missing values in data. Must be one of: ‘raise’ (default), ‘propagate’, or ‘omit’. See Notes below.
- **missing**(*str*, *optional*) – Synonym for ‘nan_policy’ for backward compatibility.
- ****kwargs**(*optional*) – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - ‘raise’ : Raise a `ValueError` (default)
 - ‘propagate’ : do nothing
 - ‘omit’ : (was ‘drop’) drop missing data
2. The *missing* argument is deprecated in `lmfit` 0.9.8 and will be removed in a later version. Use *nan_policy* instead, as it is consistent with the `Minimizer` class.

8.2.2 LinearModel

class LinearModel (*independent_vars*=['x'], *prefix*="", *nan_policy*='raise', ***kwargs*)
Linear model, with two Parameters *intercept* and *slope*.

Defined as:

$$f(x; m, b) = mx + b$$

with *slope* for *m* and *intercept* for *b*.

Parameters

- **independent_vars** (['x']) – Arguments to func that are independent variables.
- **prefix** (*str*, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (*str*, *optional*) – How to handle NaN and missing values in data. Must be one of: 'raise' (default), 'propagate', or 'omit'. See Notes below.
- **missing** (*str*, *optional*) – Synonym for 'nan_policy' for backward compatibility.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - 'raise' : Raise a ValueError (default)
 - 'propagate' : do nothing
 - 'omit' : (was 'drop') drop missing data
2. The *missing* argument is deprecated in lmfit 0.9.8 and will be removed in a later version. Use *nan_policy* instead, as it is consistent with the Minimizer class.

8.2.3 QuadraticModel

class QuadraticModel (*independent_vars*=['x'], *prefix*="", *nan_policy*='raise', ***kwargs*)
A quadratic model, with three Parameters *a*, *b*, and *c*.

Defined as:

$$f(x; a, b, c) = ax^2 + bx + c$$

Parameters

- **independent_vars** (['x']) – Arguments to func that are independent variables.
- **prefix** (*str*, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (*str*, *optional*) – How to handle NaN and missing values in data. Must be one of: 'raise' (default), 'propagate', or 'omit'. See Notes below.
- **missing** (*str*, *optional*) – Synonym for 'nan_policy' for backward compatibility.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - `'raise'` : Raise a `ValueError` (default)
 - `'propagate'` : do nothing
 - `'omit'` : (was `'drop'`) drop missing data
2. The `missing` argument is deprecated in `lmfit 0.9.8` and will be removed in a later version. Use `nan_policy` instead, as it is consistent with the `Minimizer` class.

8.2.4 PolynomialModel

class PolynomialModel (*degree*, *independent_vars*=['x'], *prefix*="", *nan_policy*='raise', ***kwargs*)
A polynomial model with up to 7 Parameters, specified by *degree*.

$$f(x; c_0, c_1, \dots, c_7) = \sum_{i=0,7} c_i x^i$$

with parameters `c0`, `c1`, ..., `c7`. The supplied *degree* will specify how many of these are actual variable parameters. This uses `numpy.polyval` for its calculation of the polynomial.

Parameters

- **`independent_vars`** (`['x']`) – Arguments to func that are independent variables.
- **`prefix`** (*str*, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **`nan_policy`** (*str*, *optional*) – How to handle NaN and missing values in data. Must be one of: `'raise'` (default), `'propagate'`, or `'omit'`. See Notes below.
- **`missing`** (*str*, *optional*) – Synonym for `'nan_policy'` for backward compatibility.
- **`**kwargs`** (*optional*) – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - `'raise'` : Raise a `ValueError` (default)
 - `'propagate'` : do nothing
 - `'omit'` : (was `'drop'`) drop missing data
2. The `missing` argument is deprecated in `lmfit 0.9.8` and will be removed in a later version. Use `nan_policy` instead, as it is consistent with the `Minimizer` class.

8.3 Step-like models

Two models represent step-like functions, and share many characteristics.

8.3.1 StepModel

class StepModel (*independent_vars*=['x'], *prefix*="", *nan_policy*='raise', *form*='linear', ***kwargs*)

A model based on a Step function, with three Parameters: amplitude (A), center (μ) and sigma (σ).

There are four choices for *form*:

- `linear` (the default)
- `atan` or `arctan` for an arc-tangent function
- `erf` for an error function
- `logistic` for a logistic function (see https://en.wikipedia.org/wiki/Logistic_function)

The step function starts with a value 0, and ends with a value of A rising to $A/2$ at μ , with σ setting the characteristic width. The functional forms are defined as:

$$\begin{aligned}f(x; A, \mu, \sigma, \text{form} = \text{'linear'}) &= A \min [1, \max (0, \alpha)] \\f(x; A, \mu, \sigma, \text{form} = \text{'arctan'}) &= A[1/2 + \arctan (\alpha)/\pi] \\f(x; A, \mu, \sigma, \text{form} = \text{'erf'}) &= A[1 + \text{erf}(\alpha)]/2 \\f(x; A, \mu, \sigma, \text{form} = \text{'logistic'}) &= A[1 - \frac{1}{1 + e^\alpha}]\end{aligned}$$

where $\alpha = (x - \mu)/\sigma$.

Parameters

- **`independent_vars`** (`['x']`) – Arguments to `func` that are independent variables.
- **`prefix`** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **`nan_policy`** (`str`, *optional*) – How to handle NaN and missing values in data. Must be one of: 'raise' (default), 'propagate', or 'omit'. See Notes below.
- **`missing`** (`str`, *optional*) – Synonym for 'nan_policy' for backward compatibility.
- **`**kwargs`** (*optional*) – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - 'raise' : Raise a `ValueError` (default)
 - 'propagate' : do nothing
 - 'omit' : (was 'drop') drop missing data
2. The *missing* argument is deprecated in `lmfit` 0.9.8 and will be removed in a later version. Use *nan_policy* instead, as it is consistent with the `Minimizer` class.

8.3.2 RectangleModel

class RectangleModel (*independent_vars*=['x'], *prefix*="", *nan_policy*='raise', *form*='linear', ***kwargs*)

A model based on a Step-up and Step-down function, with five Parameters: amplitude (A), center1 (μ_1), center2 (μ_2), *sigma1* (σ_1) and *sigma2* (σ_2).

There are four choices for *form*, which is used for both the Step up and the Step down:

- `linear` (the default)
- `atan` or `arctan` for an arc-tangent function
- `erf` for an error function
- `logistic` for a logistic function (see https://en.wikipedia.org/wiki/Logistic_function)

The function starts with a value 0, transitions to a value of A , taking the value $A/2$ at μ_1 , with σ_1 setting the characteristic width. The function then transitions again to the value $A/2$ at μ_2 , with σ_2 setting the characteristic width. The functional forms are defined as:

$$\begin{aligned} f(x; A, \mu, \sigma, \text{form} = \text{'linear'}) &= A\{\min[1, \max(0, \alpha_1)] + \min[-1, \max(0, \alpha_2)]\} \\ f(x; A, \mu, \sigma, \text{form} = \text{'arctan'}) &= A[\arctan(\alpha_1) + \arctan(\alpha_2)]/\pi \\ f(x; A, \mu, \sigma, \text{form} = \text{'erf'}) &= A[\text{erf}(\alpha_1) + \text{erf}(\alpha_2)]/2 \\ f(x; A, \mu, \sigma, \text{form} = \text{'logistic'}) &= A\left[1 - \frac{1}{1 + e^{\alpha_1}} - \frac{1}{1 + e^{\alpha_2}}\right] \end{aligned}$$

where $\alpha_1 = (x - \mu_1)/\sigma_1$ and $\alpha_2 = -(x - \mu_2)/\sigma_2$.

Parameters

- **`independent_vars`** (`['x']`) – Arguments to `func` that are independent variables.
- **`prefix`** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **`nan_policy`** (`str`, *optional*) – How to handle NaN and missing values in data. Must be one of: 'raise' (default), 'propagate', or 'omit'. See Notes below.
- **`missing`** (`str`, *optional*) – Synonym for 'nan_policy' for backward compatibility.
- **`**kwargs`** (*optional*) – Keyword arguments to pass to `Model`.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - 'raise' : Raise a `ValueError` (default)
 - 'propagate' : do nothing
 - 'omit' : (was 'drop') drop missing data
2. The `missing` argument is deprecated in `lmfit` 0.9.8 and will be removed in a later version. Use `nan_policy` instead, as it is consistent with the `Minimizer` class.

8.4 Exponential and Power law models

8.4.1 ExponentialModel

`class ExponentialModel` (`independent_vars=['x']`, `prefix=""`, `nan_policy='raise'`, `**kwargs`)

A model based on an exponential decay function (see https://en.wikipedia.org/wiki/Exponential_decay) with two Parameters: `amplitude` (A), and `decay` (τ), in:

$$f(x; A, \tau) = Ae^{-x/\tau}$$

Parameters

- **independent_vars** (`['x']`) – Arguments to func that are independent variables.
- **prefix** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`str`, *optional*) – How to handle NaN and missing values in data. Must be one of: 'raise' (default), 'propagate', or 'omit'. See Notes below.
- **missing** (`str`, *optional*) – Synonym for 'nan_policy' for backward compatibility.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - 'raise' : Raise a ValueError (default)
 - 'propagate' : do nothing
 - 'omit' : (was 'drop') drop missing data
2. The *missing* argument is deprecated in lmfit 0.9.8 and will be removed in a later version. Use *nan_policy* instead, as it is consistent with the Minimizer class.

8.4.2 PowerLawModel

class PowerLawModel (*independent_vars=['x'], prefix="", nan_policy='raise', **kwargs*)

A model based on a Power Law (see https://en.wikipedia.org/wiki/Power_law), with two Parameters: amplitude (*A*), and exponent (*k*), in:

$$f(x; A, k) = Ax^k$$

Parameters

- **independent_vars** (`['x']`) – Arguments to func that are independent variables.
- **prefix** (`str`, *optional*) – String to prepend to parameter names, needed to add two Models that have parameter names in common.
- **nan_policy** (`str`, *optional*) – How to handle NaN and missing values in data. Must be one of: 'raise' (default), 'propagate', or 'omit'. See Notes below.
- **missing** (`str`, *optional*) – Synonym for 'nan_policy' for backward compatibility.
- ****kwargs** (*optional*) – Keyword arguments to pass to Model.

Notes

1. `nan_policy` sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - 'raise' : Raise a ValueError (default)
 - 'propagate' : do nothing
 - 'omit' : (was 'drop') drop missing data
2. The *missing* argument is deprecated in lmfit 0.9.8 and will be removed in a later version. Use *nan_policy* instead, as it is consistent with the Minimizer class.

8.5 User-defined Models

As shown in the previous chapter (*Modeling Data and Curve Fitting*), it is fairly straightforward to build fitting models from parametrized Python functions. The number of model classes listed so far in the present chapter should make it clear that this process is not too difficult. Still, it is sometimes desirable to build models from a user-supplied function. This may be especially true if model-building is built-in to some larger library or application for fitting in which the user may not be able to easily build and use a new model from Python code.

The *ExpressionModel* allows a model to be built from a user-supplied expression. This uses the *asteval* module also used for mathematical constraints as discussed in *Using Mathematical Constraints*.

8.5.1 ExpressionModel

class ExpressionModel (*expr, independent_vars=None, init_script=None, nan_policy='raise', **kws*)
Generate a model from user-supplied expression.

Parameters

- **expr** (*str*) – Mathematical expression for model.
- **independent_vars** (*list of str or None, optional*) – Variable names to use as independent variables.
- **init_script** (*str or None, optional*) – Initial script to run in *asteval* interpreter.
- **nan_policy** (*str, optional*) – How to handle NaN and missing values in data. Must be one of: 'raise' (default), 'propagate', or 'omit'. See Notes below.
- **missing** (*str, optional*) – Synonym for 'nan_policy' for backward compatibility.
- ****kws** (*optional*) – Keyword arguments to pass to *Model*.

Notes

1. each instance of *ExpressionModel* will create and using its own version of an *asteval* interpreter.
2. prefix is **not supported** for *ExpressionModel*.
3. *nan_policy* sets what to do when a NaN or missing value is seen in the data. Should be one of:
 - 'raise' : Raise a *ValueError* (default)
 - 'propagate' : do nothing
 - 'omit' : (was 'drop') drop missing data
4. The *missing* argument is deprecated in *lmfit* 0.9.8 and will be removed in a later version. Use *nan_policy* instead, as it is consistent with the *Minimizer* class.

Since the point of this model is that an arbitrary expression will be supplied, the determination of what are the parameter names for the model happens when the model is created. To do this, the expression is parsed, and all symbol names are found. Names that are already known (there are over 500 function and value names in the *asteval* namespace, including most Python built-ins, more than 200 functions inherited from NumPy, and more than 20 common *lineshapes* defined in the *lineshapes* module) are not converted to parameters. Unrecognized names are expected to be names of either parameters or independent variables. If *independent_vars* is the default value of *None*, and if the expression contains a variable named *x*, that will be used as the independent variable. Otherwise, *independent_vars* must be given.

For example, if one creates an *ExpressionModel* as:

```
>>> mod = ExpressionModel('off + amp * exp(-x/x0) * sin(x*phase)')
```

The name *exp* will be recognized as the exponent function, so the model will be interpreted to have parameters named *off*, *amp*, *x0* and *phase*. In addition, *x* will be assumed to be the sole independent variable. In general, there is no obvious way to set default parameter values or parameter hints for bounds, so this will have to be handled explicitly.

To evaluate this model, you might do the following:

```
>>> x = numpy.linspace(0, 10, 501)
>>> params = mod.make_params(off=0.25, amp=1.0, x0=2.0, phase=0.04)
>>> y = mod.eval(params, x=x)
```

While many custom models can be built with a single line expression (especially since the names of the lineshapes like *gaussian*, *lorentzian* and so on, as well as many NumPy functions, are available), more complex models will inevitably require multiple line functions. You can include such Python code with the *init_script* argument. The text of this script is evaluated when the model is initialized (and before the actual expression is parsed), so that you can define functions to be used in your expression.

As a probably unphysical example, to make a model that is the derivative of a Gaussian function times the logarithm of a Lorentzian function you may could to define this in a script:

```
>>> script = """
def mycurve(x, amp, cen, sig):
    loren = lorentzian(x, amplitude=amp, center=cen, sigma=sig)
    gauss = gaussian(x, amplitude=amp, center=cen, sigma=sig)
    return log(loren) * gradient(gauss) / gradient(x)
"""
```

and then use this with *ExpressionModel* as:

```
>>> mod = ExpressionModel('mycurve(x, height, mid, wid)',
                           init_script=script,
                           independent_vars=['x'])
```

As above, this will interpret the parameter names to be *height*, *mid*, and *wid*, and build a model that can be used to fit data.

8.6 Example 1: Fit Peak data to Gaussian, Lorentzian, and Voigt profiles

Here, we will fit data to three similar line shapes, in order to decide which might be the better model. We will start with a Gaussian profile, as in the previous chapter, but use the built-in *GaussianModel* instead of writing one ourselves. This is a slightly different version from the one in previous example in that the parameter names are different, and have built-in default values. We will simply use:

```
from numpy import loadtxt

from lmfit.models import GaussianModel

data = loadtxt('test_peak.dat')
x = data[:, 0]
y = data[:, 1]
```

(continues on next page)

(continued from previous page)



```
mod = GaussianModel()

pars = mod.guess(y, x=x)
out = mod.fit(y, pars, x=x)
print(out.fit_report(min_correl=0.25))
```

which prints out the results:

```
[[Model]]
  Model(gaussian)
[[Fit Statistics]]
  # fitting method   = leastsq
  # function evals   = 27
  # data points      = 401
  # variables        = 3
  chi-square         = 29.9943157
  reduced chi-square = 0.07536260
  Akaike info crit   = -1033.77437
  Bayesian info crit = -1021.79248
[[Variables]]
  sigma:      1.23218359 +/- 0.00737496 (0.60%) (init = 1.35)
  amplitude:  30.3135620 +/- 0.15712686 (0.52%) (init = 43.62238)
  center:     9.24277047 +/- 0.00737496 (0.08%) (init = 9.25)
  fwhm:       2.90157056 +/- 0.01736670 (0.60%) == '2.3548200*sigma'
  height:     9.81457817 +/- 0.05087283 (0.52%) == '0.3989423*amplitude/max(1.e-15,
↳sigma)'
[[Correlations]] (unreported correlations are < 0.250)
  C(sigma, amplitude) = 0.577
```

We see a few interesting differences from the results of the previous chapter. First, the parameter names are longer. Second, there are `fwhm` and `height` parameters, to give the full width at half maximum and maximum peak height. And third, the automated initial guesses are pretty good. A plot of the fit:



Fit

to peak with Gaussian (left) and Lorentzian (right) models.

shows a decent match to the data – the fit worked with no explicit setting of initial parameter values. Looking more closely, the fit is not perfect, especially in the tails of the peak, suggesting that a different peak shape, with longer tails, should be used. Perhaps a Lorentzian would be better? To do this, we simply replace `GaussianModel` with `LorentzianModel` to get a `LorentzianModel`:

```
from lmfit.models import LorentzianModel
mod = LorentzianModel()
```

with the rest of the script as above. Perhaps predictably, the first thing we try gives results that are worse:

```

[[Model]]
    Model(lorentzian)
[[Fit Statistics]]
    # fitting method      = leastsq
    # function evals      = 23
    # data points         = 401
    # variables           = 3
    chi-square            = 53.7535387
    reduced chi-square    = 0.13505914
    Akaike info crit     = -799.830322
    Bayesian info crit   = -787.848438
[[Variables]]
    sigma:      1.15483925 +/- 0.01315659 (1.14%) (init = 1.35)
    center:     9.24438944 +/- 0.00927619 (0.10%) (init = 9.25)
    amplitude:  38.9727645 +/- 0.31386183 (0.81%) (init = 54.52798)
    fwhm:       2.30967850 +/- 0.02631318 (1.14%) == '2.0000000*sigma'
    height:     10.7421156 +/- 0.08633945 (0.80%) == '0.3183099*amplitude/max(1.e-15,
↪sigma)'
[[Correlations]] (unreported correlations are < 0.250)
    C(sigma, amplitude) = 0.709

```

with the plot shown on the right in the figure above. The tails are now too big, and the value for χ^2 almost doubled. A Voigt model does a better job. Using `VoigtModel`, this is as simple as using:

```

from lmfit.models import VoigtModel
mod = VoigtModel()

```



with all the rest of the script as above. This gives:

```

[[Model]]
    Model(voigt)
[[Fit Statistics]]
    # fitting method      = leastsq
    # function evals      = 23
    # data points         = 401
    # variables           = 3
    chi-square            = 14.5448627
    reduced chi-square    = 0.03654488
    Akaike info crit     = -1324.00615
    Bayesian info crit   = -1312.02427
[[Variables]]
    center:      9.24411150 +/- 0.00505482 (0.05%) (init = 9.25)
    sigma:       0.73015627 +/- 0.00368460 (0.50%) (init = 0.8775)
    amplitude:   35.7554146 +/- 0.13861321 (0.39%) (init = 65.43358)
    gamma:       0.73015627 +/- 0.00368460 (0.50%) == 'sigma'
    fwhm:        2.62951907 +/- 0.01326940 (0.50%) == '3.6013100*sigma'
    height:      10.2203969 +/- 0.03009415 (0.29%) == 'amplitude*wofz((1j*gamma)/
↪(sigma*sqrt(2))).real/(sigma*sqrt(2*pi))'
[[Correlations]] (unreported correlations are < 0.250)
    C(sigma, amplitude) = 0.651

```

which has a much better value for χ^2 and the other goodness-of-fit measures, and an obviously better match to the data as seen in the figure below (left).



Fit

to peak with Voigt model (left) and Voigt model with γ varying independently of σ (right).

Can we do better? The Voigt function has a γ parameter (gamma) that can be distinct from σ . The default behavior used above constrains gamma to have exactly the same value as σ . If we allow these to vary separately, does the fit improve? To do this, we have to change the gamma parameter from a constrained expression and give it a starting value using something like:

```
mod = VoigtModel()
pars = mod.guess(y, x=x)
pars['gamma'].set(value=0.7, vary=True, expr='')
```

which gives:

```
[[Model]]
  Model(voigt)
[[Fit Statistics]]
  # fitting method  = leastsq
  # function evals  = 23
  # data points     = 401
  # variables       = 4
  chi-square        = 10.9301767
  reduced chi-square = 0.02753193
  Akaike info crit  = -1436.57602
  Bayesian info crit = -1420.60017
[[Variables]]
  sigma: 0.89518909 +/- 0.01415450 (1.58%) (init = 0.8775)
  amplitude: 34.1914737 +/- 0.17946860 (0.52%) (init = 65.43358)
  center: 9.24374847 +/- 0.00441903 (0.05%) (init = 9.25)
  gamma: 0.52540199 +/- 0.01857955 (3.54%) (init = 0.7)
  fwhm: 3.22385342 +/- 0.05097475 (1.58%) == '3.6013100*sigma'
  height: 10.0872204 +/- 0.03482129 (0.35%) == 'amplitude*wofz((1j*gamma)/
  ↳(sigma*sqrt(2))).real/(sigma*sqrt(2*pi))'
[[Correlations]] (unreported correlations are < 0.250)
  C(sigma, gamma) = -0.928
  C(amplitude, gamma) = 0.821
  C(sigma, amplitude) = -0.651
```

and the fit shown on the right above.

Comparing the two fits with the Voigt function, we see that χ^2 is definitely improved with a separately varying gamma parameter. In addition, the two values for gamma and sigma differ significantly – well outside the estimated uncertainties. More compelling, reduced χ^2 is improved even though a fourth variable has been added to the fit. In the simplest statistical sense, this suggests that gamma is a significant variable in the model. In addition, we can use both the Akaike or Bayesian Information Criteria (see *Akaike and Bayesian Information Criteria*) to assess how likely the model with variable gamma is to explain the data than the model with gamma fixed to the value of sigma. According

to theory, $\exp(-(AIC1 - AIC0)/2)$ gives the probability that a model with AIC1 is more likely than a model with AIC0. For the two models here, with AIC values of -1436 and -1324 (Note: if we had more carefully set the value for `weights` based on the noise in the data, these values might be positive, but their difference would be roughly the same), this says that the model with `gamma` fixed to `sigma` has a probability less than $5.e-25$ of being the better model.

8.7 Example 2: Fit data to a Composite Model with pre-defined models

Here, we repeat the point made at the end of the last chapter that instances of `Model` class can be added together to make a *composite model*. By using the large number of built-in models available, it is therefore very simple to build models that contain multiple peaks and various backgrounds. An example of a simple fit to a noisy step function plus a constant:

```
#!/usr/bin/env python

# <examples/doc_builtinmodels_stepmodel.py>
import matplotlib.pyplot as plt
import numpy as np

from lmfit.models import LinearModel, StepModel

x = np.linspace(0, 10, 201)
y = np.ones_like(x)
y[:48] = 0.0
y[48:77] = np.arange(77-48)/(77.0-48)
np.random.seed(0)
y = 110.2 * (y + 9e-3*np.random.randn(len(x))) + 12.0 + 2.22*x

step_mod = StepModel(form='erf', prefix='step_')
line_mod = LinearModel(prefix='line_')

pars = line_mod.make_params(intercept=y.min(), slope=0)
pars += step_mod.guess(y, x=x, center=2.5)

mod = step_mod + line_mod
out = mod.fit(y, pars, x=x)

print(out.fit_report())

plt.plot(x, y, 'b')
plt.plot(x, out.init_fit, 'k--')
plt.plot(x, out.best_fit, 'r-')
plt.show()
# <end examples/doc_builtinmodels_stepmodel.py>
```

After constructing step-like data, we first create a `StepModel` telling it to use the `erf` form (see details above), and a `ConstantModel`. We set initial values, in one case using the data and `guess()` method for the initial step function parameters, and `make_params()` arguments for the linear component. After making a composite model, we run `fit()` and report the results, which gives:

```
[[Model]]
  (Model(step, prefix='step_', form='erf') + Model(linear, prefix='line_'))
[[Fit Statistics]]
```


(continues on next page)

(continued from previous page)

```

# fitting method = leastsq
# function evals = 49
# data points   = 201
# variables     = 5
chi-square      = 593.709622
reduced chi-square = 3.02913072
Akaike info crit = 227.700173
Bayesian info crit = 244.216698
[[Variables]]
line_intercept: 12.0964833 +/- 0.27606235 (2.28%) (init = 11.58574)
line_slope:    1.87164655 +/- 0.09318714 (4.98%) (init = 0)
step_sigma:    0.67392841 +/- 0.01091168 (1.62%) (init = 1.428571)
step_center:   3.13494792 +/- 0.00516615 (0.16%) (init = 2.5)
step_amplitude: 112.858376 +/- 0.65392949 (0.58%) (init = 134.7378)
[[Correlations]] (unreported correlations are < 0.100)
C(line_slope, step_amplitude) = -0.879
C(step_sigma, step_amplitude) = 0.564
C(line_slope, step_sigma)     = -0.457
C(line_intercept, step_center) = 0.427
C(line_intercept, line_slope)  = -0.309
C(line_slope, step_center)     = -0.234
C(line_intercept, step_sigma)  = -0.137
C(line_intercept, step_amplitude) = -0.117
C(step_center, step_amplitude) = 0.109

```



with a plot of

8.8 Example 3: Fitting Multiple Peaks – and using Prefixes

As shown above, many of the models have similar parameter names. For composite models, this could lead to a problem of having parameters for different parts of the model having the same name. To overcome this, each `Model` can have a `prefix` attribute (normally set to a blank string) that will be put at the beginning of each parameter name. To illustrate, we fit one of the classic datasets from the [NIST StRD](#) suite involving a decaying exponential and two gaussians.

```

#!/usr/bin/env python

# <examples/doc_builtinmodels_nistgauss.py>
import matplotlib.pyplot as plt

```

(continues on next page)

(continued from previous page)

```

import numpy as np

from lmfit.models import ExponentialModel, GaussianModel

dat = np.loadtxt('NIST_Gauss2.dat')
x = dat[:, 1]
y = dat[:, 0]

exp_mod = ExponentialModel(prefix='exp_')
pars = exp_mod.guess(y, x=x)

gauss1 = GaussianModel(prefix='g1_')
pars.update(gauss1.make_params())

pars['g1_center'].set(105, min=75, max=125)
pars['g1_sigma'].set(15, min=3)
pars['g1_amplitude'].set(2000, min=10)

gauss2 = GaussianModel(prefix='g2_')
pars.update(gauss2.make_params())

pars['g2_center'].set(155, min=125, max=175)
pars['g2_sigma'].set(15, min=3)
pars['g2_amplitude'].set(2000, min=10)

mod = gauss1 + gauss2 + exp_mod

init = mod.eval(pars, x=x)

out = mod.fit(y, pars, x=x)
print(out.fit_report(min_correl=0.5))

plot_components = False

plt.plot(x, y, 'b')
plt.plot(x, init, 'k--')
plt.plot(x, out.best_fit, 'r-')

if plot_components:
    comps = out.eval_components(x=x)
    plt.plot(x, comps['g1_'], 'b--')
    plt.plot(x, comps['g2_'], 'b--')
    plt.plot(x, comps['exp_'], 'k--')

plt.show()
# <end examples/doc_builtinmodels_nistgauss.py>

```

where we give a separate prefix to each model (they all have an amplitude parameter). The prefix values are attached transparently to the models.

Note that the calls to `make_param()` used the bare name, without the prefix. We could have used the prefixes, but because we used the individual model `gauss1` and `gauss2`, there was no need.

Note also in the example here that we explicitly set bounds on many of the parameter values.


The fit results printed out are:


```

[[Model]]
  ((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) +
  ↪Model(exponential, prefix='exp_'))
[[Fit Statistics]]
  # fitting method      = leastsq
  # function evals      = 48
  # data points         = 250
  # variables           = 8
  chi-square            = 1247.52821
  reduced chi-square    = 5.15507524
  Akaike info crit     = 417.864631
  Bayesian info crit   = 446.036318
[[Variables]]
  exp_decay:           90.9508860 +/- 1.10310509 (1.21%) (init = 93.24905)
  exp_amplitude:       99.0183283 +/- 0.53748735 (0.54%) (init = 162.2102)
  g1_center:           107.030954 +/- 0.15006786 (0.14%) (init = 105)
  g1_sigma:            16.6725753 +/- 0.16048161 (0.96%) (init = 15)
  g1_amplitude:        4257.77319 +/- 42.3833645 (1.00%) (init = 2000)
  g1_fwhm:             39.2609138 +/- 0.37790530 (0.96%) == '2.3548200*g1_sigma'
  g1_height:           101.880231 +/- 0.59217100 (0.58%) == '0.3989423*g1_amplitude/
  ↪max(1.e-15, g1_sigma)'
  g2_center:           153.270101 +/- 0.19466743 (0.13%) (init = 155)
  g2_sigma:            13.8069484 +/- 0.18679415 (1.35%) (init = 15)
  g2_amplitude:        2493.41771 +/- 36.1694731 (1.45%) (init = 2000)
  g2_fwhm:             32.5128783 +/- 0.43986659 (1.35%) == '2.3548200*g2_sigma'
  g2_height:           72.0455934 +/- 0.61722094 (0.86%) == '0.3989423*g2_amplitude/
  ↪max(1.e-15, g2_sigma)'
[[Correlations]] (unreported correlations are < 0.500)
  C(g1_sigma, g1_amplitude) = 0.824
  C(g2_sigma, g2_amplitude) = 0.815
  C(exp_decay, exp_amplitude) = -0.695
  C(g1_sigma, g2_center) = 0.684
  C(g1_center, g2_amplitude) = -0.669
  C(g1_center, g2_sigma) = -0.652
  C(g1_amplitude, g2_center) = 0.648
  C(g1_center, g2_center) = 0.621
  C(g1_center, g1_sigma) = 0.507
  C(exp_decay, g1_amplitude) = -0.507

```

We get a very good fit to this problem (described at the NIST site as of average difficulty, but the tests there are generally deliberately challenging) by applying reasonable initial guesses and putting modest but explicit bounds on the parameter values. This fit is shown on the left:



One final point on setting initial values. From looking at the data itself, we can see the two Gaussian peaks are

reasonably well separated but do overlap. Furthermore, we can tell that the initial guess for the decaying exponential component was poorly estimated because we used the full data range. We can simplify the initial parameter values by using this, and by defining an `index_of()` function to limit the data range. That is, with:

```
def index_of(arrval, value):
    """Return index of array *at or below* value."""
    if value < min(arrval):
        return 0
    return max(np.where(arrval <= value)[0])

ix1 = index_of(x, 75)
ix2 = index_of(x, 135)
ix3 = index_of(x, 175)

exp_mod.guess(y[:ix1], x=x[:ix1])
gauss1.guess(y[ix1:ix2], x=x[ix1:ix2])
gauss2.guess(y[ix2:ix3], x=x[ix2:ix3])
```

we can get a better initial estimate. The fit converges to the same answer, giving to identical values (to the precision printed out in the report), but in few steps, and without any bounds on parameters at all:

```
[[Model]]
  ((Model(gaussian, prefix='g1_') + Model(gaussian, prefix='g2_')) +
  ↪Model(exponential, prefix='exp_'))
[[Fit Statistics]]
  # fitting method   = leastsq
  # function evals   = 39
  # data points      = 250
  # variables        = 8
  chi-square         = 1247.52821
  reduced chi-square = 5.15507524
  Akaike info crit   = 417.864631
  Bayesian info crit = 446.036318
[[Variables]]
  exp_decay:      90.9508890 +/- 1.10310483 (1.21%) (init = 111.1985)
  exp_amplitude:  99.0183270 +/- 0.53748905 (0.54%) (init = 94.53724)
  g1_sigma:       16.6725765 +/- 0.16048227 (0.96%) (init = 14.5)
  g1_amplitude:   4257.77343 +/- 42.3836432 (1.00%) (init = 3189.648)
  g1_center:      107.030956 +/- 0.15006873 (0.14%) (init = 106.5)
  g1_fwhm:        39.2609166 +/- 0.37790686 (0.96%) == '2.3548200*g1_sigma'
  g1_height:      101.880230 +/- 0.59217233 (0.58%) == '0.3989423*g1_amplitude/
  ↪max(1.e-15, g1_sigma)'
  g2_sigma:       13.8069461 +/- 0.18679534 (1.35%) (init = 15)
  g2_amplitude:   2493.41733 +/- 36.1696911 (1.45%) (init = 2818.337)
  g2_center:      153.270101 +/- 0.19466905 (0.13%) (init = 150)
  g2_fwhm:        32.5128728 +/- 0.43986940 (1.35%) == '2.3548200*g2_sigma'
  g2_height:      72.0455948 +/- 0.61722329 (0.86%) == '0.3989423*g2_amplitude/
  ↪max(1.e-15, g2_sigma)'
[[Correlations]] (unreported correlations are < 0.500)
  C(g1_sigma, g1_amplitude) = 0.824
  C(g2_sigma, g2_amplitude) = 0.815
  C(exp_decay, exp_amplitude) = -0.695
  C(g1_sigma, g2_center) = 0.684
  C(g1_center, g2_amplitude) = -0.669
  C(g1_center, g2_sigma) = -0.652
  C(g1_amplitude, g2_center) = 0.648
  C(g1_center, g2_center) = 0.621
  C(g1_sigma, g1_center) = 0.507
```

(continues on next page)

(continued from previous page)

```
C(exp_decay, g1_amplitude) = -0.507
```

This script is in the file `doc_builtinmodels_nistgauss2.py` in the examples folder, and the fit result shown on the right above shows an improved initial estimate of the data.

CALCULATION OF CONFIDENCE INTERVALS

The `lmfit confidence` module allows you to explicitly calculate confidence intervals for variable parameters. For most models, it is not necessary since the estimation of the standard error from the estimated covariance matrix is normally quite good.

But for some models, the sum of two exponentials for example, the approximation begins to fail. For this case, `lmfit` has the function `conf_interval()` to calculate confidence intervals directly. This is substantially slower than using the errors estimated from the covariance matrix, but the results are more robust.

9.1 Method used for calculating confidence intervals

The F-test is used to compare our null model, which is the best fit we have found, with an alternate model, where one of the parameters is fixed to a specific value. The value is changed until the difference between χ_0^2 and χ_f^2 can't be explained by the loss of a degree of freedom within a certain confidence.

$$F(P_{fix}, N - P) = \left(\frac{\chi_f^2}{\chi_0^2} - 1 \right) \frac{N - P}{P_{fix}}$$

N is the number of data points and P the number of parameters of the null model. P_{fix} is the number of fixed parameters (or to be more clear, the difference of number of parameters between our null model and the alternate model).

Adding a log-likelihood method is under consideration.

9.2 A basic example

First we create an example problem:

```
>>> import lmfit
>>> import numpy as np
>>> x = np.linspace(0.3, 10, 100)
>>> np.random.seed(0)
>>> y = 1/(0.1*x) + 2 + 0.1*np.random.randn(x.size)
>>> pars = lmfit.Parameters()
>>> pars.add_many(('a', 0.1), ('b', 1))
>>> def residual(p):
...     return 1/(p['a']*x) + p['b'] - y
```

before we can generate the confidence intervals, we have to run a fit, so that the automated estimate of the standard errors can be used as a starting point:

```
>>> mini = lmfit.Minimizer(residual, pars)
>>> result = mini.minimize()
>>> print(lmfit.fit_report(result.params))
[[Variables]]
  a:  0.09943896 +/- 1.9322e-04 (0.19%) (init = 0.1)
  b:  1.98476945 +/- 0.01222678 (0.62%) (init = 1)
[[Correlations]] (unreported correlations are < 0.100)
  C(a, b) =  0.601
```

Now it is just a simple function call to calculate the confidence intervals:

```
>>> ci = lmfit.conf_interval(mini, result)
>>> lmfit.printfuncs.report_ci(ci)
  99.73%   95.45%   68.27%   _BEST_   68.27%   95.45%   99.73%
a: -0.00059 -0.00039 -0.00019  0.09944 +0.00019 +0.00039 +0.00060
b: -0.03766 -0.02478 -0.01230  1.98477 +0.01230 +0.02478 +0.03761
```

This shows the best-fit values for the parameters in the *_BEST_* column, and parameter values that are at the varying confidence levels given by steps in σ . As we can see, the estimated error is almost the same, and the uncertainties are well behaved: Going from 1 σ (68% confidence) to 3 σ (99.7% confidence) uncertainties is fairly linear. It can also be seen that the errors are fairly symmetric around the best fit value. For this problem, it is not necessary to calculate confidence intervals, and the estimates of the uncertainties from the covariance matrix are sufficient.

9.3 An advanced example

Now we look at a problem where calculating the error from approximated covariance can lead to misleading result – two decaying exponentials. In fact such a problem is particularly hard for the Levenberg-Marquardt method, so we first estimate the results using the slower but robust Nelder-Mead method, and *then* use Levenberg-Marquardt to estimate the uncertainties and correlations.

```
#!/usr/bin/env python

# <examples/doc_confidence_advanced.py>
import numpy as np
import matplotlib.pyplot as plt

import lmfit

x = np.linspace(1, 10, 250)
np.random.seed(0)
y = 3.0*np.exp(-x/2) - 5.0*np.exp(-(x-0.1)/10.) + 0.1*np.random.randn(len(x))

p = lmfit.Parameters()
p.add_many(('a1', 4.), ('a2', 4.), ('t1', 3.), ('t2', 3.))

def residual(p):
    return p['a1']*np.exp(-x/p['t1']) + p['a2']*np.exp(-(x-0.1)/p['t2']) - y

# create Minimizer
mini = lmfit.Minimizer(residual, p, nan_policy='omit')

# first solve with Nelder-Mead
```

(continues on next page)

(continued from previous page)

```

out1 = mini.minimize(method='Nelder')

# then solve with Levenberg-Marquardt using the
# Nelder-Mead solution as a starting point
out2 = mini.minimize(method='leastsq', params=out1.params)

lmfit.report_fit(out2.params, min_correl=0.5)

ci, trace = lmfit.conf_interval(mini, out2, sigmas=[1, 2],
                               trace=True, verbose=False)
lmfit.printfuncs.report_ci(ci)

plot_type = 2

if plot_type == 0:
    plt.plot(x, y)
    plt.plot(x, residual(out2.params) + y)

elif plot_type == 1:
    cx, cy, grid = lmfit.conf_interval2d(mini, out2, 'a2', 't2', 30, 30)
    plt.contourf(cx, cy, grid, np.linspace(0, 1, 11))
    plt.xlabel('a2')
    plt.colorbar()
    plt.ylabel('t2')

elif plot_type == 2:
    cx, cy, grid = lmfit.conf_interval2d(mini, out2, 'a1', 't2', 30, 30)
    plt.contourf(cx, cy, grid, np.linspace(0, 1, 11))
    plt.xlabel('a1')
    plt.colorbar()
    plt.ylabel('t2')

elif plot_type == 3:
    cx1, cy1, prob = trace['a1']['a1'], trace['a1']['t2'], trace['a1']['prob']
    cx2, cy2, prob2 = trace['t2']['t2'], trace['t2']['a1'], trace['t2']['prob']
    plt.scatter(cx1, cy1, c=prob, s=30)
    plt.scatter(cx2, cy2, c=prob2, s=30)
    plt.gca().set_xlim((2.5, 3.5))
    plt.gca().set_ylim((11, 13))
    plt.xlabel('a1')
    plt.ylabel('t2')

if plot_type > 0:
    plt.show()
# <end examples/doc_confidence_advanced.py>

```

which will report:

```

[[Variables]]
  a1:  2.98622120 +/- 0.14867187 (4.98%) (init = 2.986237)
  a2: -4.33526327 +/- 0.11527506 (2.66%) (init = -4.335256)
  t1:  1.30994233 +/- 0.13121177 (10.02%) (init = 1.309932)
  t2: 11.8240351 +/- 0.46316470 (3.92%) (init = 11.82408)
[[Correlations]] (unreported correlations are < 0.500)
  C(a2, t2) = 0.987
  C(a2, t1) = -0.925
  C(t1, t2) = -0.881

```


(continues on next page)

(continued from previous page)

| | | | | | |
|-----|--------------------|----------|----------|----------|----------|
| | C(a1, t1) = -0.599 | | | | |
| | 95.45% | 68.27% | _BEST_ | 68.27% | 95.45% |
| a1: | -0.27286 | -0.14165 | 2.98622 | +0.16353 | +0.36343 |
| a2: | -0.30444 | -0.13219 | -4.33526 | +0.10688 | +0.19683 |
| t1: | -0.23392 | -0.12494 | 1.30994 | +0.14660 | +0.32369 |
| t2: | -1.01943 | -0.48820 | 11.82404 | +0.46041 | +0.90441 |

Again we called `conf_interval()`, this time with tracing and only for 1- and 2- σ . Comparing these two different estimates, we see that the estimate for $a1$ is reasonably well approximated from the covariance matrix, but the estimates for $a2$ and especially for $t1$, and $t2$ are very asymmetric and that going from 1 σ (68% confidence) to 2 σ (95% confidence) is not very predictable.


Plots of the confidence region are shown in the figures below for $a1$ and $t2$ (left), and $a2$ and $t2$ (right):



Neither of these plots is very much like an ellipse, which is implicitly assumed by the approach using the covariance matrix.

The trace returned as the optional second argument from `conf_interval()` contains a dictionary for each variable parameter. The values are dictionaries with arrays of values for each variable, and an array of corresponding probabilities for the corresponding cumulative variables. This can be used to show the dependence between two parameters:


```
>>> x, y, prob = trace['a1']['a1'], trace['a1']['t2'], trace['a1']['prob']
>>> x2, y2, prob2 = trace['t2']['t2'], trace['t2']['a1'], trace['t2']['prob']
>>> plt.scatter(x, y, c=prob, s=30)
>>> plt.scatter(x2, y2, c=prob2, s=30)
>>> plt.gca().set_xlim((2.5, 3.5))
>>> plt.gca().set_ylim((11.5, 12.5))
>>> plt.xlabel('a1')
>>> plt.ylabel('t2')
>>> plt.show()
```

which shows the trace of values:

`emcee()` method uses Markov Chain Monte Carlo to sample the posterior probability distribution. These distributions demonstrate the range of solutions that the data supports. The following image was obtained by using `Minimizer.emcee()` on the same problem.

The `Minimizer`.



Credible intervals (the Bayesian equivalent of the frequentist confidence interval) can be obtained with this method. MCMC can be used for model selection, to determine outliers, to marginalise over nuisance parameters, etcetera. For example, you may have fractionally underestimated the uncertainties on a dataset. MCMC can be used to estimate the true level of uncertainty on each datapoint. A tutorial on the possibilities offered by MCMC can be found at¹.

9.4 Confidence Interval Functions

conf_interval (*minimizer, result, p_names=None, sigmas=(1, 2, 3), trace=False, maxiter=200, verbose=False, prob_func=None*)

Calculate the confidence interval for parameters.

The parameter for which the ci is calculated will be varied, while the remaining parameters are re-optimized

¹ <https://jakevdp.github.io/blog/2014/03/11/frequentism-and-bayesianism-a-practical-intro/>

to minimize chi-square. The resulting chi-square is used to calculate the probability with a given statistic (e.g., F-test). This function uses a 1d-rootfinder from SciPy to find the values resulting in the searched confidence region.

Parameters

- **minimizer** (*Minimizer*) – The minimizer to use, holding objective function.
- **result** (*MinimizerResult*) – The result of running `minimize()`.
- **p_names** (*list, optional*) – Names of the parameters for which the ci is calculated. If None, the ci is calculated for every parameter.
- **sigmas** (*list, optional*) – The sigma-levels to find. Default is [1, 2, 3]. See Note below.
- **trace** (*bool, optional*) – Defaults to False, if True, each result of a probability calculation is saved along with the parameter. This can be used to plot so-called “profile traces”.
- **maxiter** (*int, optional*) – Maximum of iteration to find an upper limit. Default is 200.
- **verbose** (*bool, optional*) – Print extra debugging information. Default is False.
- **prob_func** (*None or callable, optional*) – Function to calculate the probability from the optimized chi-square. Default is None and uses built-in `f_compare` (F-test).

Returns

- **output** (*dict*) – A dictionary that contains a list of (sigma, vals)-tuples for each name.
- **trace_dict** (*dict, optional*) – Only if trace is True. Is a dict, the key is the parameter which was fixed. The values are again a dict with the names as keys, but with an additional key ‘prob’. Each contains an array of the corresponding values.

Note: The values for *sigma* are taken as the number of standard deviations for a normal distribution and converted to probabilities. That is, the default `sigma=(1, 2, 3)` will use probabilities of 0.6827, 0.9545, and 0.9973. If any of the sigma values is less than 1, that will be interpreted as a probability. That is, a value of 1 and 0.6827 will give the same results, within precision.

See also:

`conf_interval2d()`

Examples

```
>>> from lmfit.printfuncs import *
>>> mini = minimize(some_func, params)
>>> mini.leastsq()
True
>>> report_errors(params)
... #report
>>> ci = conf_interval(mini)
>>> report_ci(ci)
... #report
```

Now with quantiles for the sigmas and using the trace.

```
>>> ci, trace = conf_interval(mini, sigmas=(0.5, 1, 2, 3), trace=True)
>>> fixed = trace['para1']['para1']
>>> free = trace['para1']['not_para1']
>>> prob = trace['para1']['prob']
```

This makes it possible to plot the dependence between free and fixed parameters.

conf_interval2d (*minimizer, result, x_name, y_name, nx=10, ny=10, limits=None, prob_func=None*)
Calculate confidence regions for two fixed parameters.

The method itself is explained in *conf_interval*: here we are fixing two parameters.

Parameters

- **minimizer** (*Minimizer*) – The minimizer to use, holding objective function.
- **result** (*MinimizerResult*) – The result of running *minimize()*.
- **x_name** (*str*) – The name of the parameter which will be the x direction.
- **y_name** (*str*) – The name of the parameter which will be the y direction.
- **nx** (*int, optional*) – Number of points in the x direction.
- **ny** (*int, optional*) – Number of points in the y direction.
- **limits** (*tuple, optional*) – Should have the form ((x_upper, x_lower), (y_upper, y_lower)). If not given, the default is 5 std-errs in each direction.
- **prob_func** (*None or callable, optional*) – Function to calculate the probability from the optimized chi-square. Default is *None* and uses built-in *f_compare* (F-test).

Returns

- **x** (*numpy.ndarray*) – X-coordinates (same shape as *nx*).
- **y** (*numpy.ndarray*) – Y-coordinates (same shape as *ny*).
- **grid** (*numpy.ndarray*) – Grid containing the calculated probabilities (with shape (*nx*, *ny*)).

Examples

```
>>> mini = Minimizer(some_func, params)
>>> result = mini.leastsq()
>>> x, y, gr = conf_interval2d(mini, result, 'para1', 'para2')
>>> plt.contour(x, y, gr)
```

ci_report (*ci, with_offset=True, ndigits=5*)
Return text of a report for confidence intervals.

Parameters

- **with_offset** (*bool, optional*) – Whether to subtract best value from all other values (default is *True*).
- **ndigits** (*int, optional*) – Number of significant digits to show (default is 5).

Returns Text of formatted report on confidence intervals.

Return type *str*

BOUNDS IMPLEMENTATION

This section describes the implementation of `Parameter` bounds. The `MINPACK-1` implementation used in `scipy.optimize.leastsq` for the Levenberg-Marquardt algorithm does not explicitly support bounds on parameters, and expects to be able to fully explore the available range of values for any `Parameter`. Simply placing hard constraints (that is, resetting the value when it exceeds the desired bounds) prevents the algorithm from determining the partial derivatives, and leads to unstable results.

Instead of placing such hard constraints, bounded parameters are mathematically transformed using the formulation devised (and documented) for `MINUIT`. This is implemented following (and borrowing heavily from) the `leastsqbound` from J. J. Helmus. Parameter values are mapped from internally used, freely variable values P_{internal} to bounded parameters P_{bounded} . When both `min` and `max` bounds are specified, the mapping is:

$$\begin{aligned} P_{\text{internal}} &= \arcsin\left(\frac{2(P_{\text{bounded}} - \text{min})}{(\text{max} - \text{min})} - 1\right) \\ P_{\text{bounded}} &= \text{min} + (\sin(P_{\text{internal}}) + 1) \frac{(\text{max} - \text{min})}{2} \end{aligned}$$

With only an upper limit `max` supplied, but `min` left unbounded, the mapping is:

$$\begin{aligned} P_{\text{internal}} &= \sqrt{(\text{max} - P_{\text{bounded}} + 1)^2 - 1} \\ P_{\text{bounded}} &= \text{max} + 1 - \sqrt{P_{\text{internal}}^2 + 1} \end{aligned}$$

With only a lower limit `min` supplied, but `max` left unbounded, the mapping is:

$$\begin{aligned} P_{\text{internal}} &= \sqrt{(P_{\text{bounded}} - \text{min} + 1)^2 - 1} \\ P_{\text{bounded}} &= \text{min} - 1 + \sqrt{P_{\text{internal}}^2 + 1} \end{aligned}$$

With these mappings, the value for the bounded `Parameter` cannot exceed the specified bounds, though the internally varied value can be freely varied.

It bears repeating that code from `leastsqbound` was adopted to implement the transformation described above. The challenging part (thanks again to Jonathan J. Helmus!) here is to re-transform the covariance matrix so that the uncertainties can be estimated for bounded `Parameters`. This is included by using the derivate $dP_{\text{internal}}/dP_{\text{bounded}}$ from the equations above to re-scale the Jacobin matrix before constructing the covariance matrix from it. Tests show that this re-scaling of the covariance matrix works quite well, and that uncertainties estimated for bounded are quite reasonable. Of course, if the best fit value is very close to a boundary, the derivative estimated uncertainty and correlations for that parameter may not be reliable.

The `MINUIT` documentation recommends caution in using bounds. Setting bounds can certainly increase the number of function evaluations (and so computation time), and in some cases may cause some instabilities, as the range of acceptable parameter values is not fully explored. On the other hand, preliminary tests suggest that using `max` and `min` to set clearly outlandish bounds does not greatly affect performance or results.

USING MATHEMATICAL CONSTRAINTS

Being able to fix variables to a constant value or place upper and lower bounds on their values can greatly simplify modeling real data. These capabilities are key to lmfit's Parameters. In addition, it is sometimes highly desirable to place mathematical constraints on parameter values. For example, one might want to require that two Gaussian peaks have the same width, or have amplitudes that are constrained to add to some value. Of course, one could rewrite the objective or model function to place such requirements, but this is somewhat error prone, and limits the flexibility so that exploring constraints becomes laborious.

To simplify the setting of constraints, Parameters can be assigned a mathematical expression of other Parameters, builtin constants, and builtin mathematical functions that will be used to determine its value. The expressions used for constraints are evaluated using the `asteval` module, which uses Python syntax, and evaluates the constraint expressions in a safe and isolated namespace.

This approach to mathematical constraints allows one to not have to write a separate model function for two Gaussians where the two `sigma` values are forced to be equal, or where amplitudes are related. Instead, one can write a more general two Gaussian model (perhaps using `GaussianModel`) and impose such constraints on the Parameters for a particular fit.

11.1 Overview

Just as one can place bounds on a Parameter, or keep it fixed during the fit, so too can one place mathematical constraints on parameters. The way this is done with lmfit is to write a Parameter as a mathematical expression of the other parameters and a set of pre-defined operators and functions. The constraint expressions are simple Python statements, allowing one to place constraints like:

```
pars = Parameters()
pars.add('frac_curve1', value=0.5, min=0, max=1)
pars.add('frac_curve2', expr='1-frac_curve1')
```

as the value of the `frac_curve1` parameter is updated at each step in the fit, the value of `frac_curve2` will be updated so that the two values are constrained to add to 1.0. Of course, such a constraint could be placed in the fitting function, but the use of such constraints allows the end-user to modify the model of a more general-purpose fitting function.

Nearly any valid mathematical expression can be used, and a variety of built-in functions are available for flexible modeling.

11.2 Supported Operators, Functions, and Constants

The mathematical expressions used to define constrained Parameters need to be valid python expressions. As you'd expect, the operators '+', '-', '*', '/', '**', are supported. In fact, a much more complete set can be used, including Python's bit- and logical operators:

```
+, -, *, /, **, &, |, ^, <<, >>, %, and, or,
==, >, >=, <, <=, !=, ~, not, is, is not, in, not in
```

The values for e (2.7182818...) and π (3.1415926...) are available, as are several supported mathematical and trigonometric function:

```
abs, acos, acosh, asin, asinh, atan, atan2, atanh, ceil,
copysign, cos, cosh, degrees, exp, fabs, factorial,
floor, fmod, frexp, fsum, hypot, isinf, isnan, ldexp,
log, log10, loglp, max, min, modf, pow, radians, sin,
sinh, sqrt, tan, tanh, trunc
```

In addition, all Parameter names will be available in the mathematical expressions. Thus, with parameters for a few peak-like functions:

```
pars = Parameters()
pars.add('amp_1', value=0.5, min=0, max=1)
pars.add('cen_1', value=2.2)
pars.add('wid_1', value=0.2)
```

The following expression are all valid:

```
pars.add('amp_2', expr='(2.0 - amp_1**2)')
pars.add('cen_2', expr='cen_1 * wid_2 / max(wid_1, 0.001)')
pars.add('wid_2', expr='sqrt(pi)*wid_1')
```

In fact, almost any valid Python expression is allowed. A notable example is that Python's 1-line *if expression* is supported:

```
pars.add('bounded', expr='param_a if test_val/2. > 100 else param_b')
```

which is equivalent to the more familiar:

```
if test_val/2. > 100:
    bounded = param_a
else:
    bounded = param_b
```

11.3 Using Inequality Constraints

A rather common question about how to set up constraints that use an inequality, say, $x + y \leq 10$. This can be done with algebraic constraints by recasting the problem, as $x + y = \delta$ and $\delta \leq 10$. That is, first, allow x to be held by the freely varying parameter x . Next, define a parameter δ to be variable with a maximum value of 10, and define parameter y as $\delta - x$:

```
pars = Parameters()
pars.add('x', value=5, vary=True)
pars.add('delta', value=5, max=10, vary=True)
pars.add('y', expr='delta-x')
```

The essential point is that an inequality still implies that a variable (here, δ) is needed to describe the constraint. The secondary point is that upper and lower bounds can be used as part of the inequality to make the definitions more convenient.

11.4 Advanced usage of Expressions in Imfit

The expression used in a constraint is converted to a Python [Abstract Syntax Tree](#), which is an intermediate version of the expression – a syntax-checked, partially compiled expression. Among other things, this means that Python’s own parser is used to parse and convert the expression into something that can easily be evaluated within Python. It also means that the symbols in the expressions can point to any Python object.

In fact, the use of Python’s AST allows a nearly full version of Python to be supported, without using Python’s built-in `eval()` function. The `asteval` module actually supports most Python syntax, including for- and while-loops, conditional expressions, and user-defined functions. There are several unsupported Python constructs, most notably the class statement, so that new classes cannot be created, and the import statement, which helps make the `asteval` module safe from malicious use.

One important feature of the `asteval` module is that you can add domain-specific functions into it, for later use in constraint expressions. To do this, you would use the `asteval` attribute of the `Minimizer` class, which contains a complete AST interpreter. The `asteval` interpreter uses a flat namespace, implemented as a single dictionary. That means you can preload any Python symbol into the namespace for the constraints:

```
def mylorentzian(x, amp, cen, wid):
    "lorentzian function: wid = half-width at half-max"
    return (amp / (1 + ((x-cen) / wid)**2))

fitter = Minimizer()
fitter.asteval.symtable['lorentzian'] = mylorentzian
```

and this `lorentzian()` function can now be used in constraint expressions.

RELEASE NOTES

This section discusses changes between versions, especially changes significant to the use and behavior of the library. This is not meant to be a comprehensive list of changes. For such a complete record, consult the [lmfit GitHub repository](#).

12.1 Version 0.9.10 Release Notes

Two new global algorithms were added: `basinhopping` and `AMPGO`. `Basinhopping` wraps the method present in `scipy`, and more information can be found in the documentation ([`basinhopping\(\)`](#) and `scipy.optimize.basinhopping`). The Adaptive Memory Programming for Global Optimization (AMPGO) algorithm was adapted from Python code written by [Andrea Gavana](#). A more detailed explanation of the algorithm is available in the [AMPGO paper](#) and specifics for `lmfit` can be found in the `ampgo()` function.

`lmfit` uses the external uncertainties (<https://github.com/lebigot/uncertainties>) package (available on PyPI), instead of distributing its own fork.

An `AbortFitException` is now raised when the fit is aborted by the user (i.e., by using `iter_cb`).

Bugfixes:

- all exceptions are allowed when trying to import `matplotlib`
- simplify and fix corner-case errors when testing closeness of large integers

12.2 Version 0.9.9 Release Notes

`lmfit` now uses the `asteval` (<https://github.com/newville/asteval>) package instead of distributing its own copy. The minimum required `asteval` version is 0.9.12, which is available on PyPI. If you see import errors related to `asteval`, please make sure that you actually have the latest version installed.

12.3 Version 0.9.6 Release Notes

Support for SciPy 0.14 has been dropped: SciPy 0.15 is now required. This is especially important for `lmfit` maintenance, as it means we can now rely on SciPy having code for differential evolution and do not need to keep a local copy.

A brute force method was added, which can be used either with `Minimizer.brute()` or using the `method='brute'` option to `Minimizer.minimize()`. This method requires finite bounds on all varying parameters, or that parameters have a finite `brute_step` attribute set to specify the step size.

Custom cost functions can now be used for the scalar minimizers using the *reduce_fcn* option.

Many improvements to documentation and docstrings in the code were made. As part of that effort, all API documentation in this main Sphinx documentation now derives from the docstrings.

Uncertainties in the resulting best-fit for a model can now be calculated from the uncertainties in the model parameters.

Parameters have two new attributes: *brute_step*, to specify the step size when using the *brute* method, and *user_data*, which is unused but can be used to hold additional information the user may desire. This will be preserved on copy and pickling.

Several bug fixes and cleanups.

Versioneer was updated to 0.18.

Tests can now be run either with nose or pytest.

12.4 Version 0.9.5 Release Notes

Support for Python 2.6 and SciPy 0.13 has been dropped.

12.5 Version 0.9.4 Release Notes

Some support for the new *least_squares* routine from SciPy 0.17 has been added.

Parameters can now be used directly in floating point or array expressions, so that the Parameter value does not need *sigma = params['sigma'].value*. The older, explicit usage still works, but the docs, samples, and tests have been updated to use the simpler usage.

Support for Python 2.6 and SciPy 0.13 is now explicitly deprecated and will be dropped in version 0.9.5.

12.6 Version 0.9.3 Release Notes

Models involving complex numbers have been improved.

The *emcee* module can now be used for uncertainty estimation.

Many bug fixes, and an important fix for performance slowdown on getting parameter values.

ASV benchmarking code added.

12.7 Version 0.9.0 Release Notes

This upgrade makes an important, non-backward-compatible change to the way many fitting scripts and programs will work. Scripts that work with version 0.8.3 will not work with version 0.9.0 and vice versa. The change was not made lightly or without ample discussion, and is really an improvement. Modifying scripts that did work with 0.8.3 to work with 0.9.0 is easy, but needs to be done.

12.7.1 Summary

The upgrade from 0.8.3 to 0.9.0 introduced the `MinimizerResult` class (see [MinimizerResult – the optimization result](#)) which is now used to hold the return value from `minimize()` and `Minimizer.minimize()`. This returned object contains many goodness of fit statistics, and holds the optimized parameters from the fit. Importantly, the parameters passed into `minimize()` and `Minimizer.minimize()` are no longer modified by the fit. Instead, a copy of the passed-in parameters is made which is changed and returns as the `params` attribute of the returned `MinimizerResult`.

12.7.2 Impact

This upgrade means that a script that does:

```
my_pars = Parameters()
my_pars.add('amp', value=300.0, min=0)
my_pars.add('center', value= 5.0, min=0, max=10)
my_pars.add('decay', value= 1.0, vary=False)

result = minimize(objfunc, my_pars)
```

will still work, but that `my_pars` will **NOT** be changed by the fit. Instead, `my_pars` is copied to an internal set of parameters that is changed in the fit, and this copy is then put in `result.params`. To look at fit results, use `result.params`, not `my_pars`.

This has the effect that `my_pars` will still hold the starting parameter values, while all of the results from the fit are held in the `result` object returned by `minimize()`.

If you want to do an initial fit, then refine that fit to, for example, do a pre-fit, then refine that result different fitting method, such as:

```
result1 = minimize(objfunc, my_pars, method='nelder')
result1.params['decay'].vary = True
result2 = minimize(objfunc, result1.params, method='leastsq')
```

and have access to all of the starting parameters `my_pars`, the result of the first fit `result1`, and the result of the final fit `result2`.

12.7.3 Discussion

The main goal for making this change were to

1. give a better return value to `minimize()` and `Minimizer.minimize()` that can hold all of the information about a fit. By having the return value be an instance of the `MinimizerResult` class, it can hold an arbitrary amount of information that is easily accessed by attribute name, and even be given methods. Using objects is good!
2. To limit or even eliminate the amount of “state information” a `Minimizer` holds. By state information, we mean how much of the previous fit is remembered after a fit is done. Keeping (and especially using) such information about a previous fit means that a `Minimizer` might give different results even for the same problem if run a second time. While it’s desirable to be able to adjust a set of `Parameters` re-run a fit to get an improved result, doing this by changing an internal attribute (`Minimizer.params`) has the undesirable side-effect of not being able to “go back”, and makes it somewhat cumbersome to keep track of changes made while adjusting parameters and re-running fits.

PYTHON MODULE INDEX

I

`lmfit.confidence`, [105](#)
`lmfit.minimizer`, [21](#)
`lmfit.model`, [47](#)
`lmfit.models`, [77](#)
`lmfit.parameter`, [14](#)

A

add() (Parameters method), 17
 add_many() (Parameters method), 18
 aic (in module lmfit.model), 66
 aic (MinimizerResult attribute), 29
 ampgo() (Minimizer method), 36

B

basinhopping() (Minimizer method), 36
 best_fit (in module lmfit.model), 66
 best_values (in module lmfit.model), 66
 bic (in module lmfit.model), 66
 bic (MinimizerResult attribute), 29
 BreitWignerModel (class in lmfit.models), 82
 brute() (Minimizer method), 35

C

chisqr (in module lmfit.model), 66
 chisqr (MinimizerResult attribute), 28
 ci_out (in module lmfit.model), 66
 ci_report() (in module lmfit), 112
 ci_report() (ModelResult method), 62
 Composite models, 69
 CompositeModel (class in lmfit.model), 72
 conf_interval() (in module lmfit), 110
 conf_interval() (ModelResult method), 62
 conf_interval2d() (in module lmfit), 112
 ConstantModel (class in lmfit.models), 87
 correl (Parameter attribute), 16
 covar (in module lmfit.model), 66
 covar (MinimizerResult attribute), 28

D

DampedHarmonicOscillatorModel (class in lmfit.models), 84
 DampedOscillatorModel (class in lmfit.models), 84
 data (in module lmfit.model), 66
 DonaichModel (class in lmfit.models), 86
 dump() (Parameters method), 19
 dumps() (Parameters method), 18

E

emcee() (Minimizer method), 37
 errorbars (in module lmfit.model), 66
 errorbars (MinimizerResult attribute), 28
 eval() (Model method), 51
 eval() (ModelResult method), 61
 eval_components() (ModelResult method), 61
 eval_uncertainty() (ModelResult method), 62
 ExponentialGaussianModel (class in lmfit.models), 85
 ExponentialModel (class in lmfit.models), 91
 ExpressionModel (class in lmfit.models), 93

F

fit() (Model method), 52
 fit() (ModelResult method), 61
 fit_report() (in module lmfit.printfuncs), 43
 fit_report() (ModelResult method), 62
 flatchain (MinimizerResult attribute), 29
 func (in module lmfit.model), 54

G

GaussianModel (class in lmfit.models), 77
 guess() (Model method), 53

I

ier (in module lmfit.model), 66
 ier (MinimizerResult attribute), 28
 independent_vars (in module lmfit.model), 54
 init_fit (in module lmfit.model), 66
 init_params (in module lmfit.model), 66
 init_vals (MinimizerResult attribute), 28
 init_values (in module lmfit.model), 66
 init_values (MinimizerResult attribute), 28
 iter_cb (in module lmfit.model), 66

J

jacfcn (in module lmfit.model), 67

L

least_squares() (Minimizer method), 33
 leastsq() (Minimizer method), 33

LinearModel (class in lmfit.models), 88
lmdif_message (in module lmfit.model), 67
lmdif_message (MinimizerResult attribute), 28
lmfit.confidence (module), 105
lmfit.minimizer (module), 21
lmfit.model (module), 47
lmfit.models (module), 77
lmfit.parameter (module), 14
load() (Parameters method), 19
load_model() (in module lmfit.model), 59
load_modelresult() (in module lmfit.model), 68
loads() (Parameters method), 19
LognormalModel (class in lmfit.models), 83
LorentzianModel (class in lmfit.models), 78

M

make_params() (Model method), 53
message (in module lmfit.model), 67
message (MinimizerResult attribute), 28
method (in module lmfit.model), 67
minimize() (in module lmfit.minimizer), 23
minimize() (Minimizer method), 32
Minimizer (class in lmfit.minimizer), 30
MinimizerResult (class in lmfit.minimizer), 27
Model (class in lmfit.model), 50
model (in module lmfit.model), 67
ModelResult (class in lmfit.model), 60
MoffatModel (class in lmfit.models), 80

N

name (in module lmfit.model), 54
nan_policy (in module lmfit.model), 54
ndata (in module lmfit.model), 67
ndata (MinimizerResult attribute), 28
nfev (in module lmfit.model), 67
nfev (MinimizerResult attribute), 28
nfree (in module lmfit.model), 67
nfree (MinimizerResult attribute), 28
nvarys (in module lmfit.model), 67
nvarys (MinimizerResult attribute), 28

O

opts (in module lmfit.model), 54

P

param_hints (in module lmfit.model), 54
param_names (in module lmfit.model), 54
Parameter (class in lmfit.parameter), 15
Parameters (class in lmfit.parameter), 17
params (in module lmfit.model), 67
params (MinimizerResult attribute), 28
Pearson7Model (class in lmfit.models), 81
plot() (ModelResult method), 63

plot_fit() (ModelResult method), 64
plot_residuals() (ModelResult method), 65
PolynomialModel (class in lmfit.models), 89
PowerLawModel (class in lmfit.models), 92
prefix (in module lmfit.model), 54
prepare_fit() (Minimizer method), 34
pretty_print() (Parameters method), 18
print_param_hints() (Model method), 54
PseudoVoigtModel (class in lmfit.models), 80

Q

QuadraticModel (class in lmfit.models), 88

R

RectangleModel (class in lmfit.models), 90
redchi (in module lmfit.model), 67
redchi (MinimizerResult attribute), 28
Removing a Constraint Expression, 16
residual (in module lmfit.model), 67
residual (MinimizerResult attribute), 28

S

save_model() (in module lmfit.model), 59
save_modelresult() (in module lmfit.model), 68
scalar_minimize() (Minimizer method), 33
scale_covar (in module lmfit.model), 67
set() (Parameter method), 16
set_param_hint() (Model method), 53
show_candidates() (MinimizerResult method), 29
SkewedGaussianModel (class in lmfit.models), 86
status (MinimizerResult attribute), 28
stderr (Parameter attribute), 16
StepModel (class in lmfit.models), 90
StudentsTModel (class in lmfit.models), 82
success (in module lmfit.model), 67
success (MinimizerResult attribute), 28

V

valuesdict() (Parameters method), 18
var_names (MinimizerResult attribute), 28
VoigtModel (class in lmfit.models), 79

W

weights (in module lmfit.model), 67